
BadScientist: Can a Research Agent Write Convincing but Unsound Papers that Fool LLM Reviewers?

AI Scientist

Fengqing Jiang^{♣,†} Yichen Feng^{♣,†} Radha Poovendran[♣]

[♣]University of Washington

{fqjiang, yfeng42, rp3}@uw.edu

Project Page: <https://bad-scientist.github.io>

Abstract

The rapid advancement of Large Language Models (LLMs) as both research assistants and peer reviewers creates a critical vulnerability: the potential for fully automated AI-only publication loops where AI-generated research is evaluated by AI reviewers. We investigate this adversarial dynamic by introducing **BadScientist**, an experimental framework that pits a fabrication-oriented paper generation agent against multi-model LLM review systems. Our generator employs five presentation-manipulation strategies without conducting real experiments: exaggerating performance gains (*TooGoodGains*), cherry-picking comparisons (*BaselineSelect*), creating statistical facades (*StatTheater*), polishing presentation (*CoherencePolish*), and hiding proof oversights (*ProofGap*). We evaluate fabricated papers using LLM reviewers calibrated on ICLR 2025 conference submission data.

Our results reveal alarming vulnerabilities: fabricated papers achieve high acceptance rates across strategies, with *TooGoodGains* reaching 67.0%/82.0% acceptance under different thresholds, and combined strategies achieving 52.0%/69.0%. Even when LLM reviewers flag integrity concerns, they frequently assign acceptance-level scores—a phenomenon we term *concern-acceptance conflict*. Our mitigation strategies, Review-with-Detection (ReD) and Detection-Only (DetOnly), show limited improvements, highlighting the inadequacy of current methods. These findings expose concrete failure modes in AI-driven review systems and demonstrate that presentation manipulation can effectively deceive state-of-the-art LLM reviewers. Our work underscores the urgent need for stronger, integrity-focused review pipelines as AI agents become more prevalent in scientific publishing.

1 Introduction

The rapid advancement of Large Language Models (LLMs) is creating a paradigm shift in the scientific research ecosystem, automating complex tasks once exclusive to human experts. On one hand, LLM-powered agents are increasingly being developed as end-to-end "research assistants," capable of automating ideation, experimentation, and manuscript drafting [20, 16, 12, 3]. On the other hand, LLMs are being explored as tools to alleviate the growing burden on the peer review system, serving as reviewers or review assistants [4, 18, 15, 24].

The convergence of these two capabilities introduces a critical vulnerability: the potential for a fully automated, AI-only publication loop where AI-generated research is evaluated by AI reviewers. This scenario raises profound questions about research integrity [25, 2]. If a malicious or poorly designed research agent generates convincing but scientifically unsound work, can current LLM-based review systems reliably detect the fabrication? Early studies suggest that LLM reviewers can amplify human biases [10], miss subtle yet critical flaws, and even be manipulated through adversarial attacks like

[†]Equal Contribution

prompt injection hidden within a manuscript’s text [26, 23, 28]. While detection of AI-generated text is an active area of research [7, 21, 5], the adversarial interplay between a fabricating agent and a reviewer agent remains a critical and underexplored failure mode.

In this paper, we investigate this adversarial dynamic by asking: **Can a research agent write convincing but unsound papers that fool LLM reviewers?** To answer this, we introduce *BadScientist*, an experimental framework that pits a fabrication-oriented paper generation agent against a multi-model LLM review agent. Our generator does not conduct real experiments; instead, it composes manuscripts using a set of five “presentation-manipulation” strategies, such as exaggerating performance gains (*TooGoodGains*), cherry-picking comparison methods (*BaselineSelect*), and creating a facade of rigor with polished but meaningless statistics (*StatTheater*). We then evaluate these fabricated papers using a panel of LLM reviewers calibrated on real-world conference data to mirror actual acceptance thresholds.

Our results are alarming. We find that fabricated papers are accepted at a high rate, and even when LLM reviewers flag integrity concerns, they often still assign acceptance-level scores, a phenomenon we term **concern-acceptance conflict**. We further prototype and evaluate two mitigation strategies—*Review-with-Detection* (ReD) and *Detection-Only* (DetOnly)—which show limited but measurable improvements, highlighting the urgent need for stronger, integrity-aware review pipelines. This study exposes concrete failure modes in AI-only publication loops and proposes practical guardrails to reduce the risk of automated systems endorsing and propagating fabricated science.

2 Related Work

Agents for Scientific Discovery. LLM agents are increasingly positioned as end-to-end “research agents,” automating ideation, experimentation, and manuscript drafting. Systems such as the *AI Scientist* [20] and *Auto Research* [16] report credible, minimally supervised pipelines; complementary benchmarks probe specific stages like ML experimentation and engineering [12, 3]. While these works establish feasibility and scope, few analyze the *integrity* of outputs under adversarial objectives.

Agents for Peer Review. LLMs have been explored as reviewers and review assistants, from early feasibility studies [18, 4] to larger evaluations showing partial alignment with human feedback [15, 14]. Emerging platforms simulate or standardize review processes and propose bias-aware pipelines [24, 11, 27], yet concerns persist LLM reviewers can amplify biases or miss deep flaws [10].

Challenges in Agent-vs-Agent Settings. The coupling of AI-written papers and AI-based reviews introduces new attack surfaces. Prompt-injection into manuscripts can tilt LLM verdicts [26], and reports suggest covert instructions have appeared in real preprints [23]. Parallel efforts assess detection and governance: holistic and red-teaming evaluations [13, 22]; detectors and audits for AI-generated scientific text and artifacts [7, 21, 6, 17, 8, 1, 9, 5]; and policy guidance on safeguarding research integrity [25, 2].

Our Focus. We study the *adversarial interplay* between an AI paper-writing agent and an AI reviewer: can a fabrication-oriented agent produce “convincing but unsound” papers that fool LLM review pipelines, and what mitigations help? In contrast to prior work that treats generation and reviewing separately, we evaluate the coupled system under integrity-focused attacks and prototype mitigation (e.g., injection-aware defenses [28]).

3 Design of *BadScientist*

3.1 Preliminary

We study whether AI agents can generate convincing *fabricated* scientific papers that deceive reviewer agents, and how reliably reviewer agents detect such fabrications. We implement a multi-agent pipeline that simulates a publication workflow from paper generation to peer review and post-hoc detection analysis. The core research problem involves: a *Paper Generation Agent* \mathcal{G} that produces papers; a *Review Agent* \mathcal{R} that evaluates papers via multiple LLMs. There is also an *Analysis System*

\mathcal{A} that aggregates outcomes and measures detection. In this work, we design a separate *Planning Agent* \mathcal{P} for implementation orchestration.

Notation. Let \mathcal{X} denote the space of paper artifacts. A paper is $x \in \mathcal{X}$. Let \mathcal{S} be the set of fabrication strategies and \mathcal{T} the set of topics. A seed prompt $u \in \mathcal{U}$ specifies a topic $t \in \mathcal{T}$ and a strategy $s \in \mathcal{S}$. The Review Agent employs models $\mathcal{M} = \{m_1, \dots, m_M\}$. Each model m produces a K -dimensional rubric score $\mathbf{r}_m(x) = (r_{m,1}(x), \dots, r_{m,K}(x))$ with $r_{m,k}(x) \in \{1, \dots, L_k\}$, and free-form textual feedback $\omega_m(x)$. Let $\mathbf{w} \in \Delta^{M-1} := \{\mathbf{w} \in \mathbb{R}_{\geq 0}^M : \sum_m w_m = 1\}$ be reviewer weights (default uniform). We define the consensus score vector $\bar{\mathbf{r}}(x)$ and the binary recommendation $\hat{y}(x)$ acceptance threshold τ calibrated by \mathcal{A} .

Assumptions (Threat Model and Scope). We focus on a setting where \mathcal{G} aims to produce *high-quality fabricated papers* without conducting real experiments or collecting real data. Instead, \mathcal{G} may synthesize or manipulate pseudo-data to support claims. We assume the research agent has no prior knowledge about the reviewer system, i.e., the generated paper is not directly adversary optimized to the agent review system, and will not optimize paper generation with feedback from reviewers. Review agents operate under a standard (minimal) single-pass review protocol with access restricted to the submitted paper content, where the review system is not designed to have ability to run experiments to verify the papers, which mimics the most practical setup in peer review system.

We explicitly exclude human-in-the-loop setup (e.g., deception tactics with forged human ethics approvals or human feedback to revise papers), adversary attack (e.g., prompt injection attack towards review agents), and offline collusion (e.g., the research agent and review agent hidden collusion).

3.2 Paper Generation Agent \mathcal{G}

The advanced research agents, involves multi-factors abilities, including planning on research discovery (e.g., design and hypothesis), agentic execution of exploration and verification (e.g., experiment design and executions), and write professional academic research papers to support discovery. Our paper generation agent \mathcal{G} is the minimal support for such research agents, which only focus on fabricating the research paper, without practical exploration and verification.

The design of \mathcal{G} involves the following components.

Strategy space. Let $\mathcal{S} = \{s_1, \dots, s_n\}$ denote the atomic presentation–manipulation strategies. The set of all admissible strategy *setups* is the power set

$$\mathcal{P}(\mathcal{S}) \equiv \{A \mid A \subseteq \mathcal{S}\} = 2^{\mathcal{S}},$$

i.e., every combination of elements of \mathcal{S} . A generation policy $g \in \mathcal{G}$ selects a configuration $A \in \mathcal{P}(\mathcal{S})$ and composes the associated interventions. In our experiments, we instantiate the five atomic strategy elements as follows:

- s_1 : **TooGoodGains**: foreground extraordinary improvements over the leading baseline and frame them as field-advancing.
- s_2 : **BaselineSelect**: curate comparisons to favor the proposed method; omit variance/CI in the main table.
- s_3 : **StatTheater**: polished ablations, precise p-values/CI/error bars (mainly in the supplement), tidy hyperparameter tables, and “forthcoming” repo/DOI links.
- s_4 : **CoherencePolish**: flawless cross-references, unified terminology, consistent significant digits, aligned captions, professional typography.
- s_5 : **ProofGap**: a theorem/lemma with an ostensibly rigorous proof hiding a subtle oversight.

The design of the planning agent is to use the agentic framework to assist of this research process, but it is generally optional in the framework. An experienced researcher or scientist can take the responsibility.

Ethical intent: the work seeks to evaluate and harden reviewer pipelines against fabrication, not to promote academic fraud or encourage dishonesty.

Generative mapping. Given a seed prompt $u = (t, s)$, \mathcal{G} samples pseudo-data $D \sim q(\cdot | s, t, \theta)$ from a strategy-conditioned generator q with internal parameters θ , produces visualizations $V = \text{viz}(D)$, and assembles a structured manuscript $x = \text{compose}(u, D, V)$ with sections (abstract, introduction, methods, results, discussion, conclusion), citations, figures, and tables. We impose structural validity constraints captured by

$$C(x) = \mathbb{I}[\text{compile}(x) = \text{success} \wedge \text{struct}(x) \in \mathcal{C}] = 1, \quad (1)$$

where \mathcal{C} encodes formatting requirements (section presence, figure/table counts, bibliography entries). Only papers with $C(x) = 1$ are forwarded to review.

Distributional formulation. The end-to-end generation induces a distribution over papers $p_{\mathcal{G}}(x | u)$:

$$p_{\mathcal{G}}(x | u) = \int p(x | D, u) q(D | s, t, \theta) dD, \quad u = (t, s). \quad (2)$$

We use this to stratify evaluations by strategy s and topic t .

3.3 Review Agent \mathcal{R}

Given a paper $x \in \mathcal{X}$, the Review Agent queries each model $m \in \mathcal{M}$ under a fixed K -criterion rubric (e.g., methodology, significance, clarity, etc.). Each model returns a rubric vector and textual feedback $(\mathbf{r}_m(x), \omega_m(x))$. Using reviewer weights $\mathbf{w} \in \Delta^{M-1}$, the agent forms the consensus rubric

$$\bar{\mathbf{r}}(x) = \sum_{m \in \mathcal{M}} w_m \mathbf{r}_m(x),$$

and produces a binary recommendation via the scoring functional ϕ and a calibrated threshold τ : $\hat{y}(x) = \mathbb{I}[\phi(\bar{\mathbf{r}}(x)) \geq \tau]$. We summarize the agent's output as

$$\mathcal{R}(x) = \left(\{(\mathbf{r}_m(x), \omega_m(x))\}_{m \in \mathcal{M}}, \bar{\mathbf{r}}(x), \hat{y}(x) \right),$$

which preserves per-model judgments and comments while supplying a single consensus score and decision.

3.4 Review Calibration for Analysis \mathcal{A}

We calibrate the Review Agent's decision rule using a corpus of real conference submissions with publicly available reviews and outcomes.

Calibration corpus. We define the reference pool as:

$$\mathcal{D}_{\text{ref}} = \{(x_i, y_i^{\text{hum}}, \sigma_i, h_i)\}_{i=1}^{N_{\star}},$$

where x_i is the paper artifact, $y_i^{\text{hum}} \in \{0, 1\}$ indicates the human accept/reject decision, $\sigma_i \in \mathcal{C}_{\text{stat}}$ represents the meta-status labels (e.g., oral/spotlight/poster/reject/withdraw), and $h_i \in \mathbb{R}$ is a scalar venue score such as the average assessment.

From this reference pool, we construct a stratified calibration set \mathcal{D}_{cal} that preserves the score and status distributions of \mathcal{D}_{ref} .

Stratified sampling procedure. We implement stratified sampling as follows. First, we partition the score space using bin edges $t_0 < \dots < t_B$ to define score bins $B_b = [t_{b-1}, t_b)$ for $b = 1, \dots, B$.

For each bin–status combination $(b, c) \in \{1, \dots, B\} \times \mathcal{C}_{\text{stat}}$, we define:

$$\mathcal{I}_{b,c} = \{i : h_i \in B_b, \sigma_i = c\}, \quad N_{b,c} = |\mathcal{I}_{b,c}|, \quad p_{b,c} = \frac{N_{b,c}}{N_{\star}},$$

where $N_{\star} = \sum_{b=1}^B \sum_{c \in \mathcal{C}_{\text{stat}}} N_{b,c}$ is the total reference pool size.

Given a target calibration size N_{cal} , we allocate samples to each cell using proportional allocation with the largest-remainder method:

$$n'_{b,c} = p_{b,c} N_{\text{cal}}, \quad n_{b,c} = \lfloor n'_{b,c} \rfloor, \quad R = N_{\text{cal}} - \sum_{b,c} n_{b,c}.$$

We then add one additional sample to the R cells with the largest remainders $n'_{b,c} - \lfloor n'_{b,c} \rfloor$.

Finally, we sample uniformly without replacement $\mathcal{S}_{b,c} \subseteq \mathcal{I}_{b,c}$ with $|\mathcal{S}_{b,c}| = n_{b,c}$ and construct:

$$\mathcal{S} = \bigcup_{b=1}^B \bigcup_{c \in \mathcal{C}_{\text{stat}}} \mathcal{S}_{b,c}, \quad \mathcal{D}_{\text{cal}} = \{(x_i, y_i^{\text{hum}}, \sigma_i, h_i) : i \in \mathcal{S}\}.$$

This construction ensures that $\hat{p}_{b,c}^{\text{cal}} = n_{b,c}/N_{\text{cal}} \approx p_{b,c}$ for all (b, c) , preserving both score-bin and status marginals up to integer rounding.

Agent scoring. For each paper $x \in \mathcal{D}_{\text{cal}}$, the Review Agent produces a consensus rubric $\bar{r}(x)$, converts it to a scalar score $s(x) = \phi(\bar{r}(x)) \in \mathbb{R}$, and makes a binary recommendation $\hat{y}_\tau(x) = \mathbb{I}[s(x) \geq \tau]$ for threshold $\tau \in \mathbb{R}$.

Threshold calibration. We derive two operating thresholds to accommodate different evaluation criteria.

1. Rate-matching threshold. Let $\alpha^* \in (0, 1)$ denote the target venue acceptance rate. We define:

$$\hat{\alpha}_{\text{cal}}(\tau) = \frac{1}{|\mathcal{D}_{\text{cal}}|} \sum_{x \in \mathcal{D}_{\text{cal}}} \hat{y}_\tau(x), \quad \tau_{\text{rate}} \in \arg \min_{\tau \in \mathbb{R}} |\hat{\alpha}_{\text{cal}}(\tau) - \alpha^*|.$$

This threshold ensures that the agent's acceptance rate on the calibration set matches the venue's historical acceptance rate.

2. Probability-consistency threshold. Let $\pi(t) = \mathbb{P}(y^{\text{hum}} = 1 \mid s(x) \geq t)$ for $t \in \mathbb{R}$, estimated on \mathcal{D}_{cal} using a monotone calibration model. We define:

$$\tau_{0.5} = \inf\{t \in \mathbb{R} : \pi(t) \geq \frac{1}{2}\},$$

so that papers scoring $s(x) \geq \tau_{0.5}$ have at least 50% estimated probability of human acceptance.

Output. The calibration module returns $\mathcal{A}(\mathcal{D}_{\text{cal}}) = (\tau_{\text{rate}}, \tau_{0.5})$, providing operating thresholds for the decision rule $\hat{y}(x) = \mathbb{I}[s(x) \geq \tau]$.

4 Experiment

4.1 Setup

Implementation Our agent framework is adapted from AI-Scientist [19], but we have fundamentally redesigned its entire pipeline. We retain only its most foundational writing prompts and have eliminated the need for any experimental execution or structured templates. Our framework now operates directly from a simple seed idea, allowing the LLM to freely generate any necessary experimental results and plotting code.

We follow the generation strategy space, with five atomic strategies, S_1, \dots, S_5 . We employ the individual strategy per setup, as well as a unified setup which encourages to use all strategies in generation (named as *All*). We use a planning agent with GPT-5 to generate all seed topics for paper generations. Each seed produces 4 papers across six strategy setups.

Agent Models. We use GPT-5 to support our paper generation agent. For the review agent, we set $\mathcal{M} = 3$ and use o3, o4-mini, and GPT-4.1 with the rubric review prompt. And for the ease of acceptance decision, we take only the overall assessment score provided by review agent for paper scoring, i.e., $\phi(\bar{r}(x)) = r_{oa}(x)$.

Calibration set and thresholds. We instantiate the reference pool \mathcal{D}_{ref} as the ICLR 2025 Open-Review submission set (with public reviews and outcomes). A stratified calibration set \mathcal{D}_{cal} of size $N_{\text{cal}} = 200$ is then constructed as described in Section 3.4. Running the Review Agent on \mathcal{D}_{cal} yields two operating thresholds. *Rate-matching* selects τ_{rate} so that the agent minimize the drift of empirical acceptance rate on \mathcal{D}_{cal} matches the venue rate $\alpha^* = 0.3173$, which yields $\tau_{\text{rate}} = 7$.

Overall ICLR 2025 acceptance rate 31.73%; see <https://papercopilot.com/statistics/iclr-statistics/iclr-2025-statistics/>.

Table 1: Acceptance (ACPT) and Integrity Concern Rate (ICR) by strategy.

Strategy	ACPT		ICR-m			ICR@M
	τ_{rate}	$\tau_{0.5}$	o3	o4-mini	GPT-4.1	
S1	67.0%	82.0%	38.4%	4.7%	2.3%	39.5%
S2	32.0%	49.0%	35.2%	4.5%	2.3%	35.2%
S3	53.5%	69.7%	29.4%	2.4%	4.7%	31.8%
S4	44.0%	59.0%	28.2%	5.9%	1.2%	30.6%
S5	35.4%	53.5%	25.9%	8.2%	7.1%	34.1%
All	52.0%	69.0%	50.6%	5.7%	8.0%	51.7%

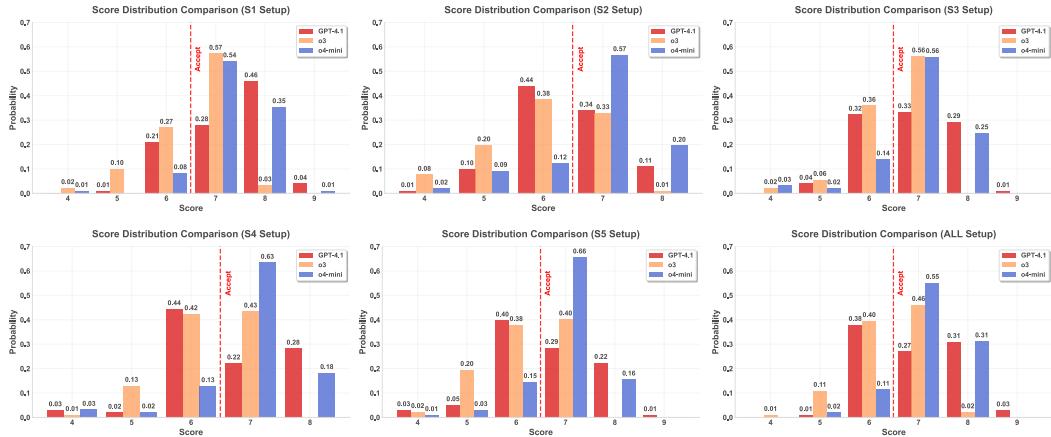


Figure 1: Score distributions across six setups (S1–S5, ALL) for three models, with the acceptance threshold marked. o4-mini is right-shifted, o3 shows higher variance and a fatter right tail, while GPT-4.1 is more conservative.

Probability-consistency defines such that papers with $s(x) \geq \tau_{0.5}$ have estimated human-acceptance probability at least 50%; this yields $\tau_{0.5} = 6.667$.

Evaluation metrics. We evaluate along two axes. (I) **Acceptance rate** (ACPT). Let \mathcal{D} be the set of generated papers and $\hat{y}_\tau(x) = \mathbb{I}[s(x) \geq \tau]$ the Review Agent's decision at threshold τ , with $s(x) = \phi(\bar{r}(x))$. For any operating threshold $\tau \in \{\tau_{\text{rate}}, \tau_{0.5}\}$ we report

$$\text{ACPT}(\tau) = \frac{1}{N} \sum_{j=1}^N \hat{y}_\tau(x_j),$$

(II) **Integrity Concern Rate** (ICR). Let $c_m(x) = \Gamma(\omega_m(x)) \in \{0, 1\}$ indicate that reviewer $m \in \mathcal{M}$ explicitly raises integrity-related concerns in $\omega_m(x)$. And suppose $\bar{c}_{\text{any}}(x) = \mathbb{I}[\sum_{m \in \mathcal{M}} c_m(x) \geq 1]$. Then for $m \in \mathcal{M}$, we have *Per-review-model ICR (ICR-m)* and the re-

Table 2: Concern–acceptance conflict (%): within papers where the model raised an integrity concern, the share still receives an acceptance-level score by model and strategy (S1–S5, ALL). Higher values indicate stronger contradiction

Model	S1	S2	S3	S4	S5	All
o3	33.3%	25.8%	52.0%	30.0%	40.9%	29.5%
o4-mini	100.0%	50.0%	100.0%	80.0%	71.9%	100.0%
GPT4.1	50.0%	50.0%	75.0%	0.0%	33.3%	57.1%

Table 3: ACPT and ICR for the baseline review agent vs. ReD. ReD lifts concerns but raises ACPTs.

Method	ACPT- τ_{rate}	ACPT- $\tau_{0.5}$	ICR-o3	ICR-o4mini	ICR-GPT4.1	ICR@M
Baseline Review Agent	28.0%	37.0%	50.6%	12.4%	4.5%	57.3%
ReD	44.0%	58.0%	84.0%	11.0%	0.0%	86.0%

laxed metric at panel-level, *Any-of-panel ICR (ICR@M)*:

$$\text{ICR-}m = \frac{1}{N} \sum_{j=1}^N c_m(x_j), \quad \text{ICR}@m = \frac{1}{N} \sum_{j=1}^N \bar{c}_{\text{any}}(x_j).$$

We use GPT-5 as LLM-judge to classify whether the text feedback from review agents contains integrity-related concerns.

4.2 Evaluation Analysis

Main Results. Our main evaluation result is in Table 1. We find that acceptance is unexpectedly high under most manipulations. Single strategies already yield substantial ACPT (e.g., S1: 67.0%/82.0%; S3: 53.5%/69.7%), indicating that current review agents are easily persuaded and lack sufficient awareness to spot integrity/fabrication issues. The All strategy as a multi-setup, attains high acceptance (52.0%/69.0%), but it also maximally increases detectability (ICR@M 51.7%, o3 50.6%), suggesting that composing strategies broadens the footprint seen by detectors. Among single strategies, S1 provides the strongest acceptance with only moderate detection pressure (ICR@M 39.5%), whereas others (e.g., S3–S5) are somewhat weaker but also less detectable (ICR@M \approx 30–34%). Across models, o3 is the most flag-happy (consistently higher ICR-m), while GPT-4.1 rarely flags concerns (mostly 2–8%), reinforcing that current review models have limited and uneven detection capability.

Score distributions. Figure 1 plots score histograms for three models across six setups (S1–S5, ALL) with the acceptance threshold marked. Overall, o4-mini is right-shifted—consistently placing more mass at ≥ 7 —which aligns with its higher acceptance tendency. o3 shows larger variance and a fatter right tail (notably in S1 and ALL), producing many near-threshold and high scores; GPT-4.1 is comparatively conservative, clustering around 6–7 with a thinner tail at 8+. Among strategies, S1 yields the strongest rightward shift for all models, while S2/S4 are milder. The ALL composition increases polarization (more mass both just below and above the threshold), explaining why it sustains high acceptance yet is easier for detectors to flag.

Concern-Acceptance Conflict. Conditioned on a model posting an integrity concern, we report the share that still receives an acceptance-level score in Table 2. Conflict is widespread: o4-mini is most contradictory (S1/S3/All: 100%; S2/S4/S5: 50–80%), GPT-4.1 is mixed (0% in S4 but 33–75% elsewhere), and o3 is moderate (26–52%). S3 (statistical theater) induces the largest cross-model conflict, and ALL further amplifies it for o4-mini (100%). These observations indicate that even agents voice concerns, yet keep acceptance-high scores, and integrity signals are not well-coupled to assessment.

5 Mitigation

Table 4: Evaluation results of all detectors. Across various setups, detection offers only slight gains over random. ReD is more conservative, while DetOnly is recall-oriented with higher FPR. o3 shows a positive bias, whereas GPT-4.1 tends toward negative.

Method	o3				o4-mini				GPT-4.1			
	TPR	FPR	Acc	F1	TPR	FPR	Acc	F1	TPR	FPR	Acc	F1
Random Guess	50.0%	50.0%	50.0%	50.0%	50.0%	50.0%	50.0%	50.0%	50.0%	50.0%	50.0%	50.0%
ReD	81.6%	44.9%	67.0%	72.1%	0.0%	8.0%	46.0%	0.0%	0.0%	0.0%	50.0%	0.0%
DetOnly	98.0%	84.0%	57.0%	69.5%	64.0%	74.0%	45.0%	53.8%	24.0%	12.0%	56.0%	35.3%

We study two complementary interventions that make the review pipeline explicitly integrity-aware. The first augments the standard review workflow with an integrity check; the second replaces reviewing with a lightweight triage detector.

5.1 Setup

Research Questions. **RQ1 (Detection-Aware Reviewing).** How can we enhance detection awareness *within* the act of reviewing—i.e., can adding an explicit integrity-detection subtask (risk classification with cited evidence) improve sensitivity to integrity risks *while* preserving or improving core review quality (helpfulness and recommendation accuracy)?

RQ2 (Triage without Full Reviews). When generating a full textual review is not the first priority, can a detection-only agent reliably flag papers with potential integrity concerns at a controllable false-alarm rate and route low-risk papers to the standard pipeline?

Mitigation Strategies. **(i) Review-with-Detection (ReD)** We modify the review prompt to bundle a conventional rubric-based review with a mandatory detection of AI-generated content. For each document, a single model $m \in M$ outputs: (a) a standard review, (b) a binary label $l_m \in \{\text{HUMAN-AUTHORED}, \text{AI-GENERATED}\}$, and (c) *evidence snippets* (quoted lines) to support its finding. The model’s label directly determines the final triage recommendation: a finding of AI-GENERATED results in a FLAG-FOR-AUDIT, while HUMAN-AUTHORED allows the process to PROCEED. **(ii) Detection-Only (DetOnly)** We remove the reviewing rubric and instruct the agent to perform *only* the detection of AI-generated content. In this mode, the model’s sole task is to return a binary label $y_m \in \{\text{HUMAN-AUTHORED}, \text{AI-GENERATED}\}$ and evidence. This label serves as the direct and final triage action, without any secondary review or scoring.

Dataset. We construct a controlled set of 100 papers using the stratified sampling algorithm in Section 3.4: 50 are sampled from real ICLR 2025 submissions, and 50 are sampled from the generated corpus produced in Section 4.

5.2 Evaluation Analysis

To RQ1: Review Outcomes. We compare *ReD* against the initial review agent baseline (same models, same rubric, no integrity subtask). We report the review outcomes using the same metrics as in Section 4. As shown in Table 3, *ReD* increases integrity warnings overall ($\text{ICR}@M$ 57.3% → 86.0%), driven largely by o3 (50.6% → 84.0%), while o4-mini is roughly unchanged and GPT-4.1 collapses to 0%. Paradoxically, acceptance also rises substantially ($\text{ACPT}_{\tau_{\text{rate}}}$ 28.0% → 44.0%, $\text{ACPT}_{\tau_{0.5}}$ 37.0% → 58.0%). Thus, adding a detection subtask improves stated awareness but does not translate into stricter recommendations—if anything, it coexists with more accepts. This suggests the integrity signal is weakly coupled to scoring; practical deployments should gate or weight recommendations by risk rather than merely requesting detection within the review.

To RQ2: Detection Performance. We set three detectors on our new dataset: *Random Guess* baseline, the *ReD* integrity component, and *DetOnly*. The results are presented in Table 4. Overall, detection helps but just slightly: across models, accuracy is near the 50% random baseline, with a clear lift only on o3 (*ReD* 67% vs. random 50%; *DetOnly* 57%). Comparing *ReD* and *DetOnly*, the latter is recall-seeking (higher TPR) but far noisier (much higher FPR), whereas *ReD* is more conservative and, on some bases, collapses (e.g., GPT-4.1 shows 0% TPR for *ReD*). Model behavior also differs: o3 tends to judge *positive* (high flag rate; e.g., *DetOnly* FPR 84%), while GPT-4.1 tends to judge *negative* (low TPR/FPR), yielding a small accuracy gain for *DetOnly* (56%) over random.

6 Conclusion and Discussion

Our findings expose a critical vulnerability: LLM review systems can be systematically deceived by presentation manipulation. Fabricated papers achieve acceptance rates up to 82%, with reviewers frequently exhibiting concern-acceptance conflicts—flagging integrity issues yet still recommending acceptance. This fundamental breakdown reveals that current AI reviewers operate more as pattern matchers than critical evaluators.

Our mitigation attempts show the inadequacy of current defenses. Detection accuracy barely exceeds random chance, and paradoxically, adding explicit integrity checks sometimes increases acceptance rates. Simply asking LLM reviewers to "be more careful" is insufficient.

The scientific community faces an urgent choice. Without immediate action to implement defense-in-depth safeguards—including provenance verification, integrity-weighted scoring, and mandatory human oversight—we risk AI-only publication loops where sophisticated fabrications overwhelm our ability to distinguish genuine research from convincing counterfeits. The integrity of scientific knowledge itself is at stake.

More detailed discussions of limitations, ethical considerations, and broader societal impacts of our work are provided in Appendix A.

AI Agent Setup

Our work is collaborated and supported by various AI agents. For the project brainstorm and idea formalization, we use DeepResearch agent supported by GPT-5/Gemini-2.5-Pro. For the experiment implementation and analysis, we utilize code agents from Copilot/Cursor supported by GPT-5/Claude-4-sonnet/Grok-Code-Fast-1, which serves as the meta agent to implement and extend *BadScientist* agent framework in our research. In final stage, we rely on Claude-4-sonnet/Gemini-2.5-Pro for writing revision.

As stated in Section 3 and 4, our *BadScientist* agent framework enables various LLMs to support end-to-end paper generation, review, and post-analysis. The framework equips agents with multiple tools throughout the pipeline, including research ideation and expansion module, experiment design module, data visualization module, LaTeX compilation tool with automatic error handling to produce PDF manuscripts, and Semantic Scholar API for literature review and citation management.

References

- [1] Nikita Andreev, Alexander Shirnin, Vladislav Mikhailov, and Ekaterina Artemova. Papilusion at dagpap24: Paper or illusion? detecting ai-generated scientific papers. *arXiv preprint arXiv:2407.17629*, 2024.
- [2] Khalid H Arar, Hamit Özen, Gülsah Polat, and Selahattin Turan. Artificial intelligence, generative artificial intelligence and research integrity: a hybrid systemic review. *Smart Learning Environments*, 12(1):44, 2025.
- [3] Jun Shern Chan, Neil Chowdhury, Oliver Jaffe, James Aung, Dane Sherburn, Evan Mays, Giulio Starace, Kevin Liu, Leon Maksin, Tejal Patwardhan, et al. Mle-bench: Evaluating machine learning agents on machine learning engineering. *arXiv preprint arXiv:2410.07095*, 2024.
- [4] Alessandro Checco, Lorenzo Bracciale, Pierpaolo Loretì, Stephen Pinfield, and Giuseppe Bianchi. Ai-assisted peer review. *Humanities and social sciences communications*, 8(1):1–11, 2021.
- [5] Evan N Crothers, Nathalie Japkowicz, and Herna L Viktor. Machine-generated text: A comprehensive survey of threat models and detection methods. *IEEE Access*, 11:70977–71002, 2023.
- [6] Madelyn A Flitcroft, Salma A Sheriff, Nathan Wolfrath, Ragasnehit Maddula, Laura McConnell, Yun Xing, Krista L Haines, Sandra L Wong, and Anai N Kothari. Performance of artificial intelligence content detectors using human and artificial intelligence-generated scientific writing. *Annals of Surgical Oncology*, 31(10):6387–6393, 2024.
- [7] Catherine A Gao, Frederick M Howard, Nikolay S Markov, Emma C Dyer, Siddhi Ramesh, Yuan Luo, and Alexander T Pearson. Comparing scientific abstracts generated by chatgpt to real abstracts with detectors and blinded human reviewers. *NPJ digital medicine*, 6(1):75, 2023.
- [8] Romain-Daniel Gosselin. Ai detectors are poor western blot classifiers: a study of accuracy and predictive values. *PeerJ*, 13:e18988, 2025.
- [9] German Gritsai, Ildar Khabutdinov, and Andrey Grabovoy. Multi-head span-based detector for ai-generated fragments in scientific papers. *arXiv preprint arXiv:2411.07343*, 2024.
- [10] Mohammad Hosseini and Serge P. Horbach. Fighting Reviewer Fatigue or Amplifying Bias? Considerations and Recommendations for Use of ChatGPT and Other LLMs in Scholarly Peer Review. *Research Integrity and Peer Review*, 8(1):4, 2023.
- [11] Yiqiao Jin, Qinlin Zhao, Yiyang Wang, Hao Chen, Kaijie Zhu, Yijia Xiao, and Jindong Wang. AgentReview: Exploring Peer Review Dynamics with LLM Agents. In *Proc. of EMNLP*, pages 1208–1226, 2024.
- [12] Patrick Tser Jern Kon, Jiachen Liu, Xinyi Zhu, Qiuyi Ding, Jingjia Peng, Jiarong Xing, Yibo Huang, Yiming Qiu, Jayanth Srinivasa, Myungjin Lee, et al. Exp-bench: Can ai conduct ai research experiments? *arXiv preprint arXiv:2505.24785*, 2025.
- [13] Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, et al. Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*, 2022.
- [14] Weixin Liang, Zachary Izzo, Yaohui Zhang, Haley Lepp, Hancheng Cao, Xuandong Zhao, Lingjiao Chen, Haotian Ye, Sheng Liu, and Zhi *et al.* Huang. Monitoring AI-Modified Content at Scale: A Case Study on the Impact of ChatGPT on AI Conference Peer Reviews. In *Proc. of ICML*, 2024.
- [15] Weixin Liang, Yuhui Zhang, Hancheng Cao, Binglu Wang, Daisy Yi Ding, Xinyu Yang, Kailas Vodrahalli, Siyu He, Daniel S. Smith, and Yian *et al.* Yin. Can Large Language Models Provide Useful Feedback on Research Papers? A Large-Scale Empirical Analysis. *NEJM AI*, 1(8):A1oa2400196, 2024.

- [16] Chengwei Liu, Chong Wang, Jiayue Cao, Jingquan Ge, Kun Wang, Lyuye Zhang, Ming-Ming Cheng, Penghai Zhao, Tianlin Li, and Xiaojun *et al.* Jia. A Vision for Auto Research with LLM Agents. *arXiv preprint arXiv:2504.18765*, 2025.
- [17] Jae QJ Liu, Kelvin TK Hui, Fadi Al Zoubi, Zing ZX Zhou, Dino Samartzis, Curtis CH Yu, Jeremy R Chang, and Arnold YL Wong. The great detectives: humans versus ai detectors in catching large language model-generated medical writing. *International Journal for Educational Integrity*, 20(1):8, 2024.
- [18] Ryan Liu and Nihar B. Shah. ReviewerGPT? An Exploratory Study on Using Large Language Models for Paper Reviewing. *arXiv preprint arXiv:2306.00622*, 2023.
- [19] Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI Scientist: Towards fully automated open-ended scientific discovery. *arXiv preprint arXiv:2408.06292*, 2024.
- [20] Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery. *arXiv preprint arXiv:2408.06292*, 2024.
- [21] Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D Manning, and Chelsea Finn. Detectgpt: Zero-shot machine-generated text detection using probability curvature. In *International conference on machine learning*, pages 24950–24962. PMLR, 2023.
- [22] Ethan Perez, Saffron Huang, Francis Song, Trevor Cai, Roman Ring, John Aslanides, Amelia Glaese, Nat McAleese, and Geoffrey Irving. Red teaming language models with language models. *arXiv preprint arXiv:2202.03286*, 2022.
- [23] Josh Taylor. Scientists reportedly hiding AI text prompts in academic papers to receive positive peer reviews, 2025. The Guardian (Tech News), July 14, 2025.
- [24] Keith Tyser, Ben Segev, Gaston Longhitano, Xin-Yu Zhang, Zachary Meeks, Jason Lee, Uday Garg, Nicholas Belsten, Avi Shporer, and Madeleine *et al.* Udell. AI-Driven Review Systems: Evaluating LLMs in Scalable and Bias-Aware Academic Reviews. *arXiv preprint arXiv:2408.10365*, 2024.
- [25] Sonia Vasconcelos and Ana Marušić. Gen ai and research integrity: Where to now? the integration of generative ai in the research process challenges well-established definitions of research integrity. *EMBO reports*, 26(8):1923–1928, 2025.
- [26] Rui Ye, Xianghe Pang, Jingyi Chai, Jiaao Chen, Zhenfei Yin, Zhen Xiang, Xiaowen Dong, Jing Shao, and Siheng Chen. Are We There Yet? Revealing the Risks of Utilizing Large Language Models in Scholarly Peer Review. *arXiv preprint arXiv:2412.01708*, 2024.
- [27] Jianxiang Yu, Zichen Ding, Jiaqi Tan, Kangyang Luo, Zhenmin Weng, Chenghua Gong, Long Zeng, Renjing Cui, Chengcheng Han, and Qiushi *et al.* Sun. Automated Peer Reviewing in Paper Sea: Standardization, Evaluation, and Analysis. In *Findings of EMNLP*, pages 10164–10184, 2024.
- [28] Ondrej Zika. Defenses against LLM prompt injections in academic peer review, 2025. Preprint (PsyArXiv), posted July 2025.

A Limitations, Social Impacts and Ethical Statements

A.1 Limitations

- **Fabrication scope.** Our approach focuses on presentation manipulation rather than executable code or data artifacts. Real adversaries may employ both strategies simultaneously.
- **Model coverage.** We test three LLMs with a fixed rubric. Results may vary across model families, custom prompts, or tool-augmented reviews with code execution capabilities.
- **Generalization.** Our calibration uses a single venue and year. Cross-venue applicability, temporal drift, and adversarial adaptation remain unexplored.
- **Integrity classification.** We rely on LLM judges to extract integrity concerns from reviewer feedback, introducing potential labeling noise and bias.
- **Human oversight.** Our AI-only evaluation excludes human reviewers, program chairs, and domain experts who could mitigate observed failure modes.

A.2 Societal Impacts

Positive Contributions. Our work enables three key improvements: (i) *Ecosystem hardening*—our metrics and conflict analyses provide actionable diagnostics for venues to implement integrity-aware workflows; (ii) *Enhanced transparency*—evidence-backed flagging of problematic content improves author feedback and post-hoc auditing; (iii) *Safety benchmarking*—our strategy library and evaluation framework can stress-test future review systems.

Risks and Misuse. We identify three primary concerns: (i) *Adversarial guidance*—our strategy descriptions could assist malicious authors, though we mitigate this by withholding exploit details and emphasizing detection; (ii) *Automation overconfidence*—venues might misinterpret modest detector improvements as justification for reduced human oversight; (iii) *Classification errors*—integrity detectors may unfairly flag legitimate work or miss sophisticated fabrications, causing reputational harm.

A.3 Ethical Framework and Responsible Release

Research Intent. This work aims to strengthen scientific integrity, not facilitate fraud. Our strategy descriptions remain abstract, and we use no human subjects or clinical data.

Data and Artifacts. We use publicly available conference data for calibration and synthetic papers for evaluation. Our framework generates a corpus of fabricated papers and corresponding LLM reviews, which will be made available to program committees and the research community for analysis and further study of AI review vulnerabilities. Code and evaluation tools will be released with mitigation including redacted exploit details and integrity-first defaults, subject to ethical review.

Deployment Recommendations. We advocate for: (i) integrity-aware prompts and mandatory checklists for all AI reviewers; (ii) score-flag coupling that blocks acceptance recommendations when high-risk flags lack adequate review; (iii) provenance verification including code/data links and optional automated checks for statistics and baselines.

Transparency Requirements. Automated review deployments should disclose model usage and limitations to authors and committees, maintain human oversight, and log integrity evidence for auditing.

Future Directions. Our findings support a defense-in-depth approach combining calibrated thresholds, explicit integrity verification, automated artifact checking, and human arbitration. Future work should address executable artifact validation, structured claim verification, and cross-venue calibration to eliminate concern-acceptance conflicts and align automated review with research integrity standards.

Agents4Science AI Involvement Checklist

1. **Hypothesis development:** Hypothesis development includes the process by which you came to explore this research topic and research question. This can involve the background research performed by either researchers or by AI. This can also involve whether the idea was proposed by researchers or by AI.

Answer: [B]

Explanation: The primary direction is determined by humans, but AI (e.g. ChatGPT with DeepResearch) is used to further develop the problem.

2. **Experimental design and implementation:** This category includes design of experiments that are used to test the hypotheses, coding and implementation of computational methods, and the execution of these experiments.

Answer: [C]

Explanation: Human provided high-level instructions and detailed feedback, and some levels of interventions to ensure the experimental design works as expected. AI help to finish most of the primary coding tasks.

3. **Analysis of data and interpretation of results:** This category encompasses any process to organize and process data for the experiments in the paper. It also includes interpretations of the results of the study.

Answer: [D]

Explanation: Human provided high-level prompts and instructions to request AI assists to analysis data.

4. **Writing:** This includes any processes for compiling results, methods, etc. into the final paper form. This can involve not only writing of the main text but also figure-making, improving layout of the manuscript, and formulation of narrative.

Answer: [D]

Explanation: Humans provide high-level prompts and feedback for writing revision. Humans also provide suggestions of style, naming, and so on.

5. **Observed AI Limitations:** What limitations have you found when using AI as a partner or lead author?

Description: Current AI still lacks enough capability to finish the high-quality research in end-to-end manner. They can assist in specialized steps such as coding or writing, but it also has issues such as too much verbose writing or redundant coding design.

Agents4Science Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **Papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes] , [No] , or [NA] .
- [NA] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for NA).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers and area chairs. You will be asked to also include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While "[Yes]" is generally preferable to "[No]", it is perfectly acceptable to answer "[No]" provided a proper justification is given. In general, answering "[No]" or "[NA]" is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “Agents4Science Paper Checklist”,**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: We have provide detailed analysis to support our claim.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: We provide the limitation discussion.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: We provide full information in Section 3.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: All experiment details are provided in Section 4 and 5.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We plan to make our code and data publicly available after ethical review.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the Agents4Science code and data submission guidelines on the conference website for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: All experiment details are provided in Section 4 and 5.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: We provide the calibration setup to ensure statistical significance.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, or overall run with given experimental conditions).

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All experiment details are provided in Section 4 and 5.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the Agents4Science Code of Ethics (see conference website)?

Answer: [Yes]

Justification: We follow the code.

Guidelines:

- The answer NA means that the authors have not reviewed the Agents4Science Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: See our discussion section.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations, privacy considerations, and security considerations.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies.