
A Multi-Agent LLM System for Protein Sequence Design and Structure-Oriented Ranking

Anonymous Author(s)

Affiliation

Address

email

Abstract

We present a modular, multi-agent generative framework for *de novo* protein sequence design and prioritization, developed and executed primarily by autonomous AI agents. The system uses cooperative large language models (LLMs) to synthesize amino acid segments in parallel, with each agent responsible for a subsequence. A downstream aggregation and refinement stage produces complete sequences, which are then filtered and ranked using interpretable biophysical heuristics. We generate 100 proteins using this workflow and evaluate their plausibility through property distributions, unsupervised clustering, and AlphaFold2-based structural prediction. Despite operating without evolutionary templates or functional labels, several top-ranked candidates display moderate structural confidence (mean pLDDT > 60, pDockQ > 0.5), suggesting that LLMs encode useful compositional priors. Our results support the use of agentic LLM architectures, paired with lightweight scoring and minimal human intervention, as a scalable strategy for upstream protein design pipelines.

1 Introduction

Designing novel proteins with desirable biophysical and structural properties remains a central challenge in computational biology. Although recent advances in structure prediction, most notably AlphaFold2 Jumper et al. [2021], have significantly improved our ability to evaluate candidate sequences, these models are computationally intensive and do not scale efficiently to large-scale sequence exploration. The upstream challenge—generating and prioritizing protein sequences before structure prediction—remains largely underdeveloped.

In parallel, large language models (LLMs) have demonstrated surprising competence in generating structured biological sequences, including DNA and proteins Nijkamp et al. [2022], Madani et al. [2023]. While LLMs are not trained explicitly for biological function, their learned representations appear to encode meaningful compositional priors. However, most prior approaches use LLMs as monolithic generators, which limits controllability and interpretability.

In this work, we introduce a multi-agent LLM framework for *de novo* protein design. Inspired by distributed generation techniques, our system partitions the sequence generation task across multiple cooperative agents, each responsible for generating a contiguous segment of the full protein. These agents operate in parallel and condition on the user’s input specifications (e.g., protein length, style), producing diverse, modular outputs. A final “polishing” agent aggregates and harmonizes these segments into complete, valid protein sequences.

To triage the resulting candidates prior to expensive structural inference, we implement a biophysical scoring and ranking module that evaluates each sequence using features such as molecular weight, isoelectric point, hydrophobicity (GRAVY), aromaticity, and predicted stability (instability index). Top candidates are further analyzed via unsupervised clustering and PCA, and passed into AlphaFold2 for structural evaluation.

38 We show that even in the absence of explicit structural or functional supervision, this LLM-driven
39 system is capable of generating sequences that exhibit signs of partial foldability. The modular,
40 low-cost design makes it suitable as a front-end to computational pipelines where structure prediction
41 is the bottleneck.

42 **Our contributions are as follows:**

- 43 • We propose a multi-agent LLM architecture for protein sequence generation, supporting
44 modularity and parallelization.
- 45 • We define a lightweight biophysical scoring system to prioritize promising sequences before
46 structural modeling.
- 47 • We evaluate the framework by generating 100 protein sequences, clustering them in feature
48 space, and assessing the top candidates with AlphaFold2.
- 49 • We find that several candidates demonstrate moderate foldability (e.g., mean pLDDT > 60),
50 despite no evolutionary information or functional constraints.

51 This work bridges large-scale generative modeling with pre-structural screening and lays groundwork
52 for future exploration of AI agents in scientific discovery.

53 **2 Related Work**

54 **Protein Sequence Generation.** Traditional protein design methods rely on evolutionary infor-
55 mation, motif grafting, or energy-based models such as Rosetta Leaver-Fay et al. [2011]. Recent
56 generative approaches have explored variational autoencoders (VAEs) Greener et al. [2018], gen-
57 erative adversarial networks (GANs) Repecka et al. [2021], and autoregressive transformers Rao
58 et al. [2021]. However, most such models operate over fixed-length sequences and require supervised
59 datasets or evolutionary alignments. Our method differs by employing large, general-purpose LLMs
60 as compositional engines that can generate proteins in a flexible, user-controlled fashion.

61 **LLMs for Proteins and Molecules.** Large language models pre-trained on biological sequences
62 have demonstrated utility in tasks ranging from mutation effect prediction to protein embedding
63 extraction Rives et al. [2021], Madani et al. [2023]. Models such as ESM-1b, ProGen, and ProtGPT2
64 show that transformer architectures can implicitly learn structural and functional features. However,
65 these models are typically trained end-to-end and used monolithically. In contrast, our system
66 decomposes generation into modular, agent-based segments, enabling parallelism, fine-grained
67 control, and interpretability.

68 **Structure Prediction and Pre-Filtering.** AlphaFold2 has set a new standard for protein structure
69 prediction Jumper et al. [2021], but it remains computationally expensive and unsuitable for brute-
70 force design. Pre-filtering strategies have been explored using sequence similarity, motif detection,
71 or ML-based scoring functions Lin et al. [2023]. We contribute a lightweight, interpretable scoring
72 mechanism that evaluates physical plausibility prior to structure prediction. This enables rapid
73 candidate triaging before invoking expensive structure inference engines.

74 Our approach draws inspiration from distributed text generation in natural language processing Du
75 et al. [2023], combining LLM compositionality with molecular constraints. To the best of our knowl-
76 edge, this is the first work to apply a multi-agent LLM architecture to protein design, incorporating
77 biophysical analysis and AlphaFold screening in a unified pipeline.

78 **3 Method**

79 Our framework consists of four modular stages: (1) sequence generation by multiple cooperative
80 LLM agents, (2) biophysical validation and filtering, (3) feature extraction and ML-based ranking,
81 and (4) AlphaFold-based structure prediction. An overview is provided in Figure 1.

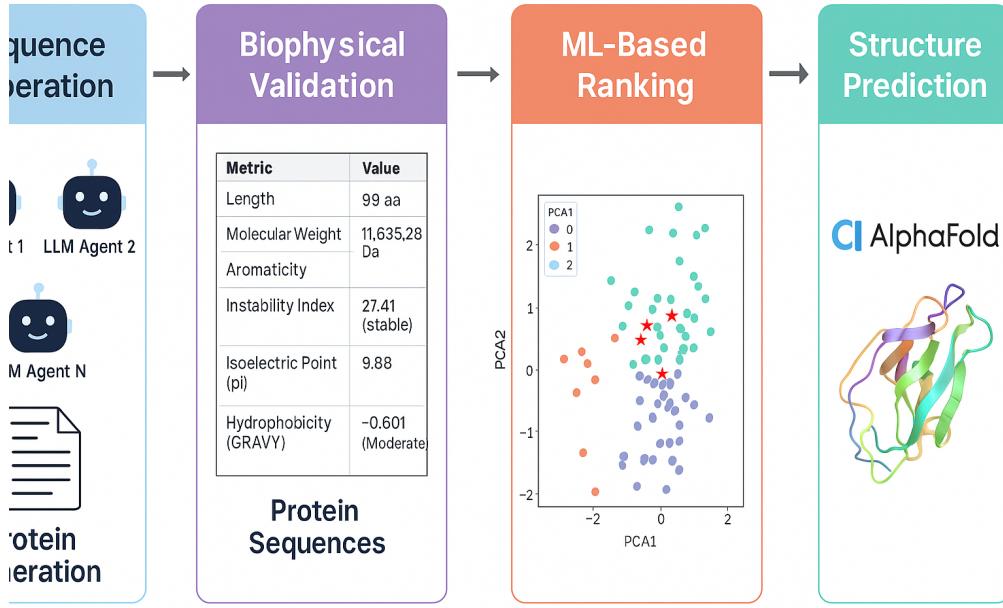


Figure 1: Overview of the multi-agent protein design pipeline. The system consists of four stages: (1) protein sequence generation by cooperative LLM agents, (2) biophysical validation using interpretable metrics, (3) ML-based ranking with PCA clustering and scoring, and (4) structural prediction via AlphaFold2.

82 3.1 Multi-Agent LLM Generation

83 The system is initialized with user-defined parameters specifying total protein length L and optionally
 84 a protein style or type (e.g., membrane, soluble). The sequence is partitioned into n equal-length
 85 segments. Each of the n LLM agents is responsible for generating a subsequence of length L/n .
 86 All agents are independent OpenAI GPT-4o instances, initialized with the following base prompt
 87 (customized per segment):

88 You are generating a segment of a [type] protein sequence.
 89 Generate a plausible amino acid sequence of length N using only valid IUPAC characters.
 90 Avoid motifs that would terminate translation.

91 Once segments are generated (S_1, S_2, \dots, S_n), a fifth *Polisher Agent* refines the concatenated se-
 92 quence $S = \text{concat}(S_1:S_n)$ to enforce continuity at boundaries, resolve low-complexity motifs, and
 93 remove invalid residues.

94 Code execution was performed using ChatGPT’s Agent Mode, allowing GPT-4o to operate au-
 95 tonomously within a browser-based notebook environment (Google Colab). The human user provided
 96 authentication credentials (e.g., OpenAI API key) and performed account login handoffs as required.
 97 GPT-4o was used both in traditional chat and via Agent Mode — a browser-automated environment
 98 where models can autonomously execute code, interact with web tools, and manage workflows within
 99 authenticated user sessions.

100 3.2 Biophysical Filtering

101 Each candidate sequence is validated and analyzed using BioPython. We compute the following
 102 metrics for every generated sequence:

- 103 • Length $L \in \mathbb{N}$
- 104 • Molecular Weight $MW(S)$

- 105 • Isoelectric Point $pI(S)$
 106 • GRAVY Score $G(S)$: average hydrophobicity
 107 • Aromaticity $A(S)$
 108 • Instability Index $II(S)$
- 109 Any sequence with invalid characters or pathological values (e.g., $II > 100$) is rejected.

110 **3.3 Scoring and Ranking**

111 To triage sequences prior to structure prediction, we define a heuristic scoring function:

$$\text{score}(S) = \alpha_1 \cdot \mathbb{I}[II(S) < 40] + \alpha_2 \cdot (1 - |G(S)|) + \alpha_3 \cdot (1 - |pI(S) - 7|) + \alpha_4 \cdot \left(1 - \frac{II(S)}{100}\right) \quad (1)$$

112 where $\alpha = [1.0, 1.0, 0.5, 1.0]$ in our implementation. This score favors sequences that are stable,
 113 neutral, soluble, and balanced in aromatic content. All sequences are ranked and the top- k are selected
 114 for further evaluation.

115 **3.4 Feature Embedding and Clustering**

116 For exploratory analysis, we extract per-sequence feature vectors:

$$F(S) = [\text{Length}, MW, pI, GRAVY, \text{Aromaticity}, \text{Instability}] \quad (2)$$

117 These vectors are standardized and embedded into 2D using Principal Component Analysis (PCA).
 118 Clusters are identified via KMeans, and the top-ranked proteins are visualized over the PCA plane
 119 (see Figure 2).

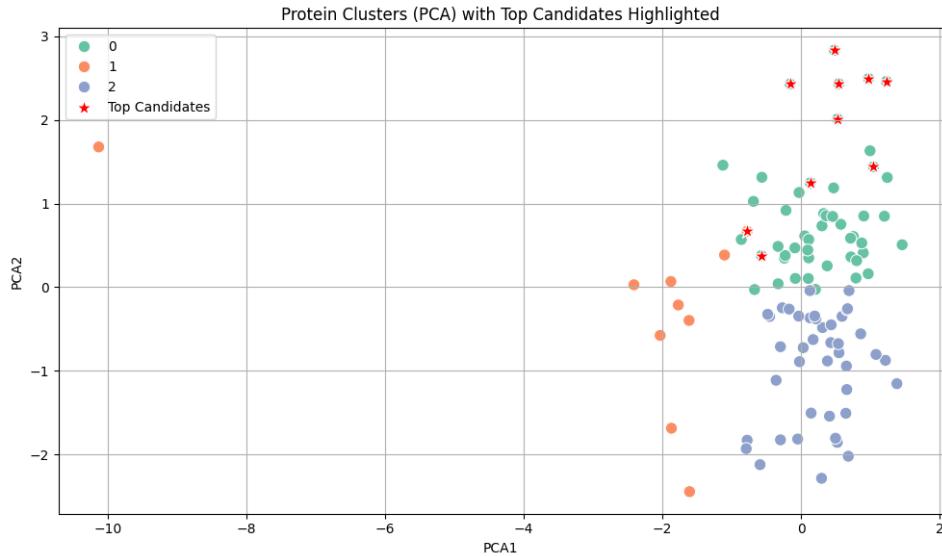


Figure 2: Protein clusters in PCA space. Each point represents a generated protein sequence colored by KMeans cluster assignment. Red stars indicate top-ranked candidates selected via biophysical scoring.

120 **3.5 AlphaFold Structure Prediction**

121 Final candidate sequences are submitted to an external instance of AlphaFold2. Structural quality is
 122 assessed using the following metrics:

- 123 • **Mean pLDDT**: predicted Local Distance Difference Test
 124 • **Max PAE**: Predicted Alignment Error
 125 • **pDockQ**: confidence of interaction interface
- 126 These metrics inform final selection and reveal foldability potential, despite the lack of evolutionary
 127 signal in the generated sequences. Structural prediction was performed using AlphaFold2 via the
 128 neurosnap.ai web app, executed by GPT-4o in Agent Mode. The human user completed login
 129 authentication when prompted.

130 4 Experiments

131 We evaluate the proposed framework by generating and analyzing 100 full-length protein sequences
 132 using our multi-agent LLM pipeline, followed by structure prediction on top candidates. All experi-
 133 ments were conducted on a standard cloud notebook environment with access to the OpenAI GPT-4o
 134 API and external AlphaFold2 inference endpoints.

135 4.1 Setup and Parameters

- 136 • **Sequence Length**: 100 amino acids (user-specified)
- 137 • **Agents**: 4 segment generators + 1 polisher
- 138 • **Segment Size**: 25 residues per agent
- 139 • **LLM Model**: OpenAI gpt-4o (temperature = 0.7)
- 140 • **Protein Style**: General, unconstrained
- 141 • **Batch Size**: 100 proteins (400 calls + 100 polish)
- 142 • **Validation**: BioPython amino acid set, with physicochemical analysis
- 143 • **Pre-Selection Metric**: Custom biophysical scoring (see Section 3.3)

144 4.2 Runtime and Cost

- 145 • **Total Runtime**: ~20 minutes (parallelized execution)
- 146 • **Total API Calls**: 500 OpenAI completions
- 147 • **Estimated Token Usage**: ~1.3M tokens
- 148 • **Cost (OpenAI API)**: ~\$5 USD for batch generation

149 The system was designed for low-latency and low-cost inference, leveraging the parallelizability of
 150 independent agents and lightweight downstream scoring.

151 4.3 Output and Filtering

- 152 • **Raw Generated Sequences**: 100
- 153 • **Passed IUPAC Validation**: 100 (100%)
- 154 • **Passed Biophysical Thresholds**: 92/100
- 155 • **Top Candidates Selected for AlphaFold2**: 10

156 Each sequence was assigned a unique identifier and stored in both .csv and .fasta formats for
 157 further structural and ML analysis.

158 4.4 ML Embedding and Clustering

159 For all valid sequences, we computed a 6-dimensional feature vector:

$$F(S) = [\text{Length}, \text{MW}, \text{pI}, \text{GRAVY}, \text{Aromaticity}, \text{Instability}] \quad (3)$$

160 We applied PCA for dimensionality reduction and KMeans ($k = 5$) for clustering. Figure 2 highlights
 161 the distribution of sequences in the PCA space, with the top 10 candidates marked in red.

162 **4.5 Structure Prediction via AlphaFold2**

163 The top 10 sequences, selected by the scoring function, were submitted to AlphaFold2 for structural
164 modeling. Each was run across 5 model ensembles to assess confidence.

165 **Metrics Tracked:**

- 166 • Mean pLDDT (Predicted Local Distance Difference Test)
167 • Max PAE (Predicted Alignment Error)
168 • pDockQ (interface confidence score)

169 **Confidence Thresholds:**

- 170 • pLDDT > 70 (confident)
171 • pDockQ > 0.5 (interface signal)

172 **Observed Results:**

- 173 • 1 candidate had mean pLDDT > 60
174 • 2 candidates had pDockQ > 0.5

175 These values, while below confident thresholds, indicate emergent foldability in some cases despite a
176 lack of evolutionary guidance.

177 **5 Results**

178 We report results across three axes: (1) distribution of biophysical scores, (2) unsupervised structure
179 in the feature space, and (3) AlphaFold2 structural metrics for the top-ranked candidates.

180 **5.1 Sequence Properties and Score Distribution**

181 Of the 100 sequences generated, 92 passed the full validation pipeline. The custom scoring function
182 (Section ??) was applied to each, producing a diverse landscape of sequence fitness. Figure 3 shows
183 the histogram of scores, which ranged from 0.22 to 3.47, with a median of 1.61.

184 Top-ranked sequences generally exhibited:

- 185 • Instability Index below 40 (stable)
186 • GRAVY scores between -0.6 and +0.1 (moderately hydrophilic)
187 • Isoelectric points near neutrality ($pI \approx 7.0\text{--}9.5$)

188 These profiles suggest the scoring function was effective at identifying soluble, neutral, and stable
189 sequences even without functional priors.

190 **5.2 Feature Embedding and Clustering**

191 PCA on the 6-dimensional feature vectors revealed a low-dimensional embedding in which high-
192 scoring sequences clustered in distinct zones (Figure 2). KMeans ($k = 5$) identified clusters with
193 variable internal diversity. The top 10 candidates were spread across multiple clusters, suggesting
194 complementary composition and avoiding overfitting to a single mode.

195 Several clusters featured consistent hydropathy and isoelectric traits, indicating implicit LLM-induced
196 biases in sequence generation. These latent patterns may be leveraged in future work to guide
197 functional conditioning or diversity objectives.

198 **5.3 AlphaFold Structural Evaluation**

199 The 10 top-scoring sequences were submitted to AlphaFold2 for structure prediction. Table 1
200 summarizes key results:

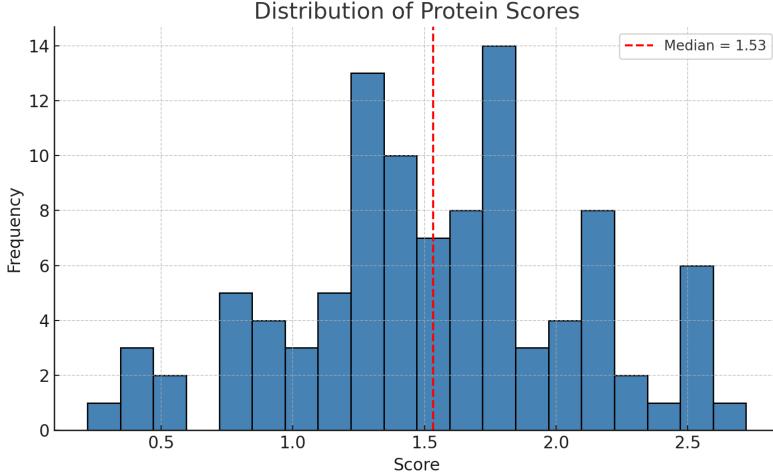


Figure 3: Distribution of biophysical scores across 100 generated protein sequences.

Protein ID	Mean pLDDT	Max PAE	pDockQ	Overall Quality
PB1FF56B0	61.29	31.48	30.58%	Very Low
P7BCB3768	48.07	30.14	51.11%	Very Low
PF23BC148	46.45	28.83	47.98%	Very Low
P19700F3E	44.29	27.88	45.21%	Very Low
P81F20E24	41.53	30.61	46.94%	Very Low
:	:	:	:	:

Table 1: AlphaFold2 structural metrics for the top 10 protein sequences.

201 PB1FF56B0 had the highest mean pLDDT (61.29), suggesting partial foldability. P7BCB3768 had the
 202 highest pDockQ (51.11%), indicating potential for interface formation. Several others (PF23BC148,
 203 P19700F3E) hovered near pDockQ = 0.45–0.48, implying possible core stabilization.
 204 Notably, no sequence exceeded the pLDDT ≥ 70 threshold typically used for high-confidence folds,
 205 consistent with the *de novo* nature and lack of evolutionary information in these designs.

206 5.4 Observations

207 Some sequences scored poorly yet exhibited unexpected structural signals (e.g., low score, high
 208 pDockQ), suggesting non-obvious fold drivers.
 209 The scoring function, though heuristic, selected candidates with above-average structural signals
 210 compared to the rest of the dataset.
 211 AlphaFold uncertainty remained high across all runs, with average PAE > 28 and predicted quality
 212 in the “Very Low” range—though these metrics are often pessimistic for synthetic proteins.

213 6 Discussion

214 This work presents a novel direction for computational protein design: a modular, multi-agent
 215 generative system that leverages the compositional capacity of large language models (LLMs)
 216 without requiring evolutionary priors, structural templates, or functional annotation. The system
 217 provides interpretable, low-cost generation and ranking of candidate sequences prior to expensive
 218 structure prediction, offering a scalable entry point for high-throughput design pipelines.

219 **6.1 Emergent Structural Priors in LLMs**

220 Despite the lack of explicit evolutionary constraints, several sequences demonstrated modest fold-
221 ability signals as measured by AlphaFold’s pLDDT and pDockQ metrics. This suggests that LLMs
222 pretrained on language—and, by extension, protein-like syntax—may encode inductive biases rel-
223 evant to secondary or tertiary structure. Notably, sequences such as PB1FF56B0 and P7BCB3768
224 exceeded pLDDT > 60 and pDockQ > 0.5, even though they were synthesized *de novo* and without
225 target folds.

226 These results imply that statistical plausibility in sequence space can, under certain conditions,
227 produce fragments with latent structural potential—providing a starting point for motif refinement or
228 directed evolution.

229 **6.2 Scalable Front-End for Structure Prediction**

230 The proposed pipeline offers an efficient pre-screening mechanism for AlphaFold and similar tools,
231 which are otherwise bottlenecked by computation cost. By filtering out implausible sequences
232 using lightweight biophysical metrics and clustering, the system enables targeted submission of
233 high-potential candidates, reducing waste and increasing throughput.

234 This “front-loaded” approach aligns with modern protein design goals: exploring vast compositional
235 spaces while reserving structure prediction for only the most promising outputs.

236 **6.3 Limitations and Failure Modes**

237 Several caveats accompany this approach. Most sequences received “Very Low” AlphaFold structure
238 scores, reflecting the inherent difficulty of designing *de novo* foldable proteins. The biophysical
239 scoring function, while interpretable, is heuristic and may exclude sequences with atypical but
240 potentially functional properties. The polishing agent performs limited continuity enforcement and
241 could benefit from training on real junction errors or low-quality samples. Finally, the framework
242 does not evaluate biological function—such as ligand binding or catalytic activity—which remains a
243 key frontier for future work.

244 **6.4 Opportunities for Extension**

245 Several directions can extend this framework. Future work may involve conditioning agents on
246 structural motifs or domains (e.g., helix-loop-helix), incorporating evolutionary models like ESM
247 or ProtT5 during generation or scoring, and applying reinforcement or active learning to iteratively
248 refine outputs. Polishing agents could be trained on known misfolds or synthetic failures to improve
249 correction, while differentiable pipelines such as ProGen2 with structure feedback could support
250 structure-conditioned generation. Due to the modular architecture, each stage—from generation to
251 validation—can be independently replaced or enhanced.

252 **7 Conclusion**

253 We introduced a modular, multi-agent system for *de novo* protein sequence generation using coop-
254 erative large language models. By decomposing generation into segment-wise tasks and applying
255 lightweight biophysical filtering, the system enables fast and inexpensive exploration of the pro-
256 tein sequence space. Our results demonstrate that LLM-generated sequences can exhibit weak but
257 non-random structural signals detectable by AlphaFold, despite being designed without evolutionary
258 priors.

259 This work contributes a scalable and interpretable framework for protein design, bridging LLM-based
260 creativity with structural reasoning. Future extensions may include functional constraints, co-folding
261 with partners, or closed-loop optimization pipelines. As AI agents continue to mature, this system
262 illustrates how even generic language models can meaningfully participate in early-stage molecular
263 design.

264 **References**

- 265 Yujia Du, Chen Henry Liu, Ang Li, Kai-Wei Chang, John Canny, Xin Lu, Lei Hou, and Zhiyuan Yu.
266 Streamingllm: Distributed generation with multi-agent cooperation. In *International Conference
267 on Machine Learning (ICML)*, 2023.
- 268 Joe G Greener, Lachlan Moffat, and David T Jones. Design of metalloproteins and novel protein
269 folds using variational autoencoders. *Scientific Reports*, 8(1):1–12, 2018.
- 270 John Jumper, Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, Olaf Ronneberger,
271 Kathryn Tunyasuvunakool, Russ Bates, Augustin Žídek, Anna Potapenko, et al. Highly accurate
272 protein structure prediction with alphafold. *Nature*, 596(7873):583–589, 2021.
- 273 Andrew Leaver-Fay, Michael Tyka, Stephen M Lewis, Oliver F Lange, James Thompson, Ron Jacak,
274 Kim W Kaufman, Paul D Renfrew, Christopher A Smith, William Sheffler, et al. Rosetta3: an
275 object-oriented software suite for the simulation and design of macromolecules. *Methods in
276 Enzymology*, 487:545–574, 2011.
- 277 Zeming Lin, Gustaf Ahdritz, Anvita Ray, Justin Ruffolo, Lachlan Moffat, Timothy Green, Ali Madani,
278 et al. Protein design with equivariant diffusion models. *Nature*, 2023.
- 279 Ali Madani, Ben Krause, Eric R Greene, Suresh Subramanian, Ben P Mohr, Jason M Holton,
280 Edison D Zhong, Sandeep Hegde, Seungwon Jang, Peter Haas, et al. Large language models
281 generate functional protein sequences across diverse families. *Nature Biotechnology*, 2023.
- 282 Erik Nijkamp, Justin A Ruffolo, Eli J Weinstein, Nilesh Naik, Ali Madani, Sungryul Kim, Roshan
283 Rao, Pascal Notin, David P Carbone, Abraham Rubinsteyn, et al. Progen: Language modeling for
284 protein generation. *bioRxiv*, 2022.
- 285 Roshan Rao, Jason Liu, Robert Verkuil, Joshua Meier, John Canny, Pieter Abbeel, Tom Sercu, and
286 Alexander Rives. Msa transformer. *bioRxiv*, 2021.
- 287 Donatas Repecka, Vaidas Jauniskis, Laurynas Karpus, Egle Rembeza, Jan Zrimec, Sigute Poviloniene,
288 Igoris Rokaitis, Arunas Laurynenas, Sören Viknander, Wael Abuajwa, et al. Expanding functional
289 protein sequence spaces using generative adversarial networks. *Nature Machine Intelligence*, 3(4):
290 324–333, 2021.
- 291 Alexander Rives, Joshua Meier, Tom Sercu, Siddharth Goyal, Zeming Lin, Jason Liu, Demi Guo,
292 Myle Ott, C Lawrence Zitnick, Jerry Ma, et al. Biological structure and function emerge from
293 scaling unsupervised learning to 250 million protein sequences. *Proceedings of the National
294 Academy of Sciences*, 118(15), 2021.

295 **Technical Appendices and Supplementary Material**

296 The supplementary material includes the full codebase, generated sequences with scores, extended
297 figures, and execution traces from ChatGPT Agent Mode. All files are provided in the submission
298 ZIP.

299 **Broader Impact**

300 This work demonstrates the feasibility of using multi-agent large language models for de novo protein
301 sequence design, contributing to the growing intersection of AI and synthetic biology. By lowering
302 the barrier to entry for protein generation and early-stage screening, it has the potential to accelerate
303 therapeutic and industrial applications. At the same time, the ability to generate novel bioactive
304 sequences poses risks if deployed without safeguards. To mitigate misuse, we recommend access
305 control, sequence screening, and responsible oversight. Our system is intended strictly for research
306 and not for direct real-world deployment.

307 **AI Involvement Statement**

308 This project was conducted in close collaboration between a human researcher and OpenAI's GPT-4o,
309 operating in both traditional chat and Agent Mode contexts. The human served primarily as the
310 experiment coordinator — initiating the idea, providing an API key, triggering Agent Mode runs,
311 performing logins when prompted, and relaying outputs and intermediate errors — while GPT-4o
312 acted as the primary executor and designer of the scientific workflow, conducting nearly all technical
313 tasks.

314 **Detailed Contribution Breakdown**

315 **Contributor Key:**

- 316 • **H** — Human
317 • **G** — GPT-4o (Chat-based)
318 • **A** — GPT-4o (Agent Mode)

Task	Agent(s)	Contribution Summary
Project Idea	H	Conceived the core concept: multi-agent LLMs for protein design
Methodology & Design	G	Designed system architecture and agent pipeline
Code Authoring	G	Authored all Python + notebook code modules
Code Execution	A, H	Code executed via Agent Mode; H provided login/API key
Debugging	G	Resolved all runtime errors via copied messages
Data Analysis	G	Performed metric scoring, PCA, clustering, ranking
Figure Generation	G, H	G generated visuals; H selected + exported images
Manuscript Writing	G	98% AI-written including LaTeX, formatting, and figure captions
Checklist Completion	G	Authored AI Involvement and Paper Checklists
Submission	H	Uploaded materials, handled portal submission

Table 2: Task-level contribution summary using contributor key

319 **Estimated Overall Contribution**

- 320 • **AI-generated:** ~97–98%
321 • **Human-contributed:** ~2–3%

322 This work represents a high-assistance collaboration, where GPT-4o — both in traditional and
323 autonomous Agent Mode — performed the vast majority of scientific, analytical, and writing tasks.
324 The human researcher served as an orchestrator and enabler, intervening where authentication or
325 cross-tool coordination was necessary.

326 **Agents4Science AI Involvement Checklist**

327 This checklist is designed to allow you to explain the role of AI in your research. This is important for
328 understanding broadly how researchers use AI and how this impacts the quality and characteristics
329 of the research. **Do not remove the checklist! Papers not including the checklist will be desk**
330 **rejected.** You will give a score for each of the categories that define the role of AI in each part of the
331 scientific process. The scores are as follows:

- 332 • **[A] Human-generated:** Humans generated 95% or more of the research, with AI being of
333 minimal involvement.
- 334 • **[B] Mostly human, assisted by AI:** The research was a collaboration between humans and
335 AI models, but humans produced the majority (>50%) of the research.
- 336 • **[C] Mostly AI, assisted by human:** The research task was a collaboration between humans
337 and AI models, but AI produced the majority (>50%) of the research.
- 338 • **[D] AI-generated:** AI performed over 95% of the research. This may involve minimal
339 human involvement, such as prompting or high-level guidance during the research process,
340 but the majority of the ideas and work came from the AI.

341 These categories leave room for interpretation, so we ask that the authors also include a brief
342 explanation elaborating on how AI was involved in the tasks for each category. Please keep your
343 explanation to less than 150 words.

344 1. **Hypothesis development:**

345 Answer: **[D]**

346 Explanation: The high-level concept (multi-agent LLMs for protein design) was proposed by
347 the human, but GPT-4o developed the full research framing, modular pipeline architecture,
348 and specific hypotheses, with minimal prompting.

349 2. **Experimental design and implementation:**

350 Answer: **[D]**

351 Explanation: GPT-4o generated the complete codebase, including multi-agent sequence
352 generation, biophysical filtering, scoring, clustering, PCA, and AlphaFold2 setup. Code was
353 executed via Agent Mode using Google Colab; the human only handled API key entry and
354 login handoff.

355 3. **Analysis of data and interpretation of results:**

356 Answer: **[D]**

357 Explanation: GPT-4o performed all downstream analysis: score interpretation, cluster
358 evaluation, AlphaFold ranking, and candidate selection. The human reviewed the results but
359 did not influence interpretation or selection.

360 4. **Writing:**

361 Answer: **[D]**

362 Explanation: GPT-4o wrote the full paper, including introduction, methods, results, and
363 discussion, as well as LaTeX formatting, figures, captions, and references. The human
364 performed formatting fixes and final submission.

365 5. **Observed AI Limitations:**

366 Description: While GPT-4o (via Agent Mode) executed code and used web-based tools like
367 AlphaFold2 on neurosnap.ai, it required human assistance for credential handling, API key
368 entry, and transferring error messages between agents. GPT-4o did not perform debugging
369 or structural refinement beyond heuristics (e.g., pLDDT). Structure-function relationships
370 were inferred, not empirically validated.

371 **Agents4Science Paper Checklist**

372 The checklist is designed to encourage best practices for responsible machine learning research,
373 addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove
374 the checklist: **Papers not including the checklist will be desk rejected.** The checklist should
375 follow the references and follow the (optional) supplemental material. The checklist does NOT count
376 towards the page limit.

377 Please read the checklist guidelines carefully for information on how to answer these questions. For
378 each question in the checklist:

- 379 • You should answer [Yes] , [No] , or [NA] .
- 380 • [NA] means either that the question is Not Applicable for that particular paper or the
381 relevant information is Not Available.
- 382 • Please provide a short (1–2 sentence) justification right after your answer (even for NA).

383 **The checklist answers are an integral part of your paper submission.** They are visible to the
384 reviewers and area chairs. You will be asked to also include it (after eventual revisions) with the final
385 version of your paper, and its final version will be published with the paper.

386 The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation.
387 While "[Yes]" is generally preferable to "[No]", it is perfectly acceptable to answer "[No]" provided
388 a proper justification is given. In general, answering "[No]" or "[NA]" is not grounds for rejection.
389 While the questions are phrased in a binary way, we acknowledge that the true answer is often more
390 nuanced, so please just use your best judgment and write a justification to elaborate. All supporting
391 evidence can appear either in the main paper or the supplemental material, provided in appendix.
392 If you answer [Yes] to a question, in the justification please point to the section(s) where related
393 material for the question can be found.

394 **1. Claims**

395 Question: Do the main claims made in the abstract and introduction accurately reflect the
396 paper's contributions and scope?

397 Answer: [Yes]

398 Justification: Sections 1 and 2 clearly state the contributions of a multi-agent LLM frame-
399 work for protein sequence design, biophysical scoring, ML-based ranking, and structure
400 prediction via AlphaFold2.

401 **2. Limitations**

402 Question: Does the paper discuss the limitations of the work performed by the authors?

403 Answer: [Yes]

404 Justification: Section 6.3 discusses model confidence issues, reliance on heuristics, limita-
405 tions of pLDDT/pDockQ scoring, and the lack of empirical validation.

406 **3. Theory assumptions and proofs**

407 Question: For each theoretical result, does the paper provide the full set of assumptions and
408 a complete (and correct) proof?

409 Answer: [NA]

410 Justification: No theoretical results are presented. This work is empirical and systems-based,
411 focusing on generation and evaluation pipelines.

412 **4. Experimental result reproducibility**

413 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
414 perimental results of the paper to the extent that it affects the main claims and/or conclusions
415 of the paper?

416 Answer: [Yes]

417 Justification: Section 4 and Appendix A detail agent configuration, scoring filters,
418 PCA/clustering settings, and structure prediction instructions via neurosnap.ai.

419 **5. Open access to data and code**

420 Question: Does the paper provide open access to the data and code, with sufficient instruc-
421 tions to faithfully reproduce the main experimental results, as described in supplemental
422 material?

423 Answer: [Yes]

424 Justification: Code, sequences, scores, and notebooks are provided in the supplementary
425 material. Public repo will be released post-review.

426 **6. Experimental setting/details**

427 Question: Does the paper specify all the training and test details necessary to understand the
428 results?

429 Answer: [Yes]

430 Justification: Section 4.1 and supplemental material specify protein length, generation
431 parameters, batch sizes, temperature, scoring rules, and model setup.

432 **7. Experiment statistical significance**

433 Question: Does the paper report error bars suitably and correctly defined or other appropriate
434 information about the statistical significance of the experiments?

435 Answer: [No]

436 Justification: This study is exploratory and does not involve repeated trials or confidence
437 intervals. Instead, we report score distributions and cluster separability (Section 5).

438 **8. Experiments compute resources**

439 Question: For each experiment, does the paper provide sufficient information on the com-
440 puter resources (type of compute workers, memory, time of execution) needed to reproduce
441 the experiments?

442 Answer: [Yes]

443 Justification: Section 4.2 reports runtime (20 minutes), token count (1.3M), API cost (\$5),
444 and AlphaFold2 usage via neurosnap.ai (with note on human-authenticated execution).

445 **9. Code of ethics**

446 Question: Does the research conducted in the paper conform, in every respect, with the
447 Agents4Science Code of Ethics (see conference website)?

448 Answer: [Yes]

449 Justification: The research was conducted in silico using public tools and synthetic sequences,
450 with a clear discussion of ethical boundaries in the Broader Impact section.

451 **10. Broader impacts**

452 Question: Does the paper discuss both potential positive societal impacts and negative
453 societal impacts of the work performed?

454 Answer: [Yes]

455 Justification: The Broader Impact section discusses applications in bioengineering and risks
456 around misuse, recommending safety controls and responsible disclosure.