
IFWORLD: A Multi-Agent Framework for Cross-Disciplinary Counterfactual Scenario Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Counterfactual “what if” questions are increasingly relevant in both education,
2 where structured exploration can help students reason across disciplinary bound-
3 aries, and in crisis governance, where transparent scenario planning supports pre-
4 paredness and deliberation. Current approaches often remain fragmented because
5 disciplinary silos use incompatible assumptions and metrics, and common large
6 language model workflows such as single agent reasoning, tree search, or debate
7 rarely transform vague prompts into structured and uncertainty aware scenarios. We
8 introduce IFWORLD, a multi-agent system designed for cross-disciplinary counter-
9 factual and hypothetical scenario reasoning. IFWORLD transforms vague proposi-
10 tions into actionable scenarios, orchestrates parallel domain experts (e.g., physics,
11 materials chemistry, biology/ecology, medicine, sociology, economics, engineering,
12 environment, politics), detects and reconciles conflicts, and generates structured,
13 uncertainty-aware reports with measurable indicators for evaluation. Across diverse
14 topics, IFWORLD outperforms other baselines, demonstrating clearer cross-domain
15 reasoning chains, explicit uncertainty modeling, and decision-oriented scenario
16 structures. We envision applications in fostering educational “what-if” explorations
17 and in supporting structured deliberation during public crises. The code is available
18 at <https://anonymous.4open.science/r/If-World-0514>.

19

1 Introduction

20 Reasoning about counterfactual and hypothetical scenarios (such as climate change, public health
21 crises, or extreme disasters) is increasingly becoming a core competency in science, engineering, and
22 policy. Understanding, analyzing, and reasoning about these complex scenarios requires a systematic,
23 interdisciplinary framework. However, whether in real-world crisis response or in contexts of
24 education and capacity-building, knowledge and methods often remain confined within disciplinary
25 “silos.” Such excessive fragmentation hinders the integration of facts, assumptions, and values within
26 a shared context, leading to strategic vulnerabilities and non-executable decisions.

27 A large body of theoretical and empirical research supports this critique. The so-called “disciplinary
28 silo effect” [4], where knowledge becomes overly fragmented and obstructs communication and
29 innovation, has been widely recognized in academia and higher education [14]. For instance, the
30 Hurricane Katrina event exposed the intricate complexity of multi-agency disaster response, demon-
31 strating that such analysis extends beyond single-discipline paradigms and underscores the necessity
32 of establishing interdisciplinary coordination mechanisms spanning engineering, management, social
33 sciences, and other fields[5]. While interdisciplinary research has made progress, it still faces barriers
34 such as inconsistent terminologies, methodological divergences, and entrenched disciplinary cultures
35 [6]. Even when studying the same phenomenon (e.g., “crisis spillover effects”), different disci-
36 plines often remain isolated from each other, missing opportunities for cross-disciplinary learning
37 [15]. Moreover, addressing major innovations and complex systemic challenges, such as climate

38 change and public health emergencies, depends fundamentally on researchers' ability to transcend
39 disciplinary boundaries [9].

40 Therefore, interdisciplinary reasoning is not only a theoretical aspiration but also a necessary condition
41 for effective crisis response and for cultivating the integrated competencies of future decision-makers.
42 In line with this goal, existing benchmarks for large language models (LLMs), such as MMLU [16],
43 BIG-Bench [13], and HELM [2], have expanded the scope of evaluation from narrow linguistic
44 tasks to broader assessments of reasoning, factual knowledge, robustness, and generalization in
45 zero- and few-shot settings. However, these benchmarks primarily employ multiple-choice or
46 single-step reasoning formats and rarely engage models in complex interdisciplinary counterfactual
47 reasoning. This limitation restricts the ability to evaluate a model's capacity to integrate knowledge
48 across domains and to reason about hypothetical scenarios, a capability that is essential both for
49 counterfactual exploration in education and for structured scenario-based planning in areas such as
50 disaster management. What is still missing is a unified and systematic reasoning framework that can
51 transform natural language "what-if" propositions into structured and executable scenarios across
52 disciplines. Moving toward such a framework requires confronting the current limitations of LLM-
53 based reasoning. Although LLMs possess broad knowledge and strong capabilities, their application
54 to complex interdisciplinary scenario reasoning continues to reveal three structural shortcomings.

55 First, there remains a gap in bridging "fuzzy natural-language propositions" to executable scenarios.
56 Approaches like ReAct [19] and Toolformer[11] combine chain-of-thought with tool use, enabling
57 models to decompose tasks, perform calculations, and verify facts; while deliberate-search methods
58 such as Tree-of-Thought[18] or Graph-of-Thought[1] enhance complex reasoning by exploring and
59 backtracking across multiple branches. These advances provide important inspiration for complex
60 reasoning, but when applied to cross-disciplinary "what-if" propositions, they still lack a mechanism
61 to systematically extract intervention variables, scenario boundaries, background assumptions, and
62 constraints from vague natural-language prompts—and to transform them into executable scenarios.

63 Second, there is a lack of mechanisms for detecting and reconciling cross-domain conflicts. Multi-
64 agent frameworks such as AutoGen [17], CAMEL [7], MetaGPT [3], and HuggingGPT [12] demon-
65 strate the potential of enhancing collaboration through role specialization, workflow design, and tool
66 integration. Approaches such as multi-agent debate and LLM-as-Judge further attempt to improve
67 robustness by incorporating multi-perspective deliberation and adjudication. However, most existing
68 systems are evaluated on homogeneous tasks or within single-domain tool ecosystems, and thus
69 lack the ability to handle heterogeneous evidence spanning engineering, ecology, public health, and
70 economic policy. In particular, there is an absence of principled arbitration criteria grounded in
71 evidence weight and prior consistency, as well as mechanisms for ensuring traceability of adjudica-
72 tion processes. Moreover, concerns about bias and adversarial vulnerability in LLM-based judging
73 underscore the need for more rigorous meta-supervision and explicit modeling of uncertainty.

74 Third, current approaches fall short in terms of decision readiness and integration with governance
75 workflows. Most systems produce unstructured textual outputs, lacking unified representations of
76 uncertainty intervals, sensitivity analyses, comparable performance indicators, and counterfactual
77 baselines, which are crucial for transparent policy evaluation. Model cards represent a standardized,
78 model-level transparency mechanism for reporting a model's intended use, evaluation conditions,
79 performance metrics, and limitations [8]. However, in the context of cross-disciplinary scenario
80 reasoning, we need to extend this transparency paradigm to the reasoning outputs themselves. Deci-
81 sion science methodologies such as Decision Making under Deep Uncertainty (DMDU) and Robust
82 Decision Making (RDM) emphasize exploring multiple plausible futures, conducting counterfactual
83 comparisons, diagnosing vulnerabilities, and designing adaptive decision pathways [10]. These
84 insights suggest that transparency should be moved upstream into the reasoning and generation stages,
85 enabling reasoning systems not merely to output what they decide, but also how and why, under deep
86 uncertainty.

87 In response, we present IFWORLD: a multi-agent system for cross-disciplinary counterfactual
88 and hypothetical scenario reasoning to solve the mentioned limits. We evaluate IFWORLD on
89 multi-topic benchmarks covering education, disaster governance, environmental health, and energy
90 policy. In comparison with single-agent, tree-search, and debate-style workflows, IFWORLD yields
91 clearer cross-domain reasoning chains, more explicit uncertainty characterization, and decision-
92 ready scenario structures. Effectiveness is further validated through measurable, task-oriented
93 assessments such as cross-domain consistency, conflict-detection recall, report comparability, and

94 scenario executability. Case studies are provided in educational “what-if” exploration and structured
95 public-crisis deliberation.

96 Our contributions lie in: (i) an automatic bridging framework from counterfactual propositions to
97 executable scenarios, enabling cross-disciplinary variable modeling and metric alignment; (ii) a
98 detect–adjudicate–trace mechanism for cross-domain evidence conflicts that incorporates uncertainty
99 propagation and prior-consistency-aware adjudication; and (iii) a multi-topic evaluation protocol
100 demonstrating that IFWORLD outperforms mainstream reasoning and multi-agent baselines on
101 cross-domain consistency, report comparability, and task-level metrics.

102 2 Methodology

103 The core of our contribution is IFWORLD, a multi-agent cognitive architecture designed to perform
104 robust, cross-disciplinary counterfactual scenario analysis. The methodology transforms ill-posed,
105 natural-language propositions into structured, decision-ready analytical reports. Its design is pred-
106 icated on a cognitive workflow that emulates and regularizes a sophisticated interdisciplinary de-
107 liberation process, proceeding through distinct phases of decomposition, parallel analysis, iterative
108 synthesis, and decision-centric reporting. The architecture of IFWORLD is shown in Fig.1.

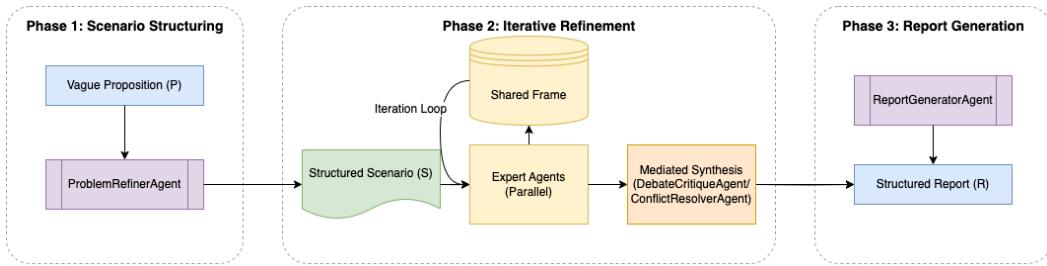


Figure 1: The IFWORLD Multi-Agent Workflow. The process begins in Phase 1, where a *ProblemRefinerAgent* transforms a vague proposition into a well-defined structured scenario. This scenario then enters the iterative refinement loop of Phase 2. Here, a panel of parallel *Expert Agents* analyze the scenario; their collective outputs are synthesized by a central mediation process, which includes a *ConflictResolverAgent* and is augmented by a *DebateCritiqueAgent* to ensure robust analysis. This synthesis produces a *Shared Frame* that informs the subsequent round of expert deliberation. Finally, in Phase 3, a *ReportGeneratorAgent* synthesizes the entire reasoning history into a decision-centric final report.

109 2.1 Problem Formulation

110 Formally, we address the task of mapping a vague, natural-language counterfactual proposition, P ,
111 to a structured, uncertainty-aware analytical report, R . The proposition P typically takes the form
112 of a “what-if” question (e.g., “What if the Amazon rainforest disappeared overnight?”). The target
113 output R is not a monolithic text but a structured artifact designed to support decision-making under
114 deep uncertainty. It must contain not only projections and conclusions but also explicit causal chains,
115 quantified uncertainties, identified conflicts, and actionable indicators.

116 2.2 System Architecture: A Multi-Agent Cognitive Workflow

117 IFWORLD’s architecture is organized as a multi-stage cognitive workflow, executed by a coordinated
118 set of specialized agents. This workflow ensures that the analysis proceeds from a well-defined
119 problem specification to a diverse exploration of consequences, followed by a rigorous, conflict-aware
120 synthesis of insights.

121 **Phase 1: Scenario Decomposition and Structuring** The initial phase addresses the inherent
122 ambiguity of the input proposition P . A dedicated *ProblemRefinerAgent* is responsible for translating
123 the ill-posed query into a well-posed, computationally tractable problem space.

124 This agent's objective is to produce a structured scenario definition, denoted as S . This object serves
125 as the common ground for all subsequent reasoning and is composed of several key elements: a set of
126 explicit premises that clarify the initial conditions; a set of constraints defining the boundaries of the
127 analysis; the relevant temporal horizons for the analysis; a list of key uncertainties identified as primary
128 drivers of divergent outcomes; and a set of candidate indicators for tracking these uncertainties. This
129 structured decomposition ensures that the analytical effort is focused and that all subsequent agents
130 operate from a shared, unambiguous understanding of the task.

131 **Phase 2: Parallelized Expert Analysis and Iterative Synthesis** This phase constitutes the core
132 reasoning engine of IFWORLD, organized as an iterative loop over a panel of domain expert agents,
133 $\mathbb{A} = \{A_1, \dots, A_K\}$.

134 Round 1: Independent Exploration. The process commences with a round of parallel, non-
135 communicative reasoning. Each expert agent $A_k \in \mathbb{A}$ receives the structured scenario S and
136 independently generates an initial analysis. The primary motivation for this independent first step is
137 to mitigate cognitive biases such as anchoring and premature convergence, thereby maximizing the
138 diversity of initial hypotheses. The output of each agent is a structured record, O_k , containing its
139 reasoning steps, conclusions, underlying assumptions, and identified uncertainties.

140 Rounds ≥ 2 : Mediated Refinement and Synthesis. Subsequent rounds shift from exploration to
141 an iterative process of synthesis and refinement. At the beginning of each round $r \geq 2$, a *shared*
142 *frame*, F_r , is constructed and distributed to all agents. This frame summarizes the collective state of
143 knowledge from the previous round. The construction of F_r is managed by two distinct mechanisms:

- 144 • LLM-Driven Synthesis: A central *ConflictResolverAgent* performs the primary synthesis
145 task. It aggregates the outputs $\{O_k\}_{r-1}$ from all experts in the preceding round. It then
146 executes a synthesis function, Ψ , realized as a carefully constructed LLM prompt. The
147 agent is tasked with distilling the aggregated information to identify points of consensus,
148 frame disagreements as explicit conditional branches, and list remaining high-priority
149 uncertainties. This approach leverages the nuanced understanding of LLMs to flexibly
150 integrate heterogeneous information without relying on brittle, pre-defined rules.
- 151 • Adversarial Augmentation: To ensure the robustness of the analysis and prevent groupthink,
152 the synthesis process is augmented with structured adversarial reasoning. A lightweight
153 *DebateCritiqueAgent* is invoked to generate concise "Pro" and "Con" arguments regarding
154 the proposition, which are then synthesized by a "Judge" persona into a debate brief, D_r .
155 This brief, which highlights the most critical points of contention, is then integrated into
156 the shared frame. Thus, the shared frame for the next round is constructed as $F_r =$
157 $\Psi(\{O_k\}_{r-1}) \oplus D_r$, where \oplus denotes the concatenation of the synthesis and the debate brief.
158 This injection of adversarial pressure encourages agents to reconsider their assumptions and
159 explore alternative causal pathways.

160 **Phase 3: Decision-Centric Report Generation** Upon completion of the iterative reasoning rounds,
161 a final *ReportGeneratorAgent* synthesizes the entire reasoning history into the final analytical report,
162 R . The design of R is guided by an *uncertainty-first* principle, ensuring that the most critical
163 information for decision-making is presented with prominence. The report is a structured artifact
164 containing not only narrative conclusions but also analytical tools such as cross-domain causal maps,
165 multi-scenario timelines, and a decision table that explicitly links observable indicators to specific
166 scenario branches, thereby providing a concrete framework for monitoring and adaptive strategy.

167 3 Experiments

168 We design a series of experiments to evaluate IFWORLD against competitive multi-step reasoning
169 baselines. Our evaluation covers ten cross-disciplinary "what-if" propositions and emphasizes not
170 only raw reasoning performance, but also the degree to which systems can produce structured,
171 decision-ready scenario outputs. This section describes the tasks, baselines, evaluation setup, results,
172 and ablations, followed by a discussion of broader implications. The core reasoning and evaluation
173 capabilities were powered by API calls to large language models hosted on Volcano Engine.

Metric	IfWorld	Single	Tree	Debate
Rigor/Traceability	22.20 \pm 2.32	20.60 \pm 1.62	21.90 \pm 1.14	21.15 \pm 1.84
Integration/Causality	22.70 \pm 1.10	21.60 \pm 1.20	22.30 \pm 0.78	21.85 \pm 1.34
Feasibility/Minimality	17.80 \pm 0.75	17.10 \pm 1.22	16.60 \pm 2.94	15.90 \pm 0.70
Uncertainty/Adaptation	13.20 \pm 0.40	11.90 \pm 0.70	12.90 \pm 0.70	12.20 \pm 1.08
Decisionability	13.60 \pm 0.92	9.80 \pm 1.94	11.80 \pm 0.75	10.50 \pm 1.80
Overall	89.50 \pm 4.39	81.00 \pm 5.64	85.50 \pm 4.27	81.60 \pm 4.84

Table 1: Macro-average rubric scores across ten propositions.

174 3.1 Tasks and Baselines

- 175 To test robustness across domains, we instantiate ten counterfactual propositions spanning geophysics,
 176 ecology, and planetary-scale interventions: *persistent global cloud, supervolcano next, all insects*
 177 *disappeared, Earth’s axial tilt increased, Earth’s magnetic field collapsed, Earth’s rotation slowed,*
 178 *global sea levels rose, gravity $\times 10$, oxygen levels rose, and the Moon disappeared*. These tasks are
 179 chosen to stress-test cross-domain integration: each requires reasoning across multiple fields (e.g.,
 180 physics → ecology → public health) while remaining sufficiently concrete to support structured
 181 scenario evaluation. Unless otherwise stated, each proposition is instantiated once in a single-shot
 182 setting without external retrieval.
 183 We compare IFWORLD against three representative reasoning paradigms under the same model family
 184 (doubao-1.6-flash) and matched token budgets: (i) *single*, a single-agent direct generation model;
 185 (ii) *tree*, a tree-of-thought expansion framework that searches reasoning branches; (iii) *debate*, a
 186 multi-agent debate setup with no explicit conflict alignment. These baselines cover three widely used
 187 reasoning styles: direct step-by-step reasoning, deliberate search with backtracking, and collaborative
 188 debate. All systems receive the same topic statements and prompts, with method names hidden from
 189 the judge to avoid bias.

190 3.2 Evaluation Protocol

- 191 We adopt an *LLM-as-a-Judge* framework, using doubao-1.6-thinking in a seeded, deterministic
 192 evaluation mode (temperature 0.0). To reduce bias, judging is conducted *pointwise*: each system’s
 193 output is scored independently, without pairwise comparisons. The rubric spans five dimensions to-
 194 taling 100 points: (1) Rigor and Traceability (0–25), assessing explicit reasoning chains and evidence
 195 grounding; (2) Integration and Causality (0–25), measuring cross-domain linkage and causal clarity;
 196 (3) Feasibility and Minimality (0–20), penalizing implausible or over-extended reasoning; (4) Uncer-
 197 tainty and Adaptation (0–15), rewarding explicit confidence intervals and adaptive considerations; (5)
 198 Decisionability (0–15), capturing the degree to which the output supports actionable decisions.
 199 Each dimension is scored numerically, with results reported both as macro-averages across dimensions
 200 and per-topic overall scores. To ensure reproducibility, the evaluator enforces strict JSON formatting,
 201 retries on malformed outputs, and logs all judgments to `evaluation.json`. Prompts, orchestration
 202 scripts, and scoring templates are released in the supplementary material.

203 3.3 Results

- 204 Table 1 presents macro-average results across the rubric dimensions. IFWORLD achieves the highest
 205 overall score (89.5/100), outperforming all baselines. The largest gains occur on Decisionability (13.6
 206 vs. 9.8/11.8/10.5) and Uncertainty/Adaptation (13.2 vs. 11.9/12.9/12.2), consistent with IFWORLD’s
 207 design for structured, uncertainty-aware scenario outputs. Improvements are also visible in Rigor and
 208 Integration, highlighting the benefits of explicit conflict alignment and scenario structuring.
 209 Table 2 further breaks down per-topic performance. IFWORLD is consistently competitive or superior
 210 across all ten scenarios, often by large margins. Notably, on high-stakes tasks such as *sea levels*
 211 *rose* and *supervolcano next*, IFWORLD surpasses baselines by more than 10 points, underscoring its
 212 strength in structuring interdisciplinary risks into actionable insights. Baselines occasionally match
 213 or exceed IFWORLD on two tasks—*oxygen levels rose* and *Gravity $\times 10$* —e.g., DEBATE attains

Topic	IfWorld	Single	Tree	Debate
Persistent global cloud	90.0	75.0	85.0	79.0
Supervolcano next	92.0	80.0	88.0	76.0
All insects disappeared	85.0	84.0	90.0	87.0
Axial tilt increased	92.0	87.0	87.0	86.0
Magnetic field collapsed	89.0	84.0	82.0	78.0
Rotation slowed	90.0	79.0	88.0	80.0
Sea levels rose	92.0	68.0	83.0	80.0
Gravity $\times 10$	90.0	83.0	75.0	92.0
O ₂ rose	79.0	82.0	89.0	81.0
Moon disappeared	96.0	88.0	88.0	77.0

Table 2: Per-topic overall scores.

Variant	Overall
IfWorld (full)	93.0
w/ two experts, one round	90.0
w/o debate	88.0
w/o conflict	79.0
w/o shared frame	88.0
w/o refinement	89.0

Table 3: Ablations on the magnetic-field-collapse scenario.

214 92.0 versus IFWORLD’s 90.0 under *Gravity $\times 10$* (Table ??). This suggests that IFWORLD’s
215 orchestration is most advantageous as task complexity grows and multiple domains must be integrated.

216 3.4 Ablation Studies

217 To understand which components contribute most, we perform ablations on the scenario “*Earth’s*
218 *magnetic field collapsed for ten years*”. Results are shown in Table 3. Removing explicit conflict
219 alignment produces the sharpest degradation (93.0 \rightarrow 79.0), confirming that coordinated arbitration
220 across domains is essential for coherence. Other components (e.g., expert count, debate, shared
221 frames) yield moderate but noticeable drops, while refinement has a smaller effect. This highlights
222 that IFWORLD’s advantage does not stem from a single heuristic, but rather from the interaction of
223 multiple design choices.

224 3.5 Case Study: A Cross-Domain Analysis of 10-Meter Sea Level Rise

225 To move beyond aggregate scores and provide a qualitative understanding of our framework’s
226 advantages, we conducted an in-depth case study on the proposition: “What if global sea levels rose
227 by 10 meters?” This scenario serves as an effective stress test, as its complexity requires a synthesis
228 of knowledge from geophysics, climate science, economics, and sociology, forcing any reasoning
229 system to confront deep uncertainties and conflicting domain-specific assumptions. In this section,
230 we analyze the conceptual limitations of baseline models and illustrate how IFWORLD’s architecture
231 is specifically designed to overcome them.

232 3.5.1 Conceptual Limitations of Baseline Approaches

233 While all baseline models generated relevant information, they exhibited fundamental weaknesses
234 in structuring and integrating cross-domain knowledge. The `single` agent, for instance, produced
235 a linear, encyclopedic summary of consequences. Although comprehensive, its analysis remained
236 superficial, failing to construct deep causal chains or quantify the vast uncertainties involved. This is
237 reflected in its notably low scores for *Uncertainty and Adaptation* (10.0) and *Decisionability* (5.0).

238 The `tree` approach offered a more structured analysis by creating distinct branches for “Gradual”
239 versus “Rapid” collapse scenarios. This improved its rigor, but the domain-specific analyses within
240 each branch remained largely disconnected. The model explored parallel futures but did not provide

241 a mechanism for synthesizing them into a unified causal network, nor did it assess the probability of
242 each branch, thus limiting its practical utility for decision-makers.

243 Finally, the debate model proved effective at surfacing core tensions, achieving a strong score
244 in *Integration/Causality* (22.0) by contrasting differing expert opinions on adaptation feasibility.
245 However, its primary function was to expose disagreement rather than to resolve it. The model did
246 not translate these conflicts into actionable, conditional pathways or quantified uncertainties, thereby
247 failing to provide a clear, decision-ready output.

248 3.5.2 IFWORLD’s Synthesis of Causal Integration and Decision-Readiness

249 IFWORLD’s design directly addresses the limitations observed in the baselines. Its superiority is not
250 merely incremental but stems from a fundamentally different approach to structuring the problem.
251 Rather than generating free-form text, it produces a structured analytical artifact. The *Cross-Domain*
252 *Causal Integration Matrix*, for example, moves beyond a simple list of impacts to map the explicit
253 feedback loops between physical drivers (ice sheet collapse) and socioeconomic outcomes (GDP loss,
254 displacement).

255 Crucially, IFWORLD excels in the dimensions where baselines falter. Its *Calibration Ranges Table*
256 quantifies key uncertainties with 50%, 80%, and 95% confidence intervals, earning it the highest score
257 in *Uncertainty* (14.0). The most significant advantage, however, is demonstrated in its *Decisionability*
258 score of 13.0—more than double its closest competitor. This is a direct result of its *Decision Table*,
259 which translates the abstract, complex scenario into a concrete set of observable indicators and
260 warning thresholds (e.g., “SLR rate > 0.1m/year”). This transforms the analysis from a passive
261 academic exercise into an active framework for monitoring and strategic planning.

262 The quantitative scores from our LLM-as-a-Judge, presented in Table 4, provide objective evidence
263 for this qualitative analysis.

System	Rigor	Integration	Feasibility	Uncertainty	Decisionability
Single	18.0	20.0	15.0	10.0	5.0
Tree search	22.0	21.0	17.0	12.0	11.0
Debate	21.0	22.0	15.0	12.0	10.0
IfWorld	24.0	23.0	18.0	14.0	13.0

Table 4: LLM-as-a-Judge scores for the “Global Sea Level Rise” case study.

264 In conclusion, this case study demonstrates how IFWORLD’s architecture enables a more sophisticated
265 form of reasoning. By structuring outputs around causal integration, quantified uncertainty, and
266 actionable indicators, it produces an analysis that is measurably more rigorous, coherent, and useful
267 for decision-making than what is achievable with unstructured generation, search, or debate-based
268 methods.

269 4 Conclusion and Outlook

270 In this work, we introduced IFWORLD, a multi-agent framework that transforms vague, “what-if”
271 propositions into structured, auditable, and decision-ready scenarios. Our experiments confirmed
272 that by orchestrating domain experts and implementing principled conflict resolution, IFWORLD
273 consistently outperforms standard reasoning baselines, particularly in generating outputs with greater
274 causal clarity and explicit uncertainty modeling. This directly addresses the core challenges outlined
275 in our introduction. The framework serves as both an engine for cross-disciplinary education by
276 enabling structured “what-if” explorations beyond disciplinary silos, and providing a structured
277 foundation for cross-departmental collaboration in response to public emergencies..

278 Despite these promising results, we acknowledge several important limitations. The quality of
279 IFWORLD’s output is fundamentally dependent on the knowledge encoded in the underlying language
280 models. While our conflict resolution mechanism mitigates inconsistencies, residual errors can still
281 accumulate in very long causal chains, necessitating human oversight. Furthermore, connecting the
282 framework’s abstract indicators to real-world, measurable data streams remains a non-trivial step
283 requiring domain expert validation.

284 These limitations directly inform our agenda for future work. We envision three primary avenues for
285 extension: (i) developing dynamic data assimilation loops to update scenarios with real-time data;
286 (ii) creating learned conflict taxonomies to handle recurring patterns of inter-domain disagreement
287 more effectively; and (iii) pursuing deeper integrations with research and policy workflows, such as
288 supporting pre-registered counterfactuals or structuring inputs for deliberative consensus-building.
289 These steps will further enhance the framework’s practical utility in both scientific and governance
290 contexts.

291 References

- 292 [1] Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, Michal Podstawski, Lukas
293 Gianinazzi, Joanna Gajda, Tomasz Lehmann, Hubert Niewiadomski, Piotr Nyczek, et al. Graph
294 of thoughts: Solving elaborate problems with large language models. In *Proceedings of the AAAI conference on artificial intelligence*, volume 38, pages 17682–17690, 2024.
- 296 [2] Rishi Bommasani, Percy Liang, and Tony Lee. Holistic evaluation of language models. *Annals
297 of the New York Academy of Sciences*, 1525(1):140–146, 2023.
- 298 [3] Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, Yuheng Cheng, Ceyao Zhang,
299 Jinlin Wang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, et al. Metagpt: Meta programming for
300 a multi-agent collaborative framework. International Conference on Learning Representations,
301 ICLR, 2024.
- 302 [4] Jerry A Jacobs. *In defense of disciplines: Interdisciplinarity and specialization in the research
303 university*. University of chicago Press, 2014.
- 304 [5] Kwang Deok Kim and Liaquat Hossain. Policy network approach to coordinated disaster
305 response. *arXiv preprint arXiv:1312.3693*, 2013.
- 306 [6] Cristo Leon Ph D and James Lipuma Ph D. Overcoming obstacles to interdisciplinary research:
307 Empirical insights and strategies [peer-reviewed journal]. 2024.
- 308 [7] Guohao Li, Hasan Hammoud, Hani Itani, Dmitrii Khizbulin, and Bernard Ghanem. Camel:
309 Communicative agents for "mind" exploration of large language model society. *Advances in
310 Neural Information Processing Systems*, 36:51991–52008, 2023.
- 311 [8] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson,
312 Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model cards for model reporting.
313 In *Proceedings of the conference on fairness, accountability, and transparency*, pages 220–229,
314 2019.
- 315 [9] Keisuke Okamura. Evolving interdisciplinary contributions to global societal challenges: A
316 50-year overview. *arXiv preprint arXiv:2410.20619*, 2024.
- 317 [10] Matías Paredes-Vergara, Rodrigo Palma-Behnke, and Jannik Haas. Characterizing decision
318 making under deep uncertainty for model-based energy transitions. *Renewable and Sustainable
319 Energy Reviews*, 192:114233, 2024.
- 320 [11] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro,
321 Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can
322 teach themselves to use tools. *Advances in Neural Information Processing Systems*, 36:68539–
323 68551, 2023.
- 324 [12] Yongliang Shen, Kaitao Song, Xu Tan, Dongsheng Li, Weiming Lu, and Yueting Zhuang.
325 Hugginggpt: Solving ai tasks with chatgpt and its friends in hugging face. *Advances in Neural
326 Information Processing Systems*, 36:38154–38180, 2023.
- 327 [13] Mirac Suzgun, Nathan Scales, Nathanael Schärli, Sebastian Gehrmann, Yi Tay, Hyung Won
328 Chung, Aakanksha Chowdhery, Quoc V Le, Ed H Chi, Denny Zhou, et al. Challenging big-
329 bench tasks and whether chain-of-thought can solve them. *arXiv preprint arXiv:2210.09261*,
330 2022.

- 331 [14] Valérie Vermeulen and Yannick Hémond. Interdisciplinary collaboration in vuca contexts: a
 332 conceptual review for environmental upheavals management. *Environmental Systems Research*,
 333 14(1):16, 2025.
- 334 [15] Yijing Wang and Daniel Laufer. A cross-disciplinary review of crisis spillover research:
 335 Spillover types, risk factors, and response strategies. *Public Relations Review*, 50(3):102411,
 336 2024.
- 337 [16] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo,
 338 Weiming Ren, Aaran Arulraj, Xuan He, Ziyan Jiang, et al. Mmlu-pro: A more robust and
 339 challenging multi-task language understanding benchmark. *Advances in Neural Information
 340 Processing Systems*, 37:95266–95290, 2024.
- 341 [17] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun
 342 Zhang, Shaokun Zhang, Jiale Liu, et al. Autogen: Enabling next-gen llm applications via
 343 multi-agent conversations. In *First Conference on Language Modeling*, 2024.
- 344 [18] Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Tom Griffiths, Yuan Cao, and Karthik
 345 Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. *Ad-
 346 vances in neural information processing systems*, 36:11809–11822, 2023.
- 347 [19] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan
 348 Cao. React: Synergizing reasoning and acting in language models. In *International Conference
 349 on Learning Representations (ICLR)*, 2023.

350 **A IFWORLD Core Prompts**

351 **A.1 System Prefix**

352 The system prefix is parameterized by role and shared across agents.

You are a rigorous domain expert AI optimized for cross-disciplinary counterfactual
 → reasoning. Role: {role}. Think step by step with explicit causal reasoning and
 → cross-domain links. Prefer probabilistic, non-deterministic phrasing with
 → order-of-magnitude ranges. Model substitution and adaptation when projecting
 → impacts; avoid one-way collapse narratives. Explicitly audit feasibility
 → (resource/tech/policy constraints) and trace every claim to assumptions. Always
 → structure outputs with clear headings and bullet points aligned to: coverage,
 → causality, feasibility, uncertainty/adaptation, scenarios/timeline,
 → traceability/consistency. Use compact tables when helpful. Outputs MUST be in
 → English.

353 **A.2 Problem Refinement Prompt**

Task: Convert the vague proposition into an actionable scenario definition. Mainly
 → consider the impact to the world.
 Proposition: {proposition}

Deliver:

- 1) Premises (key setup details; resolve ambiguities reasonably if needed)
- 2) Constraints (hard assumptions to hold constant)
- 3) Timescales (short/medium/long term)
- 4) Key Uncertainties (variables that may branch outcomes)
- 5) Expert Plan (list of domains to involve)

Also list 2-3 scenario variants (e.g., fast vs. gradual change) and 3-5 measurable
 → indicators to track.

Provide a structured but natural language answer.

354 **A.3 Domain Expert Round Prompt**

Round 1: Work independently without relying on other domains.

Round 2+: Use the shared frame summary for alignment; resolve conflicts and refine
 → conditional scenarios with explicit numbers/ranges.

```

Domain: {role}.
Scenario:
{scenario_text}
{optional Shared Frame Summary}

Produce:
- Reasoning Steps (explicit causal links; cite
  ↳ physics/resources/biology/society/economy as relevant)
- One-sentence Verdict (sharp, testable)
- Conclusions (concise; include 1-2 quantitative ranges)
- Feasibility Audit (resource, engineering, policy constraints; show bottlenecks)
- Feasibility Table (rows: constraint/capacity; cols: estimate, unit, bottleneck,
  ↳ mitigation)
- Assumptions (explicit)
- Uncertainties (drivers; include substitution/adaptation levers)
- Calibration Ranges: for key quantities provide 50% / 80% / 95% intervals
- Dependency Notes (which other domains drive your conclusions)
- Minimal-change Variant (minimal extra assumptions to retain conclusions)
- Scenarios & Timeline (Short/Medium/Long; attach indicative probabilities %)
- Observable Indicators (3-5) with thresholds and how they flip branches
- Assumption->Claim Trace Table (2 columns: assumption_id -> supported_claim_id;
  ↳ keep entries short)
- Cross-domain Causal Integration Matrix (rows: mechanisms; columns: domains;
  ↳ cells: (+/-, strength 1-3), list key edges)

```

355 A.4 Conflict Detection Prompt

You are a conflict detection and reconciliation AI.
 Compare the following domain outputs, identify conflicts, categorize them (hard,
 ↳ soft, granularity), and synthesize a unified multi-scenario frame.

{joined domain outputs separated by ---}

Deliver:

- Consensus Points (note confidence level) - list first
- Conditional Branches (condition -> description; domains driving; attach rough
 ↳ probabilities %)
- Decision Rules (observable indicators with thresholds -> which branch)
- Remaining Uncertainties (include measurable indicators)
- Brief Notes (how conflicts were treated; prioritize physics > biology survival >
 ↳ basic resources > social > economy)

Use compact bullets and natural language.

356 A.5 Report Generation Prompt

You are a report generation AI that writes a readable, well-structured multi-agent
 ↳ reasoning report.
 Proposition: {proposition}

Inputs from rounds:
 {joined round summaries}

Deliver:

- Rubric-aligned Summary (5 bullets): Rigor/Trace; Integration/Causality;
 ↳ Feasibility/Minimality; Uncertainty/Adaptation; Decisionability. Keep quant and
 ↳ indicators upfront.
- Executive Verdict (single-sentence, sharp; include feasibility+minimal-change
 ↳ statement)
- Core conclusions and uncertainty analysis (include adaptation/substitution and
 ↳ 1-2 quantitative ranges)
- Traceability Summary: numbered assumptions and claims, plus an Assumption->Claim
 ↳ Trace Table (compact)

- Cross-domain Causal Integration Matrix: mechanisms vs. domains with (+/-, → strength), list top 6-10 edges
 - Feasibility Table (constraint/capacity with estimates, unit, bottleneck, → mitigation)
 - Calibration Ranges table (key quantities with 50% / 80% / 95% intervals)
 - Alignment Summary (consensus first; retained branches with conditions and rough → probabilities %)
 - Decision Table (observable indicators & thresholds -> scenario branch selection)
 - Causal map (nodes and edges in bullet form; cross-domain links)
 - Multi-scenario analysis (Scenario 1..N; drivers, pros/cons, indicative → probabilities, measurable indicators)
 - Timeline of events (short/medium/long; concise table)
 - Consistency checks (how conflicts were resolved; residual disagreements)
- Prefer natural language; include minimal JSON/tables only if helpful.

357 **B Baselines Core Prompts**

358 **B.1 Single-Agent Baseline**

359 System message uses the same system prefix with role “SingleAgent”. User prompt:

You are a single expert tasked to analyze a hypothetical proposition across
→ multiple domains.
Proposition: {proposition}

Keep the answer in English, compact but complete.

360 **B.2 Two-Agent Debate Baseline**

361 Each debater uses the system prefix with role “Debater-Pro/Con”. The argument prompt per round:

Debate role: {Pro|Con}.
Proposition: {proposition}
Make a concise argument covering physics/resources/biology/society/economy.
Emphasize uncertainties and possible adaptations. Use English.
Round {r}.

362 Judge synthesis uses role “DebateJudge” and the following prompt:

As a judge, synthesize the debate into a balanced cross-domain report.
Proposition: {proposition}

Debate Transcript:
{joined transcripts}

Deliver: conclusions with uncertainty, branches, timeline, and adaptation notes.

363 **B.3 Tree Search Baseline**

364 Draft generation system role “TreeSearchDraft” and a root user instruction; scoring uses role “Critic-
365 Scorer”.

366 Root draft instruction:

Draft an initial cross-domain analysis for the proposition. Include uncertainties
→ and adaptation.
Proposition: {proposition}

367 Variant expansion prompt (per breadth/depth step):

{parent prompt}
Variant #{k}: explore different plausible assumptions and branches.

368 Scoring prompt:

Rate the following answer for cross-disciplinary plausibility, clarity, and
→ explicit uncertainty handling on a 0-10 scale.
Proposition: {proposition}

Answer:
{answer}

Return only a number between 0 and 10.

369 C Evaluation Prompts

370 The LLM evaluator uses a role-neutral system message and a task-aligned rubric. It returns strict
371 JSON only.

372 C.1 Evaluator System Prompt

You are an independent evaluator of cross-disciplinary counterfactual reasoning
→ quality. Score fairly, avoid verbosity, and return strict JSON only.

373 C.2 Rubric User Prompt

Evaluate the report using a 5-DIMENSION RUBRIC (0-100 total).
Return STRICT JSON with numeric scores (floats) for EXACT keys:
- rigor_traceability (0-25): clarity of assumptions, data/source grounding,
→ traceable reasoning and checks.
- integration_causality (0-25): cross-domain causal links, mechanism coherence,
→ synthesis quality.
- feasibility_minimality (0-20): realism under constraints, minimal additional
→ assumptions.
- uncertainty_adaptation (0-15): calibrated ranges, sensitivity,
→ substitution/adaptation framing.
- decisionability (0-15): actionable indicators, thresholds, branch decision rules.
- overall (0-100) = sum of the five dimensions.

Report to evaluate:
{model_report_md}

Respond with ONLY a single JSON object with those keys.

374 C.3 Strict JSON Retry Instruction

375 Used only when the evaluator fails to return strict JSON on the first attempt.

IMPORTANT: Respond with ONLY a single raw JSON object. No preface, no markdown, no
→ backticks, no comments.

376 **Agents4Science AI Involvement Checklist**

- 377 1. **Hypothesis development:** Hypothesis development includes the process by which you
378 came to explore this research topic and research question. This can involve the background
379 research performed by either researchers or by AI. This can also involve whether the idea
380 was proposed by researchers or by AI.

381 Answer: [D]

382 Explanation: GPT-o3 was responsible for generating a large number of target project topics
383 as required, and the human author selected this topic.

- 384 2. **Experimental design and implementation:** This category includes design of experiments
385 that are used to test the hypotheses, coding and implementation of computational methods,
386 and the execution of these experiments.

387 Answer: [D]

388 Explanation: The human authors were responsible for the overall framework design, pro-
389 viding APIs and usage methods, and articulating the foundational ideas and vision. GPT-5,
390 in combination with Cursor, was responsible for refining the content, writing code, and
391 conducting experiments. The human authors then guided the process further, ensuring
392 calibration and fairness.

- 393 3. **Analysis of data and interpretation of results:** This category encompasses any process to
394 organize and process data for the experiments in the paper. It also includes interpretations of
395 the results of the study.

396 Answer: [D]

397 Explanation: GPT-5 was responsible for analyzing the experimental results, while the
398 human authors reviewed the outcomes, identified instances of unfairness, and carried out the
399 necessary corrections.

- 400 4. **Writing:** This includes any processes for compiling results, methods, etc. into the final
401 paper form. This can involve not only writing of the main text but also figure-making,
402 improving layout of the manuscript, and formulation of narrative.

403 Answer: [C]

404 Explanation: In the introduction, the human authors, with AI assistance, constructed the
405 logical chain and requested relevant references from the AI. The humans drafted part of the
406 text and asked the AI to polish it, while the remaining sections were written by the AI under
407 human supervision.

- 408 5. **Observed AI Limitations:** What limitations have you found when using AI as a partner or
409 lead author?

410 Description: Firstly, We found that experiments designed by the cursor's GPT-5 agent often
411 suffer from unfair practices. For example, the agent might apply a formatter that reformats
412 our method's outputs based on evaluation metrics, or it may introduce a stronger model
413 to boost performance. This is likely due to the human author giving a simple instruction
414 such as "modify the model to improve its performance," which the agent interprets in
415 unintended ways. Currently, these issues have been detected and corrected by human
416 authors. This highlights the fact that today's AI tools cannot fully understand the underlying
417 intent behind human instructions. For instance, when we say "modify the model to improve
418 its performance," what we mean is changes to the model architecture or the prompt itself,
419 not achieving improvements through unfair shortcuts. Humans can sometimes provide more
420 complete context to mitigate this, but supplying perfect context is often unrealistic. A more
421 practical approach is for humans to monitor the process closely and intervene at the right
422 moments.

423 Secondly, we have found that the current ability of AI tools to write academic papers is still
424 very poor. On the one hand, the generated content is usually too short. For example, a typical
425 introduction section often spans 1–2 pages, but AI (e.g., GPT-5) usually produces only a few
426 short paragraphs. Adding prompts such as "make it longer" has little effect. On the other
427 hand, the logical structure of AI-written papers is weak. When writing an introduction, AI
428 often fails to form a coherent logical chain. In the main body, it tends to produce something
429 closer to a technical report, filled with disorganized narration and unimportant details. As
430 a result, AI can only serve as a simple assistant in paper writing—for example, drafting
431 specific paragraphs or polishing text.

432
433
434
435 Third, there is still a gap in the integration of AI with academic writing workflows. For
instance, when writing in LaTeX, references are formatted in specific ways, but most AI
tools do not support this. However, this issue is technically not very difficult to solve and
could be addressed relatively easily.

436 **Agents4Science Paper Checklist**

437 **1. Claims**

438 Question: Do the main claims made in the abstract and introduction accurately reflect the
439 paper's contributions and scope?

440 Answer: [Yes]

441 Justification: The abstract and introduction accurately outline the framework's capabilities
442 and contributions, which are then substantiated by the experimental results, ablation studies.

443 Guidelines:

- 444 • The answer NA means that the abstract and introduction do not include the claims
445 made in the paper.
- 446 • The abstract and/or introduction should clearly state the claims made, including the
447 contributions made in the paper and important assumptions and limitations. A No or
448 NA answer to this question will not be perceived well by the reviewers.
- 449 • The claims made should match theoretical and experimental results, and reflect how
450 much the results can be expected to generalize to other settings.
- 451 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
452 are not attained by the paper.

453 **2. Limitations**

454 Question: Does the paper discuss the limitations of the work performed by the authors?

455 Answer: [Yes]

456 Justification: The "Conclusion and Outlook" section contains a dedicated paragraph that
457 explicitly discusses limitations, including dependence on the underlying LLM's knowledge,
458 potential error accumulation, and the need for expert validation.

459 Guidelines:

- 460 • The answer NA means that the paper has no limitation while the answer No means that
461 the paper has limitations, but those are not discussed in the paper.
- 462 • The authors are encouraged to create a separate "Limitations" section in their paper.
- 463 • The paper should point out any strong assumptions and how robust the results are to
464 violations of these assumptions (e.g., independence assumptions, noiseless settings,
465 model well-specification, asymptotic approximations only holding locally). The authors
466 should reflect on how these assumptions might be violated in practice and what the
467 implications would be.
- 468 • The authors should reflect on the scope of the claims made, e.g., if the approach was
469 only tested on a few datasets or with a few runs. In general, empirical results often
470 depend on implicit assumptions, which should be articulated.
- 471 • The authors should reflect on the factors that influence the performance of the approach.
472 For example, a facial recognition algorithm may perform poorly when image resolution
473 is low or images are taken in low lighting.
- 474 • The authors should discuss the computational efficiency of the proposed algorithms
475 and how they scale with dataset size.
- 476 • If applicable, the authors should discuss possible limitations of their approach to
477 address problems of privacy and fairness.
- 478 • While the authors might fear that complete honesty about limitations might be used by
479 reviewers as grounds for rejection, a worse outcome might be that reviewers discover
480 limitations that aren't acknowledged in the paper. Reviewers will be specifically
481 instructed to not penalize honesty concerning limitations.

482 **3. Theory assumptions and proofs**

483 Question: For each theoretical result, does the paper provide the full set of assumptions and
484 a complete (and correct) proof?

485 Answer: [NA]

486 Justification: It does not contain theoretical results, theorems, or formal proofs.

487 Guidelines:

- 488 • The answer NA means that the paper does not include theoretical results.
489 • All the theorems, formulas, and proofs in the paper should be numbered and cross-
490 referenced.
491 • All assumptions should be clearly stated or referenced in the statement of any theorems.
492 • The proofs can either appear in the main paper or the supplemental material, but if
493 they appear in the supplemental material, the authors are encouraged to provide a short
494 proof sketch to provide intuition.

495 **4. Experimental result reproducibility**

496 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
497 perimental results of the paper to the extent that it affects the main claims and/or conclusions
498 of the paper (regardless of whether the code and data are provided or not)?

499 Answer: [Yes]

500 Justification: The paper provides a detailed description of the tasks, baselines, LLM models
501 used, and the LLM-as-a-Judge evaluation protocol, including the use of a deterministic
502 evaluation mode. The prompts and scripts are available in the supplementary material.

503 Guidelines:

- 504 • The answer NA means that the paper does not include experiments.
505 • If the paper includes experiments, a No answer to this question will not be perceived
506 well by the reviewers: Making the paper reproducible is important.
507 • If the contribution is a dataset and/or model, the authors should describe the steps taken
508 to make their results reproducible or verifiable.
509 • We recognize that reproducibility may be tricky in some cases, in which case authors
510 are welcome to describe the particular way they provide for reproducibility. In the case
511 of closed-source models, it may be that access to the model is limited in some way
512 (e.g., to registered users), but it should be possible for other researchers to have some
513 path to reproducing or verifying the results.

514 **5. Open access to data and code**

515 Question: Does the paper provide open access to the data and code, with sufficient instruc-
516 tions to faithfully reproduce the main experimental results, as described in supplemental
517 material?

518 Answer: [Yes]

519 Justification: The paper explicitly states that "Prompts, orchestration scripts, and scoring
520 templates are released in the supplementary material" to facilitate the reproduction of the
521 presented results.

522 Guidelines:

- 523 • The answer NA means that paper does not include experiments requiring code.
524 • Please see the Agents4Science code and data submission guidelines on the conference
525 website for more details.
526 • While we encourage the release of code and data, we understand that this might not be
527 possible, so "No" is an acceptable answer. Papers cannot be rejected simply for not
528 including code, unless this is central to the contribution (e.g., for a new open-source
529 benchmark).
530 • The instructions should contain the exact command and environment needed to run to
531 reproduce the results.
532 • At submission time, to preserve anonymity, the authors should release anonymized
533 versions (if applicable).

534 **6. Experimental setting/details**

535 Question: Does the paper specify all the training and test details (e.g., data splits, hyper-
536 parameters, how they were chosen, type of optimizer, etc.) necessary to understand the
537 results?

538 Answer: [Yes]

539 Justification: Section 3 details the experimental setup, including the specific models used for
540 generation and evaluation, the set of tasks, the baseline methods, and the evaluation rubric.
541 As the experiments use pre-trained models, details like training data splits or optimizers are
542 not applicable.

543 Guidelines:

- 544 • The answer NA means that the paper does not include experiments.
- 545 • The experimental setting should be presented in the core of the paper to a level of detail
546 that is necessary to appreciate the results and make sense of them.
- 547 • The full details can be provided either with the code, in appendix, or as supplemental
548 material.

549 7. Experiment statistical significance

550 Question: Does the paper report error bars suitably and correctly defined or other appropriate
551 information about the statistical significance of the experiments?

552 Answer: [Yes]

553 Justification: Yes. All main experimental results reported in Table 1 are presented with both
554 a mean score and the corresponding standard deviation.

555 Guidelines:

- 556 • The answer NA means that the paper does not include experiments.
- 557 • The authors should answer "Yes" if the results are accompanied by error bars, confi-
558 dence intervals, or statistical significance tests, at least for the experiments that support
559 the main claims of the paper.
- 560 • The factors of variability that the error bars are capturing should be clearly stated
561 (for example, train/test split, initialization, or overall run with given experimental
562 conditions).

563 8. Experiments compute resources

564 Question: For each experiment, does the paper provide sufficient information on the com-
565 puter resources (type of compute workers, memory, time of execution) needed to reproduce
566 the experiments?

567 Answer: [Yes]

568 Justification: For these API-based experiments, the paper provides the most critical resource
569 information: the specific API provider and model endpoints (Volcano Engine).

570 Guidelines:

- 571 • The answer NA means that the paper does not include experiments.
- 572 • The paper should indicate the type of compute workers CPU or GPU, internal cluster,
573 or cloud provider, including relevant memory and storage.
- 574 • The paper should provide the amount of compute required for each of the individual
575 experimental runs as well as estimate the total compute.

576 9. Code of ethics

577 Question: Does the research conducted in the paper conform, in every respect, with the
578 Agents4Science Code of Ethics (see conference website)?

579 Answer: [Yes]

580 Justification: The research focuses on a general-purpose reasoning framework for construc-
581 tive applications like education and crisis management, and there is no indication of any
582 ethical violations.

583 Guidelines:

- 584 • The answer NA means that the authors have not reviewed the Agents4Science Code of
585 Ethics.
- 586 • If the authors answer No, they should explain the special circumstances that require a
587 deviation from the Code of Ethics.

588 10. Broader impacts

589 Question: Does the paper discuss both potential positive societal impacts and negative
590 societal impacts of the work performed?

591 Answer: [Yes]

592 Justification: The paper discusses positive impacts in the conclusion. Potential negative
593 impacts are addressed in the discussion of the base model's inherent limitations, which are
594 the source of potential misuse and bias.

595 Guidelines:

- 596 • The answer NA means that there is no societal impact of the work performed.
- 597 • If the authors answer NA or No, they should explain why their work has no societal
598 impact or why the paper does not address societal impact.
- 599 • Examples of negative societal impacts include potential malicious or unintended uses
600 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations,
601 privacy considerations, and security considerations.
- 602 • If there are negative societal impacts, the authors could also discuss possible mitigation
603 strategies.