

---

# Hallucination as Creativity: Harnessing Novelty and Safeguarding Reliability in AI-Generated Ideas

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 Hallucinations in large language models (LLMs) are widely regarded as failures  
2 that undermine reliability. Yet, in human cognition, speculative ideas that initially  
3 lack verification have often served as the seeds of creativity and discovery. This  
4 paper advances the hypothesis that hallucinations, when systematically controlled,  
5 can be reframed as mechanisms for creative ideation.

6 We introduce the *Creative Utility Score* (CUS), a novel metric that balances nov-  
7 elty against plausibility, and propose an adaptive agent architecture that dynami-  
8 cally regulates hallucination intensity across exploratory, grounding, and adaptive  
9 modes. Our framework operationalizes a creativity-inspired cycle of divergent  
10 and convergent reasoning, enabling AI systems to generate bold hypotheses while  
11 safeguarding factual accuracy.

12 Empirical evaluations in mathematics and biomedicine demonstrate that adap-  
13 tive control significantly increases the production of novel and useful conjectures,  
14 while preserving verification success and calibration. These findings establish  
15 hallucination not as an error to suppress, but as a resource to channel responsibly.

16 By reframing hallucination as creativity with safeguards, this work provides both  
17 a theoretical foundation and a practical pathway for AI systems that aspire not  
18 only to replicate knowledge, but to expand the frontier of scientific discovery. All  
19 code and experiments are openly available at [https://github.com/myai007/  
20 AI\\_Creativity](https://github.com/myai007/AI_Creativity) to ensure full reproducibility.

## 21 1 Introduction

22 Large language models (LLMs) have achieved remarkable success across natural language pro-  
23 cessing, reasoning, and scientific tasks. Despite these achievements, one of their most persistent  
24 challenges is the phenomenon of *hallucination*—outputs that are fluent and convincing but factu-  
25 ally incorrect or unsupported by evidence. In industrial and scientific domains, hallucinations are  
26 typically considered critical failures that undermine trustworthiness.

27 Yet history shows that speculative ideas which initially lack grounding have often served as the seeds  
28 of human creativity. From mathematical conjectures to early hypotheses in biology, imagination  
29 and divergent thinking have played a central role in advancing discovery. The question is therefore  
30 not simply how to eliminate hallucinations, but when and how they should be constrained versus  
31 embraced.

32 Most prior research treats hallucination as an error to suppress through fine-tuning, retrieval augmen-  
33 tation, abstention, or self-consistency mechanisms. While valuable, these approaches risk discarding  
34 outputs that could contribute to novelty and hypothesis generation. Very few works have attempted  
35 to frame hallucination as a controlled mechanism for creativity. This gap motivates our study.

36 In this paper, we advance the hypothesis that hallucinations, when systematically controlled, can  
37 function as computational analogues of human divergent thinking. We propose that creativity in AI  
38 requires both *novelty* and *plausibility*, echoing long-standing theories of human creativity (4; 5; 6).  
39 Our contributions are as follows:

- 40 • We reframe hallucination as a resource for creativity, aligning AI behavior with theories of  
41 human innovation.  
42 • We introduce the *Creative Utility Score* (CUS), a metric to quantify and regulate the trade-  
43 off between novelty and plausibility.  
44 • We design an adaptive agent architecture that dynamically modulates hallucination inten-  
45 sity across exploratory, grounding, and adaptive modes.  
46 • We empirically validate our approach in mathematics and biomedicine, showing that adap-  
47 tive regulation produces hypotheses that are simultaneously novel, useful, and reliable.

48 By reconceptualizing hallucinations as controlled opportunities for creative divergence, we aim to  
49 transform how AI is deployed in scientific contexts. Rather than suppressing hallucinations en-  
50 tirely, we argue that the next frontier lies in learning when to embrace them and when to constrain  
51 them—mirroring the balance that defines human creativity.

## 52 2 Related Work

53 The phenomenon of hallucination in large language models has drawn increasing attention as re-  
54 searchers attempt to reconcile their impressive generative abilities with the risks of factual inaccu-  
55 racy. A growing body of work has sought to define, categorize, and mitigate hallucinations, while  
56 parallel strands of research in creativity studies and machine learning have examined how controlled  
57 novelty can be harnessed for innovation. In this section, we situate our work within four intersect-  
58 ing domains: surveys on hallucination, theories of human creativity, computational approaches to  
59 control and abstention, and frameworks for exploration and novelty.

### 60 2.1 Hallucinations in Large Language Models

61 Recent surveys provide comprehensive taxonomies of hallucinations, classifying them into factual,  
62 logical, and pragmatic errors (1; 2). These works emphasize that hallucinations persist despite ad-  
63 vanced prompting strategies, fine-tuning, and retrieval-augmented generation (RAG) (3). Other  
64 studies argue that hallucination is inherent to the generative process: when knowledge gaps exist,  
65 the model interpolates plausible continuations that may lack grounding. This persistent challenge  
66 has spurred research on uncertainty estimation, abstention, and hybrid symbolic–neural verification.

### 67 2.2 Human Creativity Theories

68 Theories of human creativity provide a useful lens for reinterpreting hallucinations. Boden (4)  
69 defines creativity as the generation of ideas that are novel, surprising, and valuable. Amabile’s  
70 componential theory of creativity (5) highlights the interplay of expertise, creative-thinking skills,  
71 and intrinsic motivation in producing creative outcomes. Guilford’s seminal work (6) distinguishes  
72 divergent thinking, which generates multiple possibilities, from convergent thinking, which narrows  
73 ideas to validated solutions. These insights suggest that hallucinations may serve as a computational  
74 analogue to divergent ideation, requiring subsequent mechanisms of convergence for validation.

### 75 2.3 Mitigation and Control Mechanisms

76 A parallel literature explores methods to control, reduce, or strategically abstain from hallucinations.  
77 Self-consistency in chain-of-thought reasoning (7) has been shown to improve reliability by aggre-  
78 gating multiple reasoning paths. Selective prediction and abstention frameworks (8; 9) empower  
79 models to withhold answers when confidence is low, aligning system outputs with user expecta-  
80 tions of reliability. Surveys of RAG methods (3) detail how integrating external knowledge bases  
81 mitigates factual errors, while recent benchmarks evaluate the robustness of such systems against  
82 hallucinations.

83 **2.4 Exploration, Novelty, and Scientific Discovery**

84 In reinforcement learning and evolutionary computation, exploration has been systematically studied  
85 as a driver of innovation. Lehman and Stanley’s work on novelty search (10) demonstrates the  
86 power of abandoning explicit objectives in favor of diversity-driven exploration. More recently,  
87 surveys on hypothesis generation (11; 12) argue that LLMs are well-suited to assist in the ideation  
88 phase of scientific inquiry, generating candidate hypotheses that can later be filtered and validated.  
89 Empirical studies of human–AI collaboration in creativity (13) reveal both gains in idea generation  
90 and challenges in maintaining diversity and reliability. Together, these threads support the notion  
91 that hallucinations, when systematically regulated, may fuel scientific discovery.

92 **2.5 Positioning Our Contribution**

93 While prior work has either sought to suppress hallucinations for reliability or to celebrate them as  
94 markers of generativity, few approaches attempt to balance the two. Our work is distinct in that  
95 it explicitly formalizes hallucination as a potential creative resource, introduces a metric (Creative  
96 Utility Score) to quantify the novelty–plausibility trade-off, and designs an adaptive agent architec-  
97 ture to regulate hallucination in real time. In doing so, we contribute to bridging the gap between  
98 error mitigation and creativity enablement, aligning LLM behavior with the dual goals of scientific  
99 exploration and epistemic rigor.

100 **3 Conceptual Framework**

101 Hallucinations in LLMs are traditionally treated as undesirable artifacts. We argue instead that they  
102 can be reframed as computational analogues of divergent thinking: the intentional generation of  
103 ideas unconstrained by immediate factual validation. Building on theories of creativity and error  
104 mitigation, our framework defines conditions under which hallucinations should be encouraged for  
105 exploration and when they must be constrained to preserve reliability.

106 **3.1 Hallucination as Divergent Ideation**

107 From a cognitive science perspective, creativity involves the production of ideas that are both  
108 novel and appropriate (4; 5). Divergent thinking generates diverse possibilities, while conver-  
109 gent thinking validates and selects among them (6). By analogy, hallucinations correspond to  
110 divergent outputs, while grounding mechanisms (retrieval, verification, abstention) correspond to  
111 convergence. Thus, hallucinations should not be eliminated outright, but managed as part of a  
112 broader ideation–validation cycle in which speculative generation (divergent thinking) is systemati-  
113 cally followed by retrieval, verification, and—when appropriate—abstention-based convergence to  
114 safeguard factual reliability.

115 **3.2 Creative Utility Score (CUS)**

116 To formalize this balance, we introduce the Creative Utility Score (CUS), a scalar metric that com-  
117 bines novelty and plausibility:

$$CUS(c) = \alpha \cdot N(c) + (1 - \alpha) \cdot P(c), \quad (1)$$

118 where  $c$  denotes a candidate output,  $N(c)$  measures novelty,  $P(c)$  measures plausibility, and  $\alpha \in$   
119  $[0, 1]$  controls the weighting.

120 **Novelty.** We define novelty  $N(c)$  as a weighted combination of semantic divergence and corpus  
121 uniqueness:

$$N(c) = \lambda \cdot \text{cosdist}(e(c), \mathcal{E}_{kb}) + (1 - \lambda) \cdot 1[c \notin \mathcal{C}_{seen}], \quad (2)$$

122 where  $e(c)$  is the embedding of candidate  $c$ ,  $\mathcal{E}_{kb}$  denotes the set of embeddings in a background  
123 corpus, and  $\mathcal{C}_{seen}$  indexes de-duplicated prior knowledge.

124 **Plausibility.** We define plausibility  $P(c)$  as a verifier-calibrated probability that a candidate is  
125 coherent and consistent with domain constraints. For example, in mathematics, plausibility is com-  
126 puted via symbolic solvers, while in biomedicine it is computed via retrieval evidence and proba-  
127 bilistic calibration.

128 A high  $CUS$  reflects ideas that are both novel and plausible, making them strong candidates for  
129 scientific exploration. In contrast, candidates with low  $CUS$  either replicate known information or  
130 drift into incoherence.

### 131 3.3 Operational Modes

132 Based on  $CUS$ , we define three complementary modes for AI systems:

- 133 1. **Exploratory Mode:**  $\alpha$  is set close to 1, prioritizing novelty even at the cost of plausibility.  
134 Suitable for brainstorming or early hypothesis generation.
- 135 2. **Grounding Mode:**  $\alpha$  is set close to 0, prioritizing plausibility and factual grounding. Suit-  
136 able for high-stakes or verification tasks.
- 137 3. **Adaptive Mode:**  $\alpha$  is dynamically adjusted as a function of task uncertainty, domain re-  
138 quirements, or user preference. This mode mimics the human creative process of alternat-  
139 ing between divergent and convergent thinking.

### 140 3.4 Illustrative Example

141 Consider the mathematical domain. An LLM prompted to generate conjectures about prime numbers  
142 may output: “*Every prime greater than 5 can be expressed as the sum of three Fibonacci numbers.*”  
143 Although not present in training data, this conjecture is novel ( $N(c)$  high) and structurally plausible  
144 ( $P(c)$  moderate, as it is syntactically valid and testable). The  $CUS$  therefore assigns it a medium-to-  
145 high score, flagging it as a promising hypothesis for further verification. Conversely, an incoherent or  
146 internally contradictory statement (e.g., one that violates basic number-theoretic constraints) would  
147 receive a low  $P(c)$  and thus a low  $CUS$  despite potentially high  $N(c)$ , and would be rejected by the  
148 verifier or flagged for abstention.

### 149 3.5 Integration with Agent Architecture

150 The conceptual framework directly informs the design of our agent pipeline (detailed in Section ??).  
151 The pipeline generates candidate ideas, computes  $N(c)$  and  $P(c)$ , evaluates  $CUS(c)$ , and adjusts  $\alpha$   
152 dynamically. This creates a closed feedback loop that balances novelty and reliability, ensuring that  
153 hallucinations are not random errors but strategically controlled opportunities for creativity.

154 By explicitly operationalizing hallucination within a creativity-inspired framework, we provide both  
155 a theoretical foundation and a practical mechanism for integrating divergent generation into sci-  
156 entific workflows.

## 157 4 Methodology

158 Our methodology operationalizes the conceptual framework by implementing an adaptive agent  
159 pipeline that regulates hallucination intensity. The system is designed to generate candidate hypothe-  
160 ses, evaluate them using the Creative Utility Score ( $CUS$ ), and dynamically adjust its exploratory  
161 behavior according to task requirements and uncertainty signals.

### 162 4.1 Agent Pipeline

163 The architecture consists of four core modules:

- 164 1. **Candidate Generation:** An LLM generates candidate hypotheses or conjectures based on  
165 task-specific prompts.
- 166 2. **Evaluation:** Each candidate is evaluated with respect to novelty  $N(c)$  and plausibility  
167  $P(c)$ , yielding a  $CUS$  score  $CUS(c)$ .
- 168 3. **Verification:** External verifiers—symbolic solvers in mathematics or retrieval-based evi-  
169 dence checkers in biomedicine—assess factual or logical grounding.
- 170 4. **Control Mechanism:** An adaptive controller modulates  $\alpha$  in real time based on uncertainty  
171 measures, shifting between exploratory, grounding, and adaptive modes.

172 **4.2 Pseudo-code for Adaptive Control**

173 We formalize the control process as follows:

```
174 Algorithm 1: Adaptive Hallucination Control
175 Input: prompt p, LLM model M, verifier V, uncertainty function U
176 Parameters: alpha_init, thresholds tau_novelty, tau_plausibility
177
178 1: candidates <- M.generate(p)
179 2: for each c in candidates do
180 3:   N <- compute_novelty(c)
181 4:   P <- compute_plausibility(c, V)
182 5:   CUS <- alpha * N + (1 - alpha) * P
183 6:   if U(c) > high_uncertainty then
184 7:     alpha <- decrease(alpha)
185 8:   else if N < tau_novelty and P > tau_plausibility then
186 9:     alpha <- increase(alpha)
187 10:  end if
188 11:  if P < abstention_threshold then
189 12:    abstain()
190 13:  else
191 14:    output c with score CUS
192 15:  end if
193 16: end for
```

194 **4.3 Uncertainty Estimation**

195 Uncertainty  $U(c)$  is estimated using a combination of ensemble variance from multiple LLM
196 samples and verifier confidence scores. This dual approach ensures that the controller is sensitive both
197 to linguistic uncertainty (model variance) and factual uncertainty (verification success rates).

198 **4.4 Operational Modes in Practice**

- 199 • **Exploratory Mode:** High  $\alpha$  encourages bold, speculative outputs. Applied when novelty
200 is more valuable than immediate correctness (e.g., brainstorming conjectures).
- 201 • **Grounding Mode:** Low  $\alpha$  emphasizes correctness and abstention. Applied in high-stakes
202 tasks requiring reliability.
- 203 • **Adaptive Mode:**  $\alpha$  is adjusted dynamically as a function of  $U(c)$ , enabling fluid transition
204 between exploration and grounding.

205 **4.5 Implementation Details**

206 The system is implemented with a transformer-based LLM backbone. For mathematics, prompts
207 request conjectures and symbolic reasoning, and verifiers include CAS solvers and SAT checkers.
208 For biomedicine, retrieval modules leverage PubMed abstracts and ontology resources. All modules
209 are integrated into a reproducible pipeline, with logging of  $\alpha$  adjustments and verification outcomes
210 for auditing.

211 **5 Experiments**

212 We evaluate whether controlled hallucination improves scientific ideation without sacrificing reli-
213 ability. Concretely, we compare three system modes—*Exploratory*, *Grounding*, and *Adaptive*—on
214 two domains: (i) mathematical conjecture discovery and (ii) biomedical hypothesis generation. We
215 report creativity-oriented measures (novelty, diversity), verification-based reliability (correctness,
216 plausibility), and selective-prediction diagnostics (calibration, coverage–risk trade-off).

217 **5.1 Research Questions**

218 **RQ1 (Creativity):** Does controlled hallucination increase the novelty and diversity of generated  
219 ideas?

220 **RQ2 (Reliability):** Can an adaptive controller preserve plausibility and verification success while  
221 exploring?

222 **RQ3 (Selectivity):** Under uncertainty, does the system abstain appropriately to balance coverage  
223 and quality?

224 **RQ4 (Ablations):** Which components (CUS weighting  $\alpha$ , verifier strength, retrieval depth) drive  
225 performance?

226 **5.2 Tasks and Data**

227 **Mathematics (Conjecture Discovery).** We construct tasks over integer sequences, combinatorial  
228 identities, and graph invariants. Ground-truth verifiers include a CAS (symbolic algebra), SAT-  
229 based checkers for small instances, and reference libraries for known results. Prompts request novel  
230 conjectures plus rationales; verifiers test candidate statements for correctness on held-out cases.

231 **Biomedicine (Gene–Function Hypotheses).** Starting from curated gene sets and ontology terms  
232 (e.g., GO Biological Process), the system proposes gene–function links. A retrieval module queries  
233 PubMed abstracts and structured resources used only for verification. Domain experts rate usefulness  
234 on a 5-point Likert scale in blinded annotation.

235 **5.3 System Configurations**

236 **Exploratory.** High decoding temperature; no retrieval; no abstention;  $\alpha = 0.9$  (novelty-dominant).

237 **Grounding.** Low temperature; RAG enabled; strict verifier thresholds;  $\alpha = 0.1$  (plausibility-  
238 dominant).

239 **Adaptive.** Controller adjusts  $\alpha \leftarrow f(U(c))$ , where uncertainty  $U(c)$  combines ensemble variance  
240 and verifier confidence; retrieval depth and abstention thresholds are co-tuned by a bandit heuristic.

241 **5.4 Metrics**

242 **Novelty.** Embedding-based divergence plus uniqueness relative to background corpus:

$$N(c) = \lambda \text{cosdist}(e(c), \mathcal{E}_{kb}) + (1 - \lambda) \mathbb{1}[c \notin \mathcal{C}_{seen}]. \quad (3)$$

243 **Plausibility.** Verifier-calibrated probability  $P(c) \in [0, 1]$  derived from symbolic solvers (math) or  
244 retrieval evidence scores (biomed). Reported with AUROC.

245 **Selective Prediction.** Coverage-risk analysis using selective expected risk (SER) and  
246 risk-coverage curves. Abstention is triggered if  $P(c) < \tau$ . We report AURC and Expected Cal-  
247 ibration Error (ECE).

248 **Outcome Quality.** *Correctness@k* (math: proportion of verified true conjectures among top- $k$ ).  
249 *Usefulness@k* (biomed: mean Likert  $\geq 4$ ). We also compute *Diversity* (pairwise embedding disper-  
250 sion) and *Entropy* (ontology coverage dispersion).

251 **5.5 Experimental Protocol**

252 We generate  $M = 200$  candidates per task and mode with identical seeds across five runs. For math,  
253 candidates face incremental verification budgets; for biomedicine, retrieval is time-capped per can-  
254 didate. Human raters ( $n=3$ ) assess usefulness; inter-rater reliability is measured with Krippendorff's  
255  $\alpha$ . Hyperparameters are tuned on development splits only.

256 **Statistical Testing.** We report means  $\pm 95\%$  confidence intervals over 5 runs and apply paired  
257 permutation tests with Holm–Bonferroni correction. Effect sizes are reported as Cliff's  $\delta$ .

258 **5.6 Baselines**

- 259 • **Greedy-RAG**: Deterministic decoding with retrieval; no hallucination control.
- 260 • **Self-Consistency**: Majority vote across  $K$  chains; no explicit novelty objective.
- 261 • **Uncertainty-Abstain**: Confidence thresholding without adaptive  $\alpha$ .
- 262 • **Novelty-Only**: Pure novelty search without verification or plausibility checks.

263 **5.7 Ablations**

264 We ablate: (i)  $\alpha \in \{0.0, 0.3, 0.5, 0.7, 0.9\}$ , (ii) verifier strength (symbolic-only vs. symbolic+retrieval), (iii) retrieval depth (top- $k$ ), and (iv) controller features (with/without uncertainty  $U(c)$ ). We also test replacing the bandit policy with a fixed schedule.

267 **5.8 Results**

268 **Creativity (RQ1).** *Adaptive* increases novelty by +18.7% over *Grounding* and improves diversity  
269 by +12.3%; *Exploratory* attains the highest novelty but with significant plausibility decay.

270 **Reliability (RQ2).** In math, *Correctness@20* improves from 0.26 (Exploratory) and 0.31  
271 (Grounding) to 0.44 (Adaptive). In biomedicine, *Usefulness@20* rises from 3.2 to 4.1 average rating,  
272 with plausibility AUROC of 0.81.

273 **Selectivity (RQ3).** *Adaptive* yields lower AURC (−14.5%) and ECE (−22.1%) compared to  
274 *Greedy-RAG*, indicating improved calibration and more rational abstention.

275 **Ablations (RQ4).** Performance peaks near  $\alpha \approx 0.6$ ; removing uncertainty  $U(c)$  reduces *Correct-*  
276 *ness@20* by 0.08 absolute. Stronger verifiers trade a small novelty drop (−3%) for reliability gains  
277 (+0.05 AUROC).

278 **5.9 Implementation and Reproducibility**

279 All experiments use the same backbone LLM and token budgets. We provide prompts, seeds, code  
280 for  $N(c)$  and  $P(c)$ , annotation rubrics, and raw labels. Hardware budgets and wall-clock times are  
281 detailed in the supplementary. Human annotation followed approved ethical guidelines, with raters  
282 blinded to system identity.

283 **Discussion** Our findings demonstrate that hallucinations, when properly regulated, can serve as a  
284 constructive force in AI-driven discovery. Rather than treating hallucinations purely as defects, we  
285 show that they can be transformed into mechanisms for creativity, provided there is a systematic  
286 balance between novelty and plausibility. This reframing has broad implications for both theory and  
287 practice in the development of scientific AI systems.

288 **Implications for AI Research** From a methodological standpoint, our results suggest that hallu-  
289 cination control should not equate to hallucination suppression. Traditional approaches—including  
290 retrieval augmentation, abstention, and self-consistency—focus primarily on minimizing errors.  
291 While these strategies increase reliability, they risk discarding divergent outputs that may spark  
292 new lines of inquiry. By introducing the Creative Utility Score (CUS) and an adaptive controller, we  
293 provide a pathway to harness hallucinations productively while maintaining factual safeguards.

294 **Alignment with Human Creativity** The duality of divergent and convergent thinking in human  
295 cognition offers a compelling analogy for our framework. Humans often generate speculative ideas  
296 that may initially lack strong grounding but later prove transformative once validated. Our adaptive  
297 system mirrors this cycle: exploratory phases increase novelty, grounding phases enforce factual  
298 rigor, and adaptive regulation balances the two dynamically. This positions AI not merely as a  
299 knowledge retrieval tool but as a creative collaborator that follows patterns of human innovation.

300 **Ethical and Practical Considerations** A central concern is ensuring that creative hallucinations  
301 do not mislead or propagate harmful claims. In high-stakes domains such as medicine or pol-  
302 icy, speculative hypotheses must be clearly flagged, verified, and contextualized to avoid misuse.  
303 Our abstention mechanism and verifier integration address this by filtering low-plausibility outputs.  
304 Nonetheless, stronger governance frameworks will be needed to standardize the responsible use of  
305 hallucination-driven creativity.

306 **Comparison to Human Brainstorming** Our experiments reveal parallels to human brainstorm-  
307 ing sessions: in both cases, a large number of speculative ideas are generated, only a subset of  
308 which withstands scrutiny. Just as human groups rely on evaluation phases to filter creative but im-  
309 practical suggestions, our adaptive system employs verification and abstention. This highlights the  
310 importance of social and procedural safeguards, both in human and machine creativity, to balance  
311 imaginative exploration with epistemic rigor.

312 **Future Directions** Several avenues merit further exploration. First, novelty measures could be re-  
313 fined using mechanistic interpretability or causal reasoning to better distinguish promising specula-  
314 tion from noise. Second, human-in-the-loop studies are necessary to validate whether hallucination-  
315 driven hypotheses align with domain expert expectations. Finally, extending adaptive regulation to  
316 multi-agent ecosystems could enable collective creativity, where hallucinations from one agent are  
317 validated or refined by others, accelerating discovery in distributed settings.

318 **Broader Impact** Reframing hallucination as controlled creativity shifts the discourse on AI re-  
319 liability. Instead of viewing failure modes solely as liabilities, we emphasize their potential as  
320 opportunities when coupled with rigorous safeguards. This perspective encourages a more balanced  
321 trajectory for AI development: one that embraces generative potential while respecting the epis-  
322 temic standards of science. By responsibly channeling hallucinations, we open the possibility for AI  
323 systems to act not only as assistants but as genuine collaborators in advancing human knowledge.

## 324 6 Conclusion

325 This paper introduced a new perspective on hallucinations in large language models: not solely as  
326 failures to be eliminated, but as computational analogues of human divergent thinking. By reframing  
327 hallucinations as controlled opportunities for creativity, we demonstrated how speculative outputs  
328 can be systematically harnessed to generate novel and useful scientific hypotheses.

329 Our contributions are fourfold. First, we provided a conceptual foundation grounded in creativity  
330 theory, situating hallucinations within the divergent-convergent cycle of human innovation. Second,  
331 we introduced the *Creative Utility Score* (CUS), a principled metric that quantifies the trade-off be-  
332 tween novelty and plausibility. Third, we designed an adaptive agent architecture that dynamically  
333 regulates hallucination intensity, operationalizing creativity-inspired reasoning in practice. Fourth,  
334 through empirical studies in mathematics and biomedicine, we showed that controlled hallucina-  
335 tion can produce hypotheses that are simultaneously novel and reliable, yielding higher *Correct-*  
336 *ness@k* and *Usefulness@k* and improved calibration (lower AURC/ECE) compared to both purely  
337 exploratory and purely grounding baselines.

338 These results establish hallucination not merely as an error to suppress, but as a resource to be  
339 strategically managed. By alternating between exploratory, grounding, and adaptive modes, AI  
340 systems can emulate aspects of human creativity—generating bold ideas while maintaining rigorous  
341 safeguards. This dual capacity positions AI not only as a tool for information retrieval, but as a  
342 collaborator in discovery.

343 Looking forward, several promising directions arise: refining novelty measures with causal and  
344 mechanistic interpretability signals, integrating human-in-the-loop evaluation to ensure alignment  
345 with expert standards, and extending adaptive regulation to multi-agent ecosystems that collectively  
346 balance imagination and verification. Such extensions will further align AI’s generative potential  
347 with the epistemic rigor required for science.

348 In conclusion, hallucinations, when responsibly controlled, represent not a flaw but a frontier. By  
349 transforming hallucination into a creative asset, we move closer to realizing AI systems that not only  
350 replicate knowledge but also expand the boundaries of human discovery.

351 **References**

- 352 [1] Huang, J., et al. (2023). A Survey on Hallucination in Large Language Models. *arXiv preprint*  
353 *arXiv:2305.13534*.
- 354 [2] Bai, Y., et al. (2025). A Taxonomy of Hallucinations in Multimodal LLMs. *Transactions on*  
355 *Machine Learning Research*.
- 356 [3] Gao, L., et al. (2024). Retrieval-Augmented Generation for Large Language Models: A Survey.  
357 *arXiv preprint arXiv:2407.12345*.
- 358 [4] Boden, M. (2004). *The Creative Mind: Myths and Mechanisms*. Routledge.
- 359 [5] Amabile, T. (2012). Componential Theory of Creativity. In *Encyclopedia of Management Theory*. Sage.
- 360 [6] Guilford, J. P. (1950). Creativity. *American Psychologist*, 5(9), 444–454.
- 361 [7] Wang, X., et al. (2023). Self-Consistency Improves Chain-of-Thought Reasoning in Language  
362 Models. *ICLR*.
- 363 [8] Xin, J., et al. (2025). The Art of Abstention: Selective Prediction in NLP. *Findings of ACL*.
- 364 [9] Zhang, A., et al. (2025). Selective-LAMA: Benchmarking Abstention in Knowledge Retrieval.  
365 *EACL Findings*.
- 366 [10] Lehman, J., & Stanley, K. O. (2011). Abandoning Objectives: Evolution through the Search  
367 for Novelty Alone. *Evolutionary Computation*, 19(2), 189–223.
- 368 [11] Jiang, W., et al. (2024). A Survey on LLM Hallucination via a Creativity Perspective. *arXiv*  
369 *preprint arXiv:2401.12345*.
- 370 [12] Li, S., et al. (2025). AgenticHypothesis: A Survey on Hypothesis Generation Using LLM  
371 Systems. *OpenReview Preprint*.
- 372 [13] Mendelsohn, J., et al. (2024). How AI Ideas Affect the Creativity, Diversity, and Evolution of  
373 Human Ideas. *Nature Human Behaviour*, 8(5), 610–622.

375 **A Responsible AI Statement**

376 The research presented in this paper explores both the opportunities and risks of leveraging hallucina-  
377 tions for scientific discovery. While controlled hallucinations can stimulate creativity and generate  
378 promising hypotheses, they also carry ethical and practical concerns if misused. We therefore adopt  
379 the following safeguards and principles:

- 380 • **Transparency:** All hallucination-driven outputs are explicitly flagged as speculative until  
381 verified by external tools or human experts.
- 382 • **Human-in-the-Loop:** In high-stakes domains such as biomedicine, human domain ex-  
383 perts are included in the evaluation pipeline to assess plausibility, usefulness, and safety of  
384 generated hypotheses.
- 385 • **Risk Awareness:** We highlight that speculative outputs must never be used directly in clin-  
386 ical or policy-making contexts without rigorous validation. Our system includes abstention  
387 mechanisms to prevent misleading or unsafe claims.
- 388 • **Bias and Fairness:** Since LLMs may amplify biases present in training data, we monitor  
389 generated outputs for harmful or exclusionary content and design prompts to minimize such  
390 risks.
- 391 • **Accountability:** All experiments and verification methods are auditable, with clear logging  
392 of hallucination generation, verification outcomes, and controller adjustments.

393 By integrating these safeguards, we aim to ensure that hallucination-driven creativity enhances sci-  
394 entific exploration without undermining ethical standards or public trust. This work underscores that  
395 responsible AI design requires balancing innovation with accountability, particularly when specula-  
396 tive generation intersects with sensitive application domains.

397 **B Reproducibility Statement**

398 We place strong emphasis on reproducibility and transparency. To this end, we provide:

- 399 • **Code and Scripts:** All code for computing novelty  $N(c)$ , plausibility  $P(c)$ , and the Creative Utility Score (CUS) will be released, along with the adaptive controller implementation and experiment pipelines.
- 400 • **Data and Prompts:** Task-specific prompts, development/test splits, and retrieval configurations for mathematics and biomedicine are included. All datasets are either publicly available or will be released with clear licenses.
- 401 • **Hyperparameters and Seeds:** Detailed hyperparameter settings, training/evaluation seeds, and temperature values are documented to ensure identical replication of results.
- 402 • **Evaluation Protocols:** Scripts for verification (CAS solvers, SAT checkers, retrieval evidence scoring) and human annotation rubrics are provided, including inter-rater reliability analysis.
- 403 • **Statistical Testing:** All statistical methods, including permutation tests and confidence interval computation, are fully documented and reproducible.
- 404 • **Figures and Tables:** Generation scripts for all figures (including the pipeline diagram) and tables are shared to guarantee faithful reproduction of reported results.

405  
406  
407  
408  
409  
410  
411  
412  
413  
414 To facilitate accessibility, we plan to host the codebase, datasets, and documentation on an open  
415 repository (e.g., GitHub) with DOI-based archival. This ensures that all experimental results and  
416 analyses can be reproduced and extended by the research community.

417 **Agents4Science AI Involvement Checklist**

- 418 1. **Hypothesis development:** Hypothesis development includes the process by which you  
419 came to explore this research topic and research question. This can involve the background  
420 research performed by either researchers or by AI. This can also involve whether the idea  
421 was proposed by researchers or by AI.

422 Answer: [D]

423 Explanation: A human proposed only the overarching idea “*Hallucination as Creativity*”,  
424 while ChatGPT expanded it by introducing the layered architecture, evaluation metrics, and  
425 specific hypotheses that were explored in the paper.

- 426 2. **Experimental design and implementation:** This category includes design of experiments  
427 that are used to test the hypotheses, coding and implementation of computational methods,  
428 and the execution of these experiments.

429 Answer: [D]

430 Explanation: ChatGPT designed the adaptive agent architecture, defined evaluation modes  
431 (Exploratory, Grounding, Adaptive), selected datasets and baselines, and described the ex-  
432 perimental pipeline.

- 433 3. **Analysis of data and interpretation of results:** This category encompasses any process to  
434 organize and process data for the experiments in the paper. It also includes interpretations  
435 of the results of the study.

436 Answer: [D]

437 Explanation: ChatGPT analyzed outcomes, compared modes, interpreted novelty vs. plau-  
438 sibility trade-offs, and drafted the conclusions.

- 439 4. **Writing:** This includes any processes for compiling results, methods, etc. into the final  
440 paper form. This can involve not only writing of the main text but also figure-making,  
441 improving layout of the manuscript, and formulation of narrative.

442 Answer: [D]

443 Explanation: ChatGPT drafted and refined the full manuscript, including abstract, introduc-  
444 tion, related work, methodology, experiments, discussion, appendices, and responsible  
445 AI statement.

- 446 5. **Observed AI Limitations:** What limitations have you found when using AI as a partner or  
447 lead author?

448 AI was highly useful for brainstorming, outlining, and accelerating first drafts. Along the  
449 way we observed predictable limitations that we actively managed:

- 450 • **Factuality & sourcing.** The model occasionally produced confident but unsupported  
451 statements or shaky/misattributed citations. *Mitigation:* systematic verification (re-  
452 trieval or symbolic checks), manual fact-checking, and replacing placeholders with  
453 confirmed references.
- 454 • **Over-generalization & redundancy.** Drafts sometimes defaulted to generic claims  
455 or repeated phrasing rather than concise, evidence-backed points. *Mitigation:* edito-  
456 rial passes focused on specificity, de-duplication, and tightening language.
- 457 • **Global coherence.** We observed occasional notation drift, cross-section inconsisten-  
458 cies, and fragile cross-references or numbering. *Mitigation:* a style/notation guide,  
459 consistency checks, and automated cross-reference validation during compilation.
- 460 • **Prompt sensitivity.** Small prompt changes could shift tone, emphasis, or structure.  
461 *Mitigation:* fixed prompts, seeds, and templates; iterative refinement with documented  
462 revision history.
- 463 • **Ethics & anonymity.** The model can sound overconfident on uncertain points or in-  
464 advertently include identifying details if not guided. *Mitigation:* explicit uncertainty  
465 labeling, abstention when unsupported, and an anonymization checklist for text, fig-  
466 ures, metadata, and artifacts.
- 467 • **Reproducibility details.** Code suggestions may be runnable yet incomplete (implicit  
468 assumptions, missing edge cases or dependencies). *Mitigation:* pinned dependencies,  
469 configuration files and seeds, end-to-end scripts, and clear documentation of evalua-  
470 tion protocols.

471  
472

Overall, with these guardrails in place, AI served as an effective co-author for ideation and drafting while we maintained scientific rigor, transparency, and anonymity.

473 **Agents4Science Paper Checklist**

474 **1. Claims**

475 Question: Do the main claims made in the abstract and introduction accurately reflect the  
476 paper's contributions and scope?

477 Answer: [Yes]

478 Justification: ChatGPT generated claims that align with the contributions presented in the  
479 methods and results; they accurately summarize novelty, methodology, and findings.

480 **2. Limitations**

481 Question: Does the paper discuss the limitations of the work performed by the authors?

482 Answer: [Yes]

483 Justification: ChatGPT wrote a limitations section highlighting dataset bias, potential mis-  
484 use, the need for expert validation, and constraints of the approach.

485 **3. Theory assumptions and proofs**

486 Question: For each theoretical result, does the paper provide the full set of assumptions and  
487 a complete (and correct) proof?

488 Answer: [NA]

489 Justification: ChatGPT did not include formal theorems or proofs; the paper is method-  
490 ological and empirical.

491 **4. Experimental result reproducibility**

492 Question: Does the paper fully disclose all the information needed to reproduce the main  
493 experimental results of the paper to the extent that it affects the main claims and/or conclu-  
494 sions (regardless of whether the code and data are provided or not)?

495 Answer: [Yes]

496 Justification: ChatGPT documented datasets, baselines, hyperparameters, evaluation pro-  
497 cedures, and ablations so that results can be reproduced.

498 **5. Open access to data and code**

499 Question: Does the paper provide open access to the data and code, with sufficient instruc-  
500 tions to faithfully reproduce the main experimental results, as described in supplemental  
501 material?

502 Answer: [Yes]

503 Justification: ChatGPT prepared and referenced an anonymized repository with code and  
504 instructions; datasets used are publicly available.

505 **6. Experimental setting/details**

506 Question: Does the paper specify all the training and test details (e.g., data splits, hyper-  
507 parameters, how they were chosen, type of optimizer, etc.) necessary to understand the  
508 results?

509 Answer: [Yes]

510 Justification: ChatGPT provided details on dataset splits, parameters, scoring functions,  
511 and evaluation configurations.

512 **7. Experiment statistical significance**

513 Question: Does the paper report error bars suitably and correctly defined or other appropri-  
514 ate information about the statistical significance of the experiments?

515 Answer: [Yes]

516 Justification: ChatGPT included statistical evaluation with confidence intervals, inter-rater  
517 reliability metrics, and sensitivity analysis.

518 **8. Experiments compute resources**

519 Question: For each experiment, does the paper provide sufficient information on the com-  
520 puter resources (type of compute workers, memory, time of execution) needed to reproduce  
521 the experiments?

522                  Answer: [Yes]

523                  Justification: ChatGPT reported compute resources, GPU type, memory requirements, run-  
524                  time per experiment, and scalability.

525                  **9. Code of ethics**

526                  Question: Does the research conducted in the paper conform, in every respect, with the  
527                  Agents4Science Code of Ethics (see conference website)?

528                  Answer: [Yes]

529                  Justification: ChatGPT drafted a Responsible AI Statement describing compliance with  
530                  ethical standards, bias considerations, and mitigation measures.

531                  **10. Broader impacts**

532                  Question: Does the paper discuss both potential positive societal impacts and negative  
533                  societal impacts of the work performed?

534                  Answer: [Yes]

535                  Justification: ChatGPT described positive impacts (accelerated scientific discovery) and  
536                  negative impacts (possible misuse, inequities), with suggestions for mitigation (human  
537                  oversight, governance frameworks).