
Comparative Analysis of Metaheuristic and Heuristic Strategies in Forest Fire Suppression

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Forest fires represent a significant and escalating global threat, necessitating the
2 development of effective suppression strategies. This paper investigates the ap-
3 plication of computational intelligence, specifically comparing a metaheuristic
4 approach, Ant Colony Optimization (ACO), with a simpler heuristic, a Greedy
5 algorithm, for the strategic placement of firebreaks. Although metaheuristics like
6 ACO are generally anticipated to yield superior solutions for complex optimization
7 problems, simulation results under a specific, constrained scenario—a centrally
8 located fire on a 20x20 grid with a high density of firebreaks—demonstrate that
9 the Greedy strategy unexpectedly outperformed ACO in both minimizing the area
10 burned and the time required for containment. This report analyzes this counterin-
11 tuitive outcome, providing theoretical explanations grounded in the principles of
12 local versus global optimization and contextualizing the findings within the broader
13 optimization literature.

14

1 Introduction

15 Forest fires, exacerbated by global climate change, pose an increasing threat to ecological systems,
16 human populations, and economic stability globally. The escalating frequency and intensity of these
17 catastrophic events necessitate the development of highly effective and efficient suppression strate-
18 gies. Traditional firefighting methods, often reliant on expert judgment and reactive decision-making,
19 frequently contend with cognitive biases, incomplete information, and the inherent complexity of
20 dynamic fire behavior, leading to suboptimal resource allocation and increased operational risks.
21 Computational intelligence, particularly metaheuristic algorithms, offers a promising avenue for
22 overcoming these limitations by providing robust solutions to complex, dynamic optimization prob-
23 lems through extensive exploration of vast solution spaces (Carta et al., 2023). Among these, Ant
24 Colony Optimization (ACO), inspired by the collective foraging behavior of ants, has demonstrated
25 significant potential for identifying optimal paths in dynamic networks. This paper previously in-
26 troduced a conceptual framework that leverages ACO to strategically deploy firefighting resources,
27 with the aim of minimizing overall fire damage by optimizing firebreak placements. For comparative
28 evaluation, a simpler Greedy algorithm was also implemented. Although metaheuristics like ACO
29 are generally expected to outperform simpler heuristics in complex problems due to their global
30 optimization capabilities, recent simulations reveal a scenario where the Greedy strategy achieved
31 superior performance in terms of both area burned and containment time. This report provides a
32 detailed analysis of this outcome, offering theoretical explanations and contextualizing the findings
33 within the broader literature on local versus global optimization in spatial containment problems.

34 **2 Background on optimization algorithms**

35 **2.1 Metaheuristics: ant colony optimization (ACO)**

36 Metaheuristics are high-level frameworks designed to guide the search for solutions to optimization
37 problems, allowing them to escape local optima and explore larger solution spaces. These algorithms
38 are particularly effective for solving NP-hard problems, which are characterized by their computa-
39 tional intractability for exact solutions within practical timeframes. Ant Colony Optimization (ACO),
40 a prominent metaheuristic, derives its inspiration from the ability of ant colonies to find the shortest
41 paths between their nest and food sources through pheromone deposition. In the context of fire
42 suppression, virtual "ants" construct paths representing potential firebreak locations. The probability
43 of an ant at node i choosing to move to node j is determined by the pheromone level (τ_{ij}) and
44 heuristic information (η_{ij}), governed by Equation 1:

$$P_{ij} = \frac{(\tau_{ij}^\alpha)(\eta_{ij}^\beta)}{\sum_{k \in \text{allowed}} (\tau_{ik}^\alpha)(\eta_{ik}^\beta)} \quad (1)$$

45 where α and β are parameters that control the influence of the pheromone trail and heuristic infor-
46 mation, respectively. The heuristic information, η_{ij} , is a composite value based on factors like fuel
47 load and proximity to the fire. Pheromone levels are updated iteratively based on the quality of the
48 solutions found, encouraging convergence towards optimal paths. This adaptive mechanism enables
49 ACO to conduct global searches, making it a robust choice for complex, dynamic problems.

50 **2.2 Simple heuristics: the Greedy approach**

51 In contrast to metaheuristics, simple heuristics, such as the Greedy algorithm, make decisions based
52 on immediate, local information at each step without considering the global implications of these
53 choices. A Greedy algorithm selects the locally optimal choice at each stage with the expectation of
54 finding a global optimum. For instance, in firebreak placement, our Greedy strategy prioritizes nodes
55 based on their Euclidean distance to the fire's origin. The distance d between a node (x_1, y_1) and the
56 fire start (x_f, y_f) is calculated as shown in Equation 2:

$$d = \sqrt{(x_1 - x_f)^2 + (y_1 - y_f)^2} \quad (2)$$

57 While computationally less intensive and faster, Greedy algorithms are susceptible to converging to
58 local optima. Despite this, they can be highly effective in specific scenarios, particularly when the
59 problem structure favors local decisions.

60 **3 Simulation methodology and scenario design**

61 **3.1 Forest model and fire spread dynamics**

62 The forest environment is modeled as a discrete, grid-based graph, $G = (V, E)$, on a 20×20 grid
63 (Alexandridis et al., 2011). The fire spread model is probabilistic. A burning node can ignite an
64 unburned adjacent neighbor with an ignition probability, P_{ignite} , determined by the neighbor's fuel
65 load (f_l) and a wind factor (w_f), as defined in Equation 3:

$$P_{\text{ignite}} = \frac{f_l + w_f}{2.0} \quad (3)$$

66 The `wind_factor` was set to 0.8. Nodes designated with a firebreak status cannot burn, representing
67 defensive lines. The simulation terminates after 50 time steps or when the fire is contained.

68 **3.2 Firebreak placement strategies implemented**

- 69 • **Ant colony optimization (ACO):** The ACO algorithm seeks an optimal set of 25 firebreak
70 nodes. The objective is to minimize a weighted cost function combining burned area
71 (A_{burned}) and containment time (T_{contain}), shown in Equation 4:

$$\text{Cost} = w_1 \cdot A_{\text{burned}} + w_2 \cdot T_{\text{contain}} \quad (4)$$

72 where $w_1 = 0.7$ and $w_2 = 0.3$. The heuristic information incorporates the inverse of a
73 node's fuel load and its proximity to the fire start. A beta parameter of 1.5 was used to
74 emphasize this heuristic information.

- 75 • **Greedy algorithm:** The Greedy algorithm selects 25 firebreak nodes based on their Eu-
76 clidean distance to the fire's starting point, prioritizing the closest nodes to encircle the fire
77 rapidly.

78 **3.3 Specific scenario parameters and implications - scenario 1**

79 The simulation was configured with a unique, constrained scenario:

- 80 • **Grid size:** A 20x20 grid (400 nodes).
81 • **Fire start position:** The fire initiates at the center of the grid (10, 10), creating a symmetric
82 problem.
83 • **Number of firebreaks:** 25 firebreaks are allocated (6.25% of total nodes), a high density
84 for containment.
85 • **Wind factor:** A moderate `wind_factor` of 0.8.
86 • **Fuel load distribution:** Randomly assigned fuel loads between 0.1 and 1.0.

87 **3.4 Specific scenario parameters and implications - scenario 2**

88 After conducting the first scenario, the need for a second experiment was identified(see section 4).
89 Thus, an altered version of the first experiment was reexamined.

- 90 • **Grid size:** A 20x20 grid (400 nodes).
91 • **Fire start position:** **Fire starts at random position in the 20x20 grid.**
92 • **Number of firebreaks:** 25 firebreaks are allocated (6.25% of total nodes), a high density
93 for containment.
94 • **Wind factor:** A moderate `wind_factor` of 0.8.
95 • **Fuel load distribution:** Randomly assigned fuel loads between 0.1 and 1.0.

96 **3.5 Specific scenario parameters and implications - scenario 3**

97 After conducting the second scenario, the need for a third experiment was identified(see section 4).
98 Thus, an altered version of the second experiment was reexamined.

- 99 • **Grid size:** A 20x20 grid (400 nodes).
100 • **Fire start position:** **7 different fires each start at random positions in the 20x20 grid.**
101 • **Number of firebreaks:** 25 firebreaks are allocated (6.25% of total nodes), a high density
102 for containment.
103 • **Wind factor:** A moderate `wind_factor` of 0.8.
104 • **Fuel load distribution:** Randomly assigned fuel loads between 0.1 and 1.0.

105 **3.6 Control of stochastic error**

106 Since the ACO function creates stochastic error by the random functions, the experiment was
107 conducted 100 times, and mean data was extracted from experiments.

108 **4 Simulation results and analysis**

109 **4.1 Simulation results - scenario 1**

110 Execution of the simulation under the specified scenario reveals a notable difference in performance
111 between the ACO and Greedy strategies.

Table 1: Comparison of fire suppression strategies(100 executions)

Metric	ACO Strategy	Greedy Strategy
Total Area Burned (Average, nodes)	88.74	1.00
Total Area Burned (Standard Deviation, nodes)	147.90	0.00
Time to Containment (Average, steps)	10.97	1.00
Time to Containment (Standard Deviation, steps)	16.95	0.00

112 As detailed in Table 1, the Greedy strategy significantly outperformed the ACO strategy. The Greedy
 113 algorithm limited the total area burned to a single node and achieved containment within one time
 114 step. In contrast, the ACO strategy resulted in approximately 89 nodes burned and required 11 steps
 115 for containment, with a high margin of error. As later mentioned in section 5, the highly symmetric
 116 design of the map was suspected as the cause of error.

117 4.2 Simulation results - scenario 2

118 Execution of the simulation under the specified scenario reveals no difference in the heuristic Greedy
 algorithm, while the data for ACO showed some improvement.

Table 2: Comparison of fire suppression strategies(100 executions)

Metric	ACO Strategy	Greedy Strategy
Total Area Burned (Average, nodes)	45.60	1.00
Total Area Burned (Standard Deviation, nodes)	114.31	0.00
Time to Containment (Average, steps)	6.09	1.00
Time to Containment (Standard Deviation, steps)	12.70	0.00

119
 120 As detailed in Table 2, it is clearly visible that, although the ACO method did improve in efficiency, it
 121 is still far behind the efficiency of the simple Greedy method, which hasn't changed in value. Thus,
 122 it was clear that the heuristic algorithm was too optimized for a single-fire task. Since the Greedy
 123 algorithm can simply "surround" the fire source with 4 firebreaks, any one-fire case can be easily
 124 handled by the Greedy algorithm.

125 4.3 Simulation results - scenario 3

126 Execution of the simulation under the specified scenario reveals drastically different data, but no
 difference in trend; the Greedy algorithm outperforms the ACO algorithm.

Table 3: Comparison of fire suppression strategies(100 executions)

Metric	ACO Strategy	Greedy Strategy
Total Area Burned (Average, nodes)	340.50	147.46
Total Area Burned (Standard Deviation, nodes)	31.16	164.42
Time to Containment (Average, steps)	24.63	17.21
Time to Containment (Standard Deviation, steps)	6.49	18.83

127
 128 As detailed in Table 3, since the number of fire sources exceeds $6(\lfloor \frac{25}{4} \rfloor)$, we can know that fire
 129 sources must be connected by a side in order to be contained. This has caused a drastic change in
 130 data, leading to large error in the Greedy strategy. The ACO strategy can be seen taking care of the
 131 situation without much deviation from the mean, however takes too long to contain and loses a lot of
 132 land.

133 The Greedy strategy has preformed much worse compared to scenarios 1 and 2 in scenario 3. This is
 134 most likely due to the randomness of fire origins requiring a more global resource distribution.

135 **5 Explaining Greedy's superior performance**

136 The observed performance of the Greedy algorithm can be attributed to the fundamental dichotomy
137 between local and global optimization strategies, coupled with the unique characteristics of the
138 problem instance.

139 **5.1 Local vs. global optimization principles**

140 Optimization problems typically involve finding the global optimum among all feasible solutions.
141 Greedy algorithms inherently pursue local optima, making the best choice at each step based on
142 immediate information. This approach is computationally efficient but does not guarantee a globally
143 optimal solution. Metaheuristics, including ACO, are designed to traverse the solution space more
144 thoroughly, balancing exploration and exploitation to find global optima.

145 **5.2 Specific factors contributing to Greedy's success**

- 146 • **Centralized Fire Start and Symmetric Problem Structure:** The fire originating at the
147 center of the grid creates a highly symmetric problem where the most effective containment
148 strategy is to establish a perimeter in adjacent nodes. The Greedy algorithm's focus on
149 proximity aligns perfectly with this optimal initial move.
- 150 • **High Density of Firebreaks:** With 25 firebreaks on a 400-node grid, a significant proportion
151 of the area can be converted into containment lines. This abundance allows even a simple
152 proximity-based strategy to quickly form an effective perimeter.
- 153 • **Greedy Heuristic's Direct Relevance:** The Greedy heuristic of selecting nodes by Eu-
154 clidean distance was optimally aligned with the ideal strategy for this specific centralized
155 fire scenario.
- 156 • **ACO's Exploration Overhead:** ACO inherently incurs computational overhead for ex-
157 ploration and pheromone updating. In a scenario where the optimal solution is immediate
158 and localized, ACO's broader search may delay the concentration of resources on the most
159 critical nodes.

160 **5.3 Literature context for simpler heuristics outperforming metaheuristics**

161 The observed outcome is not an anomaly but is well-documented in the optimization literature.
162 Some research indicates that simpler local search heuristics can prove highly competitive or even
163 superior to more complex metaheuristics, especially when constraints limit the effective search space.
164 This phenomenon underscores the importance of aligning algorithm selection with the problem's
165 underlying structure. The simulation results confirm this understanding, illustrating that for problems
166 with inherent symmetry and localized optimal solutions, a Greedy strategy can indeed be more
167 effective.

168 **6 Limitations of the simulation and interpretation**

169 While the simulation provides valuable insights, it is important to acknowledge its inherent limitations:

- 170 • **Model simplifications:** The 20x20 grid is a significant abstraction of a real-world forest.
- 171 • **Stochastic nature:** The stochastic property of the simulation generates large error. Robust
172 conclusions will require more efficient algorithms.
- 173 • **Simplified objective function:** The ACO's objective function is a simplification of real-
174 world fire suppression costs.
- 175 • **Scenario specificity:** The most crucial limitation is the high degree of scenario specificity.
176 The results may not be generalizable to all forest fire scenarios where the optimal solution is
177 not readily apparent.

178 **7 Future work and broader implications**

179 The findings underscore several important directions for future research:

- 180 • **Extensive scenario diversity:** Future experiments must include a broader range of scenarios
181 (e.g., varying fire start locations, grid topologies) to test algorithms across a spectrum of
182 challenges.
- 183 • **Comprehensive parameter sensitivity analysis:** A systematic analysis of ACO's parame-
184 ters is essential for robust configurations.
- 185 • **Hybrid approaches:** Exploring hybrid algorithms that combine Greedy's speed for initial
186 containment with ACO's strategic optimization for evolving fires is a promising avenue
187 (Aranzazu-Suescun et al., 2014).
- 188 • **Integrating real-world complexity:** Future research should incorporate more sophisti-
189 cated models for fire spread, accounting for dynamic wind, heterogeneous fuel types, and
190 topography.
- 191 Despite its performance in this constrained scenario, ACO's potential for complex, large-scale forest
192 fire problems remains highly relevant. Its ability to identify non-obvious, globally optimal paths
193 offers a significant advantage over local methods in truly challenging environments.

194 **References**

- 195 Alexandridis, A., Russo, L., and Vakalis, D. (2011). Wildland fire spread modelling using cellular
196 automata: evolution in large-scale spatially heterogeneous environments under fire suppression
197 tactics. *International Journal of Wildland Fire*.
- 198 Aranzazu-Suescun, N. R., Abdulla, A. E. A. A., and Aguilar, A. A. A. (2014). Smartpath: An efficient
199 hybrid aco-ga algorithm for optimal path planning for autonomous guided vehicles. In *Recent
200 Advances in Robotics and Automation*. InTech.
- 201 Carta, F., Zidda, C., Putzu, M., Loru, D., Anedda, M., and Giusto, D. D. (2023). Advancements in
202 forest fire prevention: A comprehensive survey. *Sensors*, 23(14):6635.

203 **A Appendix : code and experiment reproducibility**

204 This appendix provides the necessary details to reproduce the experiments in the paper. Necessary
205 code is included in the supplementary material. The supplementary repository can be accessed at :
206 <https://github.com/codingneerd/forestfire>.

207 **A.1 System requirements and dependencies**

208 The simulation was executed on a standard personal computer and does not require specialized
209 hardware. It has, however, been tested on the following two devices :

- 210 • **Device:** Samsung Galaxy Book 5 Pro 360
211 • **CPU:** Intel(R) Core(TM) Ultra 7 256V
212 • **GPU:** Intel(R) Arc(TM) 140V GPU (8GB)
213 • **RAM:** 16GB
214 • **OS:** Windows 11 Home

215 • **Device:** Apple Macbook Air 15" (2025)
216 • **CPU:** Apple Silicon M4
217 • **GPU:** Apple Silicon M4
218 • **RAM:** 16GB
219 • **OS:** MacOS Sequoia

220 The script is written in Python v3.9.10 and relies on the following major libraries:

- 221 • **NumPy:** For numerical operations, particularly distance calculations.
222 • **NetworkX:** To create and manage the grid-based graph representing the forest.
223 • **Matplotlib:** For plotting the results.

224 To ensure compatibility, it is recommended to install the specific versions of these packages using the
225 following command:

226 `pip install numpy networkx matplotlib`

227 **A.2 Execution instructions**

228 To run the full simulation comparing the Ant Colony Optimization (ACO) and Greedy strategies,
229 each scenario file is provided under the naming convention

230 `scenario-n.py`

231 and execute it from your terminal with the following command:

232 `python scenario-n.py`

233 The script will print the final comparison of the total burned area and the time to containment for
234 both strategies to the console, along with the standard deviation of each value.

235 **A.3 Algorithm and simulation parameters**

236 All parameters used in the code are separated as variables on the top of the code. Altering these
237 values will result in different simulation results.

238 **Agents4Science AI Involvement Checklist**

- 239 1. **Hypothesis development:** Hypothesis development includes the process by which you
240 came to explore this research topic and research question. This can involve the background
241 research performed by either researchers or by AI. This can also involve whether the idea
242 was proposed by researchers or by AI.

243 Answer: **[B]**

244 Explanation: The topic related to 'Ant Colony Optimization Algorithm' was chosen by
245 humans. AI helped us brainstorm specific details and determine the problem this paper
246 attempted to solve, which was forest fires.

- 247 2. **Experimental design and implementation:** This category includes design of experiments
248 that are used to test the hypotheses, coding and implementation of computational methods,
249 and the execution of these experiments.

250 Answer: **[C]**

251 Explanation: To test our hypothesis, we wrote a python code to simulate the difference
252 between the ACO algorithm and the Greedy algorithm. Most of the code was designed and
253 generated by AI. However, there were some assistance from humans, as the initial code
254 included multiple errors that would result in false observations.

- 255 3. **Analysis of data and interpretation of results:** This category encompasses any process to
256 organize and process data for the experiments in the paper. It also includes interpretations of
257 the results of the study.

258 Answer: **[C]**

259 Explanation: The analysis of the simulation was mostly done by AI. Humans examined the
260 possible limitations of the ACO algorithm, making a significant contribution to the final
261 conclusion.

- 262 4. **Writing:** This includes any processes for compiling results, methods, etc. into the final
263 paper form. This can involve not only writing of the main text but also figure-making,
264 improving layout of the manuscript, and formulation of narrative.

265 Answer: **[C]**

266 Explanation: A large part of the text was generated by AI; humans fixed and rewrote some
267 phrases for a better explanation of the topic.

- 268 5. **Observed AI Limitations:** What limitations have you found when using AI as a partner or
269 lead author?

270 Description: The main difficulty was inspecting the paper for possible errors or citations
271 with low credibility. As efficient the writing process was, there were some issues that needed
272 to be addressed. We believe this requirement of a survey by humans is the current limitation
273 of using AI for research.

274 **Agents4Science Paper Checklist**

275 **1. Claims**

276 Question: Do the main claims made in the abstract and introduction accurately reflect the
277 paper's contributions and scope?

278 Answer: [Yes]

279 Justification: The abstract and the introduction accurately represent the main idea of the
280 paper.

281 Guidelines:

- 282 • The answer NA means that the abstract and introduction do not include the claims
283 made in the paper.
- 284 • The abstract and/or introduction should clearly state the claims made, including the
285 contributions made in the paper and important assumptions and limitations. A No or
286 NA answer to this question will not be perceived well by the reviewers.
- 287 • The claims made should match theoretical and experimental results, and reflect how
288 much the results can be expected to generalize to other settings.
- 289 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
290 are not attained by the paper.

291 **2. Limitations**

292 Question: Does the paper discuss the limitations of the work performed by the authors?

293 Answer: [Yes]

294 Justification: There is a section dedicated to the discussion of the limitations of our research.

295 Guidelines:

- 296 • The answer NA means that the paper has no limitation while the answer No means that
297 the paper has limitations, but those are not discussed in the paper.
- 298 • The authors are encouraged to create a separate "Limitations" section in their paper.
- 299 • The paper should point out any strong assumptions and how robust the results are to
300 violations of these assumptions (e.g., independence assumptions, noiseless settings,
301 model well-specification, asymptotic approximations only holding locally). The authors
302 should reflect on how these assumptions might be violated in practice and what the
303 implications would be.
- 304 • The authors should reflect on the scope of the claims made, e.g., if the approach was
305 only tested on a few datasets or with a few runs. In general, empirical results often
306 depend on implicit assumptions, which should be articulated.
- 307 • The authors should reflect on the factors that influence the performance of the approach.
308 For example, a facial recognition algorithm may perform poorly when image resolution
309 is low or images are taken in low lighting.
- 310 • The authors should discuss the computational efficiency of the proposed algorithms
311 and how they scale with dataset size.
- 312 • If applicable, the authors should discuss possible limitations of their approach to
313 address problems of privacy and fairness.
- 314 • While the authors might fear that complete honesty about limitations might be used by
315 reviewers as grounds for rejection, a worse outcome might be that reviewers discover
316 limitations that aren't acknowledged in the paper. Reviewers will be specifically
317 instructed to not penalize honesty concerning limitations.

318 **3. Theory assumptions and proofs**

319 Question: For each theoretical result, does the paper provide the full set of assumptions and
320 a complete (and correct) proof?

321 Answer: [No]

322 Justification: No direct mathematical proof was given, as the goal of this paper was the
323 simulation of the ACO algorithm compared to the Greedy algorithm.

324 Guidelines:

- 325 • The answer NA means that the paper does not include theoretical results.
326 • All the theorems, formulas, and proofs in the paper should be numbered and cross-
327 referenced.
328 • All assumptions should be clearly stated or referenced in the statement of any theorems.
329 • The proofs can either appear in the main paper or the supplemental material, but if
330 they appear in the supplemental material, the authors are encouraged to provide a short
331 proof sketch to provide intuition.

332 **4. Experimental result reproducibility**

333 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
334 perimental results of the paper to the extent that it affects the main claims and/or conclusions
335 of the paper (regardless of whether the code and data are provided or not)?

336 Answer: [Yes]

337 Justification: Please refer to the appendix for experiment reproducibility, and the necessary
338 resources.

339 Guidelines:

- 340 • The answer NA means that the paper does not include experiments.
341 • If the paper includes experiments, a No answer to this question will not be perceived
342 well by the reviewers: Making the paper reproducible is important.
343 • If the contribution is a dataset and/or model, the authors should describe the steps taken
344 to make their results reproducible or verifiable.
345 • We recognize that reproducibility may be tricky in some cases, in which case authors
346 are welcome to describe the particular way they provide for reproducibility. In the case
347 of closed-source models, it may be that access to the model is limited in some way
348 (e.g., to registered users), but it should be possible for other researchers to have some
349 path to reproducing or verifying the results.

350 **5. Open access to data and code**

351 Question: Does the paper provide open access to the data and code, with sufficient instruc-
352 tions to faithfully reproduce the main experimental results, as described in supplemental
353 material?

354 Answer: [Yes]

355 Justification: We provide the codes written for the simulation.

356 Guidelines:

- 357 • The answer NA means that paper does not include experiments requiring code.
358 • Please see the Agents4Science code and data submission guidelines on the conference
359 website for more details.
360 • While we encourage the release of code and data, we understand that this might not be
361 possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not
362 including code, unless this is central to the contribution (e.g., for a new open-source
363 benchmark).
364 • The instructions should contain the exact command and environment needed to run to
365 reproduce the results.
366 • At submission time, to preserve anonymity, the authors should release anonymized
367 versions (if applicable).

368 **6. Experimental setting/details**

369 Question: Does the paper specify all the training and test details (e.g., data splits, hyper-
370 parameters, how they were chosen, type of optimizer, etc.) necessary to understand the
371 results?

372 Answer: [Yes]

373 Justification: Please refer to the appendix for details of the simulation.

374 Guidelines:

- 375 • The answer NA means that the paper does not include experiments.

- 376 • The experimental setting should be presented in the core of the paper to a level of detail
377 that is necessary to appreciate the results and make sense of them.
378 • The full details can be provided either with the code, in appendix, or as supplemental
379 material.

380 **7. Experiment statistical significance**

381 Question: Does the paper report error bars suitably and correctly defined or other appropriate
382 information about the statistical significance of the experiments?

383 Answer: [No]

384 Justification: Our conclusions were mainly based on our simulation, rather than other
385 relevant information.

386 Guidelines:

- 387 • The answer NA means that the paper does not include experiments.
388 • The authors should answer "Yes" if the results are accompanied by error bars, confi-
389 dence intervals, or statistical significance tests, at least for the experiments that support
390 the main claims of the paper.
391 • The factors of variability that the error bars are capturing should be clearly stated
392 (for example, train/test split, initialization, or overall run with given experimental
393 conditions).

394 **8. Experiments compute resources**

395 Question: For each experiment, does the paper provide sufficient information on the com-
396 puter resources (type of compute workers, memory, time of execution) needed to reproduce
397 the experiments?

398 Answer: [No]

399 Justification: Please refer to the appendix for details on our compute resources.

400 Guidelines:

- 401 • The answer NA means that the paper does not include experiments.
402 • The paper should indicate the type of compute workers CPU or GPU, internal cluster,
403 or cloud provider, including relevant memory and storage.
404 • The paper should provide the amount of compute required for each of the individual
405 experimental runs as well as estimate the total compute.

406 **9. Code of ethics**

407 Question: Does the research conducted in the paper conform, in every respect, with the
408 Agents4Science Code of Ethics (see conference website)?

409 Answer: [Yes]

410 Justification: We have reviewed the Agents4Science Code of Ethics.

411 Guidelines:

- 412 • The answer NA means that the authors have not reviewed the Agents4Science Code of
413 Ethics.
414 • If the authors answer No, they should explain the special circumstances that require a
415 deviation from the Code of Ethics.

416 **10. Broader impacts**

417 Question: Does the paper discuss both potential positive societal impacts and negative
418 societal impacts of the work performed?

419 Answer: [Yes]

420 Justification: In the section 'Future Work and Broader Implications', we discuss the potential
421 effects of our research.

422 Guidelines:

- 423 • The answer NA means that there is no societal impact of the work performed.
424 • If the authors answer NA or No, they should explain why their work has no societal
425 impact or why the paper does not address societal impact.

- 426 • Examples of negative societal impacts include potential malicious or unintended uses
427 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations,
428 privacy considerations, and security considerations.
429 • If there are negative societal impacts, the authors could also discuss possible mitigation
430 strategies.