

---

# Research on CAPTCHAs Targeted at AI: Human–Easy, Model–Hard Visual Tasks for the AI Era

---

**Anonymous Author(s)**

Affiliation  
Address  
email

## Abstract

1 As generative artificial intelligence advances rapidly, traditional CAPTCHAs de-  
2 signed to distinguish humans from machines are losing efficacy. State-of-the-art  
3 deep learning systems now solve conventional challenges such as distorted text  
4 recognition with near-perfect accuracy. Motivated by this, we explore “AI-targeted  
5 CAPTCHAs”—challenge tasks that humans pass easily but multimodal models  
6 find difficult. Building on a review of prior work, we posit two cognitive path-  
7 ways that contemporary models rely on: a “linguistic path” versus a “perceptual  
8 path.” Guided by these hypotheses, we design five simple, highly intuitive visual  
9 question–answer tasks to systematically compare humans with leading multimodal  
10 models. Each task pairs a single image with a single question and covers color  
11 discrimination, size comparison, combined distractors, and counting legs on birds  
12 or fingers on human hands. We evaluate five mainstream multimodal systems  
13 under the same conditions as human participants and test two main hypotheses  
14 plus one sub-hypothesis. Results show: (1) on single-feature tasks such as color  
15 discrimination, top models approach human-level performance; however, for size  
16 comparison and combined-difference tasks that require low-level visual perception,  
17 model accuracy collapses to near zero while humans perform almost perfectly; (2)  
18 in bird-leg and hand-finger counting tasks, models frequently default to stereotyped  
19 prior knowledge, achieving 20% accuracy or lower, whereas humans rely on the  
20 image and score near 100%; (3) models recognize abnormal human fingers slightly  
21 better than abnormal bird legs, supporting sub-hypothesis H1.1 that models handle  
22 common human limbs better than non-typical species. These differences confirm  
23 that current vision–language models primarily follow a linguistic path for image  
24 understanding and lack a human-like low-level perceptual path for processing  
25 obvious visual information. The findings quantify the limits of current multimodal  
26 systems and demonstrate the feasibility of constructing new CAPTCHAs from  
27 such “AI-hard” tasks.

28 **Keywords:** AI-targeted CAPTCHA; vision–language models (VLMs); visual perception; common-  
29 sense/confirmation bias; human–AI performance gap; multimodal evaluation.

30 

## 1 Introduction

31 Artificial intelligence has made remarkable gains in visual recognition and multimodal understanding.  
32 Deep neural networks now rival or surpass humans on standard benchmarks [LeCun et al., 2015,  
33 Russakovsky et al., 2015]. At the same time, the long-standing need to distinguish humans from  
34 machines in security-critical settings has not diminished. Classic CAPTCHAs—particularly text-  
35 distortion variants—were originally predicated on tasks that were considered easy for humans but  
36 hard for machines. Yet reCAPTCHA-style schemes have steadily lost security margin as machine

37 vision improved [von Ahn et al., 2008]. This raises a central question for the multimodal era: Which  
38 challenge types can sustainably remain human-easy but model-hard?

39 We study that question by proposing an AI-targeted CAPTCHA design space grounded in a two-  
40 pathway account of image question answering. Contemporary vision–language systems often appear  
41 to follow a primarily *linguistic path*: rather than adducing visual evidence, they default to high-  
42 frequency language patterns and commonsense templates learned during pretraining—e.g., assuming  
43 that a bird has two legs or a hand has five fingers—even when the image shows otherwise [Nickerson,  
44 1998, Geirhos et al., 2020, Schulze-Buschhoff et al., 2024]. Humans, by contrast, typically rely first  
45 on a *perceptual path*: fast, low-level mechanisms encode basic features such as color, size, location,  
46 and numerosity in parallel before higher-level inference engages [Treisman and Gelade, 1980, Itti  
47 and Koch, 2001, Rosenholtz, 2014, Dehaene, 2011, Marr, 1982, Wolfe and Horowitz, 2017]. This  
48 asymmetry resonates with Moravec’s paradox, whereby abstract reasoning can be comparatively easy  
49 for machines while low-level perception remains difficult [Moravec, 1988].

50 Guided by this framework, we articulate two main hypotheses and one sub-hypothesis. **H1 (linguistic-  
51 bias)**: When a task requires direct inspection of low-level visual evidence, current models will perform  
52 poorly because they over-weight linguistic priors. **H2 (perceptual-deficit)**: Multimodal systems  
53 lack a robust human-like perceptual pathway and will systematically err when image evidence  
54 conflicts with commonsense. **H1.1 (sub-hypothesis)**: On limb-counting tasks, models will do slightly  
55 better with human hands than with birds, reflecting training exposure and tuning to human imagery  
56 [Schulze-Buschhoff et al., 2024].

## 57 2 Related Work

### 58 2.1 CAPTCHAs and the human-machine gap

59 Early CAPTCHAs exploited gaps between human and machine vision to authenticate users at scale.  
60 reCAPTCHA famously coupled security with human-based character recognition, leveraging human  
61 effort to transcribe ambiguous text [von Ahn et al., 2008]. The advent of modern deep learning,  
62 however, has eroded the hardness assumptions of text-based CAPTCHAs [LeCun et al., 2015,  
63 Russakovsky et al., 2015], motivating new challenge families that do not hinge on distorted text.

### 64 2.2 Low-level perception and attention in human vision

65 A large literature shows that humans can extract basic visual features (e.g., color, size, orientation,  
66 location, and numerosity) rapidly and in parallel, often with little or no attentional load [Treisman  
67 and Gelade, 1980, Itti and Koch, 2001, Wolfe and Horowitz, 2017, Rosenholtz, 2014, Dehaene, 2011,  
68 Marr, 1982]. Bottom-up salience and top-down task goals jointly guide selection [Desimone and  
69 Duncan, 1995], and even in clutter, peripheral summary statistics support efficient search for obvious  
70 outliers [Rosenholtz, 2014]. These properties predict that tasks grounded in low-level evidence will  
71 remain trivial for healthy adults.

### 72 2.3 Shortcut learning and linguistic bias in modern models

73 Despite strong benchmark performance, deep networks often exploit statistical shortcuts that correlate  
74 with labels without supporting genuine perception or reasoning [Geirhos et al., 2020]. In multimodal  
75 systems, the language component can dominate, yielding answers that align with high-frequency  
76 commonsense rather than the specific image—a hallmark of confirmation bias [Nickerson, 1998].  
77 Recent evaluations document persistent gaps in intuitive visual cognition and perception across  
78 state-of-the-art models [Schulze-Buschhoff et al., 2024].

### 79 2.4 Security applications leveraging human–AI asymmetries

80 A complementary literature in machine learning security studies adversarial failures of vision models  
81 [e.g., Akhtar and Mian, 2018, Goodfellow et al., 2015]. While many such perturbations are not  
82 user-friendly for authentication, the broader lesson is salient: today’s systems exhibit systematic  
83 perceptual weaknesses. Our proposal diverges by enforcing a human-ease constraint and by favoring  
84 visible, low-level cues and light commonsense conflict as the lever for separating humans from  
85 machines.

86 **2.5 Summary and positioning**

87 The decline of text-based CAPTCHAs, the psychology of pre-attentive perception, and shortcut learning  
88 in modern networks converge on the same design principle: build challenges that require direct,  
89 low-level inspection of the image, particularly in scenarios that conflict mildly with commonsense.  
90 The present study operationalizes this principle through five minimal task families and human–model  
91 comparisons, providing an immediately usable recipe for AI-targeted CAPTCHAs.

92 **3 Research Hypotheses**

- 93 • **H1 (Linguistic-bias hypothesis).** Multimodal models primarily rely on linguistic patterns  
94 and memorized commonsense when answering visual questions; thus they will perform  
95 poorly on tasks that require direct, low-level perceptual processing, while potentially appear-  
96 ing competitive on tasks with clear semantic cues.
- 97 • **H2 (Perceptual-deficit hypothesis).** Current models lack a robust human-like percep-  
98 tual pathway; when visual evidence conflicts with their prior commonsense, they display  
99 systematic biases, often producing answers consistent with memory rather than the image.
- 100 • **H1.1 (Sub-hypothesis).** On tasks invoking commonsense about familiar human anatomy  
101 versus non-human species, models will perform slightly better on human hands than on  
102 birds, reflecting greater exposure and potential tuning toward human imagery in training  
103 [Schulze-Buschoff et al., 2024].

104 **4 Methods**

105 **4.1 Task Design**

106 We created five visual question–answer task families, each presented as a single image paired with a  
107 single question. Participants answered based solely on the provided image. The task families are: (A) size  
108 difference (odd-one-out: one item slightly larger), (B) color difference (odd-one-out: one item  
109 subtly different color; RGB distance  $\approx 45$ ), (C) combined distractors (one size-different item and one  
110 color-different item), (D) counterfactual animals (counting visible bird legs), and (E) counterfactual  
111 human hands (counting visible fingers).

112 **4.2 Participants and Models**

113 Five adult volunteers completed 50 items each (10 per task family) with no time limit and a single  
114 allowed response per item. We evaluated five multimodal models under identical conditions: OpenAI  
115 GPT-5, Anthropic Claude-Sonnet-4 (20250514), Google Gemini 2.5-Pro, ByteDance Doubao-seed-  
116 1.6 (250615), and Alibaba Qwen 2.5-VL-72B-Instruct. All prompts were in Chinese to avoid language  
117 confounds, and models returned constrained-format answers (e.g., “r3c4”, or an integer count).

118 **4.3 Stimuli**

119 For A–C we generated  $5 \times 5$  arrays with a program that guaranteed a single valid odd-one-out per  
120 image and auto-recorded keys. For D–E we curated AI-generated photos depicting controlled limb  
121 abnormalities while avoiding privacy concerns. All materials adhered to ethical guidelines.

122 **4.4 Procedure**

123 Humans solved all 50 items once each. For models, we ran up to 10 independent answer attempts  
124 per item to gauge response stability. We computed average accuracy per task family for humans  
125 and for each model, and we ran Fisher’s exact tests and  $\chi^2$  tests to assess significance of human–AI  
126 differences.

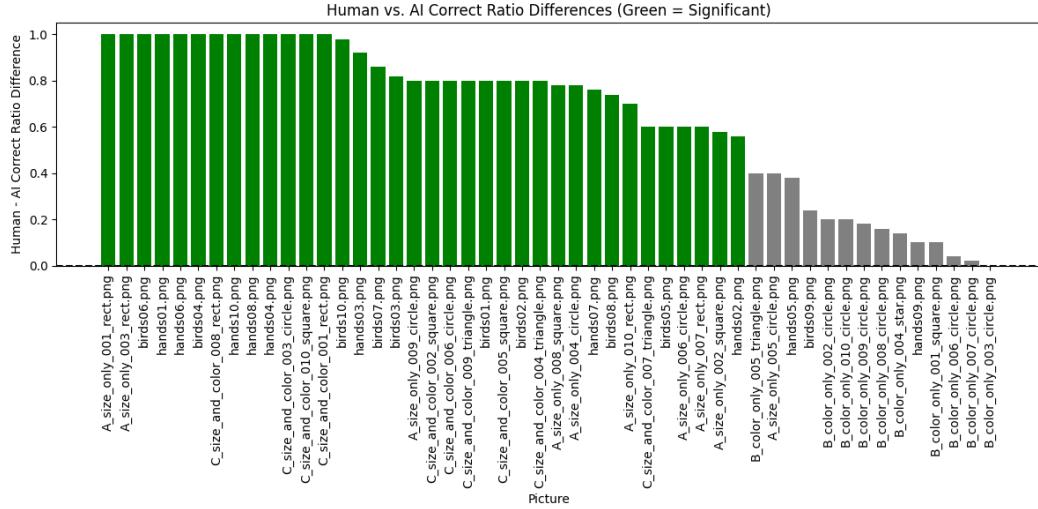


Figure 1: Human vs. AI correct-ratio differences (green bars are significant).

Table 1: Human vs. AI accuracy by task family.

Task family	Human	AI	Gap (points)
Odd-one-out (size)	78%	6%	+72
Odd-one-out (color)	100%	86%	+14
Combined (size+color)	86%	0%	+86
Bird-leg counting	94%	16%	+78
Hand-finger counting	98%	21%	+77

## 127 4.5 Data Capture and Analysis

128 Data capture and analysis were scripted with LLM assistance (for API calls, answer parsing, and  
129 scoring), then manually verified by the authors. Statistical summaries included per-task accuracies,  
130 overall human–model gaps, and model-wise totals.

## 131 5 Results

132 **Overall performance.** Figure 1 summarizes human minus AI accuracy per item (green markers:  
133  $p < 0.05$ ). Humans were near-ceiling on most items (task-family averages: A 78%, B 100%, C 86%,  
134 D 94%, E 98%). Models varied sharply by task type.

135 **(1) Color discrimination (B).** Models approached humans: five models averaged  $\approx 85.6\%$  accuracy.

136 **(2) Size differences (A).** Performance collapsed: models averaged  $\approx 5.6\%$  accuracy; humans 78%.

137 **(3) Combined distractors (C).** Models failed completely: across ten images, aggregate accuracy  
138 of all five models was 0%. Answers typically noticed the color outlier but ignored size. Humans  
139 achieved 86%.

140 **Commonsense-conflict tasks.** For bird-leg counting (D), models averaged  $\approx 15.6\%$ ; for hand-finger  
141 counting (E),  $\approx 20.8\%$ . Models often defaulted to “two legs” or “five fingers” regardless of evidence.  
142 Humans were  $\sim 94\text{--}98\%$ .

## 143 6 Discussion and Conclusions

144 The results support both H1 and H2. Models that excel at semantically cued color differences  
145 nonetheless fail on purely perceptual geometry (size) and on multi-feature integration (combined

Table 2: Overall accuracy across 50 items by model.

Model (VLM)	Overall accuracy
GPT-5	33.4%
Claude-Sonnet-4-20250514	24.8%
Gemini 2.5-Pro	31.4%
Doubao-seed-1.6-250615	20.4%
Qwen2.5-VL-72B-Instruct	17.6%

146 distractors), indicating an absent or weak low-level perceptual path. When visual evidence conflicts  
 147 with prior knowledge, models strongly favor linguistic priors, echoing confirmation bias [Nickerson,  
 148 1998]. Accuracy is slightly higher on human hands than on birds (H1.1), plausibly reflecting training-  
 149 data prevalence and targeted tuning for hands in modern generative pipelines. These findings align  
 150 with evaluations reporting large human–model gaps on intuitive visual cognition [Schulze-Buschoff  
 151 et al., 2024] and with visual-illusion CAPTCHAs that reliably trip models while remaining human-  
 152 easy [Ding et al., 2025].

## 153 7 Implications for AI-Targeted CAPTCHAs

154 The task families here are simple for humans yet reliably difficult for today’s models. A practical  
 155 CAPTCHA could ask users to click the color outlier in a grid, identify the larger item, or count visible  
 156 limbs in a photo. Curating a diverse bank of such items would create an effective human–AI separator  
 157 in the near term. As models improve, the bank will require updates, but the present gap is substantive  
 158 enough to be useful for security and for probing multimodal cognition.

## 159 8 Limitations and Future Work

160 Tasks and sample size are limited (five families, ten items each). Failures here do not imply universal  
 161 deficits on all “intuitive” problems; prompt engineering, auxiliary perception modules, or multi-turn  
 162 clarification might improve results. Future work should broaden task coverage (e.g., dynamic videos,  
 163 3D perception, cluttered scenes) and probe pathway usage more directly by manipulating prompts to  
 164 emphasize visual evidence versus prior knowledge.

## 165 9 Ethical Statement

166 All images were AI-generated or drawn from public resources with no personal data and no animal  
 167 experimentation. Five adult volunteers provided informed consent. LLM assistance was used for  
 168 stimulus generation, data processing, and drafting under human supervision; final analyses and claims  
 169 remain the authors’ responsibility.

## 170 10 Reproducibility Statement

171 We provide details to enable independent replication: scripts for generating  $5 \times 5$  grids (tasks A–  
 172 C), prompt templates (Chinese), answer-parsing code, and scoring utilities, along with seeds and  
 173 parameter ranges for color distances and size deltas. For tasks D–E, we release procedures and  
 174 prompts that reproduce comparable AI-generated images while avoiding redistribution of any single  
 175 long-lived CAPTCHA item. Each participant/model solved 50 items (10 per family); models were  
 176 queried up to 10 attempts per item to assess stability. We report versions, prompts, decoding  
 177 parameters, and significance tests (Fisher’s exact,  $\chi^2$ ). API-only inference; no training.

## 178 References

- 179 Naveed Akhtar and Ajmal Mian. Threat of adversarial attacks on deep learning in computer vision:  
 180 A survey. *IEEE Access*, 6:14410–14430, 2018.

- 181 Stanislas Dehaene. *The Number Sense: How the Mind Creates Mathematics*. Oxford University  
182 Press, revised and expanded edition, 2011.
- 183 Robert Desimone and John Duncan. Neural mechanisms of selective visual attention. *Annual Review*  
184 *of Neuroscience*, 18:193–222, 1995.
- 185 Zhe Ding, Yan Wang, Xinyu Chen, Hui Liu, and Ming Zhang. Illusioncaptcha: A captcha based on  
186 visual illusion. In *Proceedings of the ACM Web Conference 2025*, pages 1–12. ACM, 2025.
- 187 Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias  
188 Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine*  
189 *Intelligence*, 2(11):665–673, 2020.
- 190 Ian J. Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial  
191 examples. In *ICLR*, 2015.
- 192 Laurent Itti and Christof Koch. Computational modelling of visual attention. *Nature Reviews*  
193 *Neuroscience*, 2(3):194–203, 2001.
- 194 Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *Nature*, 521(7553):436–444,  
195 2015.
- 196 David Marr. *Vision: A computational investigation into the human representation and processing of*  
197 *visual information*. W. H. Freeman, 1982.
- 198 Hans P. Moravec. *Mind Children: The Future of Robot and Human Intelligence*. Harvard University  
199 Press, 1988.
- 200 Raymond S. Nickerson. Confirmation bias: A ubiquitous phenomenon in many guises. *Review of*  
201 *General Psychology*, 2(2):175–220, 1998.
- 202 Ruth Rosenholtz. What your visual system sees where you are not looking. *Interface Focus*, 4(3):  
203 2014007, 2014.
- 204 Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang,  
205 Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg, and Li Fei-Fei. Imagenet  
206 large scale visual recognition challenge. *International Journal of Computer Vision*, 115(3):211–252,  
207 2015.
- 208 Lea Marie Schulze-Buschoff, Zeynep Akata, Matthias Bethge, and Eric Schulz. Visual cognition in  
209 multimodal large language models. *Nature Machine Intelligence*, 6(12):1228–1239, 2024.
- 210 Anne Treisman and Garry Gelade. A feature-integration theory of attention. *Cognitive Psychology*,  
211 12(1):97–136, 1980.
- 212 Luis von Ahn, Benjamin Maurer, Colin McMillen, David Abraham, and Manuel Blum. recaptcha:  
213 Human-based character recognition via web security measures. *Science*, 321(5895):1465–1468,  
214 2008.
- 215 Jeremy M. Wolfe and Todd S. Horowitz. Five factors that guide attention in visual search. *Nature*  
216 *Human Behaviour*, 1(3):0058, 2017.

217 **A Technical Appendices and Supplementary Material**

- 218 Technical appendices with additional results, figures, graphs and proofs may be submitted with the  
219 paper submission before the full submission deadline, or as a separate PDF in the ZIP file before the  
220 supplementary material deadline. There is no page limit for the technical appendices.

221 **Agents4Science AI Involvement Checklist**

222     1. **Hypothesis development**

223       Answer: **[B]**

224       Explanation: Humans proposed the research idea and two-pathway framework (linguistic vs.  
225       perceptual) and formulated H1/H2/H1.1. AI helped with phrasing and literature reminders  
226       but did not originate the framing.

227     2. **Experimental design and implementation**

228       Answer: **[B]**

229       Explanation: Humans designed five task families, evaluation protocol, and scoring rules. AI  
230       assisted as a coding helper for grid generation/prompts under human specification; outputs  
231       were checked by humans.

232     3. **Analysis of data and interpretation of results**

233       Answer: **[B]**

234       Explanation: Humans ran Fisher's exact and  $\chi^2$  tests, verified the outputs, and drew conclu-  
235       sions. AI provided small parsing/plotting snippets that were reviewed and edited.

236     4. **Writing**

237       Answer: **[B]**

238       Explanation: The manuscript, tables, and figure captions were drafted and finalized by  
239       humans; AI was used for line-editing. Stimulus images for counting tasks were AI-generated  
240       under human prompts and curation.

241     5. **Observed AI Limitations**

242       Description: AI editors are unreliable for precise statistical choices and consistent parsing;  
243       multimodal models defaulted to commonsense priors (e.g., "five fingers") over image  
244       evidence, mirroring our findings.

245 **Agents4Science Paper Checklist**

246 **1. Claims**

247 Question: Do the main claims made in the abstract and introduction accurately reflect the  
248 paper's contributions and scope?

249 Answer: [Yes]

250 Justification: Claims match the demonstrated gaps between humans and VLMs and clearly  
251 state scope (five task families, offline evaluation).

252 **2. Limitations**

253 Question: Does the paper discuss the limitations of the work performed by the authors?

254 Answer: [Yes]

255 Justification: We discuss limited task coverage/sample size and generalization caveats in the  
256 "Limitations and Future Work" section.

257 **3. Theory assumptions and proofs**

258 Question: For each theoretical result, does the paper provide the full set of assumptions and  
259 a complete (and correct) proof?

260 Answer: [NA]

261 Justification: The paper is empirical and does not include formal theorems.

262 **4. Experimental result reproducibility**

263 Question: Does the paper fully disclose information needed to reproduce the main results?

264 Answer: [Yes]

265 Justification: Tasks, prompts, attempts, scoring, and statistical tests are specified; an artifact  
266 plan is described in the Reproducibility Statement.

267 **5. Open access to data and code**

268 Question: Does the paper provide open access to data/code?

269 Answer: [No]

270 Justification: To avoid seeding long-lived CAPTCHA answers, we will release generators/prompts/code upon acceptance; keys are withheld at submission.

272 **6. Experimental setting/details**

273 Question: Are training/test details (or inference protocols) specified?

274 Answer: [Yes]

275 Justification: We report model versions, attempts per item, constrained answer formats, and  
276 evaluation protocol.

277 **7. Experiment statistical significance**

278 Question: Are error bars or significance tests reported?

279 Answer: [Yes]

280 Justification: We use Fisher's exact and  $\chi^2$  tests and mark significant items.

281 **8. Experiments compute resources**

282 Question: Is compute information provided?

283 Answer: [Yes]

284 Justification: Inference-only via public APIs; versions/attempt counts allow compute estimation.

286 **9. Code of ethics**

287 Question: Does the research conform to the Agents4Science Code of Ethics?

288 Answer: [Yes]

289 Justification: Non-identifiable AI images; adult consent; security/fairness considerations  
290 documented.

291      **10. Broader impacts**

292      Question: Does the paper discuss positive and negative societal impacts?

293      Answer: [Yes]

294      Justification: We discuss security utility vs. accessibility/privacy risks and mitigations.