
Synthetic Realism in Finance: A High-Fidelity Market Simulation with Adaptive LLM Agents

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agent-based modeling offers a powerful framework for understanding complex
2 systems, from social interactions to ecological networks. In finance, such models
3 are crucial for stress testing, risk assessment, and the development of intelligent
4 trading strategies. However, the fidelity of these simulations is often limited by the
5 simplicity of their agents, which typically operate on fixed, rule-based heuristics.
6 Here we present a novel, high-fidelity financial market simulation environment
7 that integrates Large Language Models (LLMs) to create adaptive, reasoning-
8 based agents. We augment the well-established ABIDES simulator with a novel
9 Brain-Memory Architecture. This features a core LLM “Brain” for complex,
10 real-time decision-making, supported by a persistent memory module that records
11 the outcomes of its past actions. This enables a powerful online learning loop,
12 allowing the agent to adapt its strategy based on immediate feedback within a
13 single simulation run. By ingesting real-world NASDAQ data at the message level
14 and introducing a calibrated population of background traders, we demonstrate
15 that our LLM-enhanced system accurately reproduces key market microstructure
16 phenomena, including price discovery and trade intensity. This work establishes a
17 new paradigm for synthetic financial environments, providing an unprecedented
18 platform to study the emergent dynamics of human-like intelligence in complex,
19 competitive systems and enabling the development of next-generation, explainable
20 AI agents for financial applications.

21

1 Introduction

22 The modern financial landscape is a complex adaptive system, characterized by intricate, non-linear
23 dynamics that are difficult to model with traditional econometric methods [3, 11]. The need for high-
24 fidelity synthetic data has become paramount for developing robust trading strategies, stress-testing
25 financial systems, and training sophisticated artificial intelligence (AI) agents [19, 10]. While real-
26 world market data is a cornerstone of financial research, its inherent privacy concerns, inaccessibility,
27 and often incomplete nature limit its utility for large-scale, controlled experimentation [19].

28 Synthetic data, artificially generated to mimic the statistical properties and patterns of real-world data,
29 is emerging as a powerful solution to these challenges. In a highly regulated industry like finance, it
30 provides a means to address critical issues of privacy, fairness, and explainability. It allows financial
31 institutions to train and validate AI models without compromising sensitive information, to create
32 balanced datasets that mitigate algorithmic bias, and to model rare but high-impact events that are
33 underrepresented in historical records.

34 Agent-based simulations [5, 2, 6, 15, 14, 17, 20], such as the ABIDES platform, provide a controlled
35 environment for studying market behavior via exchange-style messaging, configurable latency, and
36 limit-order-book mechanics. However, a fundamental limitation remains: most simulations rely

37 on pre-defined, static strategies (e.g., Zero-Intelligence or Momentum) [9, 12] that cannot process
38 unstructured information, reason about sentiment, or adapt online to shifting regimes.

39 **Our approach.** We present a high-fidelity market simulator that integrates Large Language Models
40 (LLMs) as adaptive, reasoning agents within an enhanced ABIDES environment. Concretely, we
41 augment ABIDES with (i) a *news-aware LLM layer* that converts structured news events into
42 actionable signals for trading and market-making agents, and (ii) a *Brain–Memory architecture* that
43 supports within-episode adaptation by retrieving prior experiences and writing back outcomes to
44 persistent memory. We replay real NASDAQ TotalView-ITCH data (via LOBSTER) at the message
45 level, calibrate a heterogeneous population of background agents, and validate microstructure realism
46 against historical benchmarks.

47 At a high level, the EnhancedLLMNewsAnalyzer normalizes LLM analyses of structured news
48 into sentiment, confidence, and risk annotations that directly drive two ABIDES-native agents:
49 ABIDESLLMTradingAgent (direction, strength, duration, and size under position caps) and
50 ABIDESLLMMarketMaker (spread and inventory control that widen/recenter under uncertainty).
51 This coupling yields interpretable dynamics where sentiment drives directional pressure and uncer-
52 tainty modulates liquidity.

53 The *Brain–Memory* design separates strategic reasoning from persistent learning: the LLM “Brain”
54 acts on a real-time state vector (market + news) while a file-backed *Memory* retrieves similar past
55 contexts and logs outcomes (parameters, actions, P&L). The loop operates online within a single
56 session: decisions at hour one are evaluated, written back to memory, and inform strategy at hour
57 two.

58 To ensure *message-level fidelity*, our pipeline parses and replays ITCH events and calibrates agent
59 populations so simulated order flow, volume, and volatility align with historical NASDAQ data.
60 For external validation, we align simulated executions to a per-second ITCH clock and compute
61 price-error in basis points, alongside visual overlays of mid-price and trades.

62 **Contributions.** This work makes the following contributions:

- 63 **1. LLM agents with online adaptation.** A Brain–Memory architecture that retrieves prior
64 experiences and updates memory with realized outcomes, enabling within-episode strategy
65 adaptation.
- 66 **2. News-aware intelligence that modulates liquidity and risk.** A structured news interface
67 whose outputs drive an LLM trading agent and an LLM market maker, coupling sentiment
68 to directional pressure and uncertainty to resiliency.
- 69 **3. Message-level realism and calibration.** A data pipeline that ingests NASDAQ ITCH (via
70 LOBSTER), calibrates heterogeneous background agents, and validates against historical
71 microstructure using execution-level metrics.
- 72 **4. Experimental evidence under matched conditions.** On AMZN (June 21, 2012), we
73 compare LLM-ON, LLM-OFF, and Baseline populations under identical seeds and configu-
74 rations; LLM-ON perceives injected news and reproduces key microstructure phenomena.

75 Together, these elements bridge the gap between rule-based simulators and adaptive, explainable
76 agents, yielding a platform for studying emergent behavior in realistic, event-driven markets and for
77 generating synthetic data that captures microstructure-level nuance.

78 2 Related Work

79 2.1 High-Fidelity Agent-Based Market Simulation

80 Agent-based market simulators enable controlled studies of trading strategies and market dynamics.
81 A prominent example is ABIDES (Agent-Based Interactive Discrete Event Simulation) [5], an open-
82 source framework designed for high-fidelity equity market simulation. ABIDES [5] supports tens of
83 thousands of trading agents interacting with an exchange agent and enforces realistic “market physics”,
84 e.g. nanosecond timestamp resolution, configurable network latencies, and standardized exchange
85 messaging protocols. These features allow researchers to replicate continuous double auctions and
86 even replay specific historical trading days with fine-grained control. Built on this foundation, recent
87 works have extended ABIDES for new research purposes. ABIDES-Gym [2] provides an OpenAI




Figure 1: End-to-end pipeline of the LLM-Enhanced ABIDES architecture. Real NASDAQ TotalView-ITCH data (via LOBSTER) are parsed and replayed to drive and calibrate the enhanced ABIDES core (kernel & messaging; limit-order-book and price-time matching; heterogeneous baseline agents). A news-aware LLM layer (EnhancedLLMNewsAnalyzer) converts structured news into sentiment/confidence signals that condition ABIDES agents (ABIDESLLMTradingAgent and ABIDESLLMMarketMaker) operating under exchange constraints. A Brain-Memory module enables within-episode adaptation by retrieving past experiences and writing back outcomes (*news, params, action, P&L*) after evaluation. The simulator emits OrderBookSnapshot/Trade streams which are aligned with ITCH to compute microstructure metrics (e.g., `price_error_bps`) and to render overlay visuals.

88 Gym interface to the discrete-event simulator, making it easier to train Reinforcement Learning (RL)
 89 agents [21] in a multi-agent market environment. For instance, Amrouni et al.[2] wrap the ABIDES-
 90 Markets simulator as Gym environments to benchmark daily investor and order-execution tasks,
 91 enabling RL agents to interact with a realistic limit order book (LOB) market as their training ground.
 92 Likewise, ABIDES-Economist [6] adapts the ABIDES platform to macroeconomics, simulating an
 93 entire economy of heterogeneous households, firms, a central bank, and a government. This platform
 94 introduce learning-capable agents into this economic simulator, allowing agents to either follow
 95 rule-based policies or learn optimal behaviors via interactions in an OpenAI Gym style environment.
 96 This extension showcases the versatility of ABIDES beyond financial markets, as it can model
 97 complex economic systems with agent learning and exogenous shocks for AI-economics research.

98 2.2 Multi-Agent Reinforcement Learning in Market Simulations

99 The integration of reinforcement learning [21] with high-fidelity market simulators has been an active
 100 research area [17, 20, 15, 14]. Karpe et al. [13] demonstrated a multi-agent RL approach to optimal
 101 trade execution using an ABIDES-based LOB simulator. They configured a historical order book
 102 scenario in ABIDES and trained a trading agent with Double Deep Q-Learning [16] to execute a large
 103 order optimally. Notably, the learned RL agent in some cases converged to a classic Time-Weighted
 104 Average Price (TWAP) strategy, a sensible baseline, indicating that the simulator’s realism can foster
 105 plausibly optimal behaviors. The performance of the RL policy was evaluated against real-market
 106 data via market replay, showing the importance of a high-fidelity environment for trustworthy policy
 107 testing. This study illustrate how combining ABIDES’ realistic market mechanics with multi-agent
 108 RL techniques can yield insights into algorithmic trading strategies under near-real conditions.

109 2.3 Simulation Realism and Transferability

110 A persistent challenge in market simulation research is ensuring that agent behaviors learned in
 111 simulation generalize to real markets. Vyettrenko et al. [24] highlight that many simulators [20, 24,
 112 17, 4] “fail to reproduce statistics and stylized facts seen in real markets,” undermining the robustness
 113 of strategies validated purely in silico. Their work, “Get Real,” [24] catalogs a comprehensive set of
 114 realism metrics from distributional properties of returns and order flows to higher-order stylized facts
 115 to quantitatively compare simulated and real LOB data. The results revealed significant discrepancies
 116 between even sophisticated simulations and actual market data, providing a benchmark for improving
 117 fidelity. Along similar lines, Pino et al. [18] propose similarity metrics for transfer learning [23] in
 118 trading agents to bridge the “sim-to-real” gap [22]. They evaluate conceptual similarity (comparing
 119 stylized statistical features of price series) [24, 7], structural similarity (comparing agents’ experience
 120 trajectories) [1], and performance similarity between different market MDPs [8]. These metrics
 121 guide when a policy learned in one market (or a simulator) can be safely transferred to another. By
 122 employing techniques like Probabilistic Policy Reuse [8] based on such similarity measures, one can
 123 adapt or reuse trading policies from a simulator in the real market context more effectively. This line

124 of work underlines the importance of quantifying fidelity and alignment between simulated and real
125 environments as a prerequisite for reliable AI trading research.

126 **3 LLM-Enhanced ABIDES Architecture**

127 Our system is built on an extended version of the ABIDES simulator, tailored to support and integrate
128 LLM agents. The core innovation lies in a modular architecture that separates the roles of decision-
129 making and memory, leveraging the strengths of different LLMs.

130 **3.1 The Limit Order Book and Matching Engine:**

131 At the core of our system is a clean and robust LOB engine, implemented in `order_book.py`. It
132 maintains two distinct sides of the book—bids and asks—as price-indexed FIFO queues and uses a
133 global order map for lifecycle tracking. The engine employs a strict price-time-priority matching
134 algorithm: market orders consume liquidity from the top of the book, while marketable limit orders
135 behave similarly until their limit is reached, with any leftover quantity resting as displayed liquidity.
136 Resting orders preserve their FIFO timestamp at each price level. Every match produces a Trade
137 record with detailed information, including timestamp, symbol, price, quantity, aggressor side, and
138 participating order/agent IDs. After each trade, the engine updates key market metrics, including best
139 bid/ask, spread, mid-price (the average of the best bid and ask), and last trade price. It also emits
140 periodic OrderBookSnapshot rows, capturing the state of the book and key metrics for subsequent
141 analysis.

142 **3.2 ABIDES Agent Population**

143 A Heterogeneous Environment: To create a realistic, multi-agent environment, our system utilizes a
144 diverse population of pre-existing ABIDES agent types as background traders. These agents, which
145 have been validated in prior research, provide a heterogeneous and dynamic market context for our
146 LLM agents. The key agent types include:

- 147 • **Zero-Intelligence (ZI) Agents:** These agents, governed by simple stochastic rules,
148 provide a fundamental layer of random, unpredictable behavior, mimicking market "noise."
- 149 • **Fundamental Value Agents:** These traders are informed by real or simulated historical
150 price data, placing orders based on their belief in a stock's intrinsic value.
- 151 • **Market Study Agents:** This category includes agents designed for specific research tasks,
152 such as those that execute pre-defined algorithms (e.g., TWAP or VWAP) or those that
153 model market impact to test the effects of large trades.
- 154 • **Market Makers:** These sophisticated agents contribute to market liquidity by simultane-
155 ously quoting bid and ask prices, dynamically adjusting their spreads to manage inventory
156 and risk.

157 These diverse, rule-based agents form the a robust baseline, allowing us to isolate and study the
158 unique contributions of our LLM-enhanced agents within a complex, realistic ecosystem.

159 **3.3 The LLM Agent Layer:**

160 The LLM layer introduces a news-aware reasoning component that influences agent behavior and
161 liquidity. The `EnhancedLLMNewsAnalyzer` accepts structured NewsEvents (category, content,
162 affected symbols, sentiment, and importance). It calls the LLM via an `LLMInterface` to analyze the
163 news and normalizes the responses into a quantifiable sentiment and confidence with risk annotations.
164 This output directly informs two key agent types:

- 165 • **ABIDESLLMTradingAgent:** This agent translates news analysis into strategy-specific sig-
166 nals (direction, strength, confidence, and duration). It then calculates position deltas using
167 its portfolio value and predefined position caps, placing market or limit orders in the LOB.
- 168 • **ABIDESLLMMarketMaker:** This agent quotes around the latest price and dynamically
169 adjusts its strategy. For instance, it widens spreads as uncertainty rises (based on
170 $1 - \text{confidence}$ in recent analyses) and re-centers its quotes to manage inventory.

171 Both agents expose ABIDES-style lifecycle hooks to ensure seamless integration with the simulation
172 kernel and message-passing system. This architecture creates a coupling that is both simple and
173 interpretable: news sentiment drives directional trading pressure, while uncertainty modulates market
174 resiliency and liquidity, producing realistic microstructure responses.

175 **3.4 The Brain-Memory Architecture**

176 Our adaptive agent operates on a Brain-Memory Architecture, separating strategic decision-making
177 from persistent, long-term learning.

178 The "Brain" is the strategic decision-maker, a large, high-capacity model (e.g., Gemini 1.5 Pro). This
179 agent uses a structured prompt to analyze a real-time state vector, including dynamic news headlines
180 and market data, before generating a response.

181 To overcome context window limitations and enable learning, the "Brain" is supported by a persistent
182 memory module, powered by a file-based database system. Before making a new decision, the
183 "Brain" agent retrieves relevant past experiences from this module. Each memory entry contains the
184 historical news event, the parameters the agent used, the action it took, and the resulting profit or loss.
185 This history is then formatted and injected directly into its prompt. This architecture allows the agent
186 to reflect on its past performance and learn from specific outcomes, a capability absent in traditional
187 agent-based models.

188 **3.5 The Multi-Agent Ecosystem: A Three-Layered Approach**

189 Our simulation operates as a multi-layered ecosystem, combining traditional rule-based agents with
190 our novel LLM-powered agents to create a realistic and dynamic market environment.

- 191 • **Foundation Layer:** This layer is composed of a diverse population of pre-existing
192 ABIDES agent types, which form the "baseline market." These agents, such as Zero-
193 Intelligence (ZI) agents, Fundamental Value agents, and algorithmic Market Makers, are not
194 intelligent in the human sense but follow predictable, algorithmic behaviors. This population
195 generates a complex, dynamic, and realistic synthetic market microstructure—the liquidity,
196 spreads, and order flow that an agent would see in a real-world exchange.
- 197 • **The LLM Integration:** Into this pre-existing, realistic environment, we introduce our
198 LLM-enhanced agents. These agents are not acting in a vacuum or a simplified setting.
199 Instead, they are placed within a competitive and realistic ecosystem populated by dozens of
200 other agents, whose collective actions create a high-fidelity market. The LLM agents' per-
201 formance and emergent behavior are thus validated against a robust and complex backdrop.
- 202 • **The Interaction Loop:** The "Brain" agent's primary function is to interpret the complex,
203 emergent dynamics of the market, which are generated by the diverse population of ABIDES
204 agents. It receives a comprehensive state vector that includes not only high-level information
205 like price and volume but also microstructural details such as order book depth and recent
206 trade prints. The "Brain" then uses its reasoning capabilities to synthesize this information
207 and form a trading hypothesis. The persistent memory module provides a crucial temporal
208 dimension, allowing the "Brain" to ground its real-time decisions in historical context from
209 within the simulated environment. For example, it might recall that a similar news headline
210 in the past led to a specific market reversal, enabling a more nuanced and strategic response.
211 Crucially, this feedback loop operates online, within a single simulation run. The outcome
212 of a decision made at hour one is evaluated after a fixed time window and becomes a learned
213 experience—written back to `memory.jsonl`—that informs the agent's strategy at hour two.
214 The orders generated by the "Brain" agent are then fed back into the LOB engine, influencing
215 the market and creating a dynamic feedback loop with the rest of the agent population.

216 To make the control flow explicit, we summarize the interaction loop in Algorithm 1. At each step
217 $t \in [0, T]$ with step size Δt , the agent aggregates the market snapshot m_t , news items n_t , and
218 retrieved memories to form $context_t$, then the Brain proposes a concrete ABIDES configuration C_t
219 which is validated and executed in the exchange-faithful core. The resulting actions and fills feed
220 metrics that are written back to Memory for within-episode adaptation. The next section details how
221 these outputs are aligned to NASDAQ's ITCH protocol for calibration and external validation.

Algorithm 1 Overall workflow: LLM \leftrightarrow ABIDES with Brain–Memory (condensed)

Require: H (history), M (mkt stream), N (news); Env, C_0 ; LLM, Memory
Ensure: logs, memory snapshot, performance summary

```
1:  $C \leftarrow C_0$ ; ABIDES  $\leftarrow$  Init(Env); Memory.load( $H$ ); set  $T, \Delta t$ 
2: for  $t = 0 : \Delta t : T$  do
3:    $m_t \leftarrow M.read()$ ;  $n_t \leftarrow N.read()$ ;  $ctx \leftarrow Agg(m_t, n_t, Memory.peek())$ 
4:   Memory.update( $ctx$ );  $mem \leftarrow Memory.retrieve\_relevant\_entries(ctx)$ 
5:    $C_t \leftarrow LLM.generateConfig(C, mem, ctx)$ ;  $C_t \leftarrow ValidateTranslate(C_t)$ 
6:   ABIDES.applyConfig( $C_t$ );  $(a_t, f_t) \leftarrow ABIDES.runStep(\Delta t)$ 
7:    $met_t \leftarrow Metrics(a_t, f_t, m_t)$ ; Memory.update_entry_with_outcome( $ctx, met_t, f_t$ )
8:   if learning enabled then
9:     LLM.onlineUpdate( $met_t, f_t, mem$ )
10:  end if
11:  Log( $t, C_t, a_t, met_t, Memory.digest()$ );  $C \leftarrow Carry(C_t)$ 
12: end for
13: return logs, Memory.snapshot(), PerfSummary()
```

222 **3.6 Data Ingestion, Calibration, and Validation:**

To ensure the realism of our synthetic environment, we ingest and replay real-world financial data at the message level from sources like NASDAQ’s ITCH protocol via the LOBSTER dataset. Our system is not merely reproducing historical price curves; it is accurately simulating the underlying market microstructure. The process, governed by scripts like `itch_data_parser.py` and `real_data_ingestion.py`, ensures that our synthetic data (`AAPL_2012-06-21...csv`) maintains the statistical properties of the original. We carefully calibrate the synthetic environment by validating our generated order flow, trade volume, and price volatility against the ground truth of historical NASDAQ data. This calibration relies on an optimization routine (`calibrate_simulator.py`) to find the optimal agent population parameters, stored in `optimal_config.json`, that best reproduce historical market behavior. For external validation, ITCH trades are loaded and aggregated to a per-second baseline. The system then aligns simulated execution timestamps to this per-second clock and compares prices, computing

$$price_error_bps = \frac{sim - ITCH}{ITCH} \times 10,000$$

223 and exporting a per-execution CSV and a metrics JSON. Finally, we plot the smoothed simulated
224 mid-price as a blue line with simulated trades as red scatter over time, using the ITCH first price and
225 start time for a direct visual comparison.226 **4 Results and Discussion**227 We conducted a series of experiments with our LLM enhanced ABIDES simulator, comparing three
228 conditions under identical seeds, book parameters, market configuration, schedules, and windows, all
229 calibrated to AMZN on June 21, 2012. The three agent populations were LLMON with an adaptive
230 Brain Memory agent, LLMOFF composed of traditional heuristic agents such as simple momentum
231 and contrarian traders, and a Baseline population dominated by noise traders.232 In LLMON the news analyzer queries the LLM to generate sentiment, confidence, and risk annotations
233 from structured news events; the trading agent maps these signals to direction and size within a
234 fixed risk budget with position caps and convex scaling; and the market maker widens spreads when
235 confidence is low and recenters quotes based on inventory pressure, yielding news conditioned order
236 flow, adaptive liquidity, and short horizon impact around events. In LLMOFF the analyzer returns
237 deterministic rule based sentiment and confidence without LLM calls, and trading and market making
238 follow the same mappings and risk controls but are driven by heuristic signals, serving as an ablation
239 that isolates the incremental value of LLM reasoning under identical settings. In BASELINE no
240 news is ingested or analyzed, agents run without event driven signals, trading relies on static strategy
241 parameters, and the market maker uses a fixed spread policy, providing a microstructure only baseline
242 that quantifies the contribution of news coupling.243 To probe sensitivity to exogenous information, we introduced four synthetic news events via a news-
244 injection interface. As shown in Figure 2, only the LLMON configuration consistently perceived




Figure 2: We evaluate the simulator on the LOBSTER dataset using AMZN on June 21, 2012 under three modes: LLM-ON, LLM-OFF, and Baseline. Results with LLM-ON confirm that the simulator can perceive news events and inject their effects into the simulated market.




Figure 3: We extract the execution prices of LLM-ON executed trades within a 1-minute interval to illustrate its high-frequency characteristics and rapid temporal dynamics.

245 these injections and propagated their effects into the simulated price dynamics, whereas LLMOFF
246 and Baseline showed little to no response.

247 The results demonstrate that the LLMON agent, leveraging its online learning capabilities via the
248 memory module, generates market dynamics that are qualitatively and quantitatively distinct. The
249 introduction of the learning LLM agent was observed to influence the market in ways that traditional
250 agents could not, such as dynamically altering its reaction strength to shifting sentiment from news
251 headlines. The agent's ability to adjust its parameters mid-simulation based on the profitability of its
252 prior actions showcases a strategic evolution absent in the static LLMOFF and Baseline conditions.

253 Furthermore, as shown in Figure 3, our analysis of the simulated order book and micro structural
254 1-minute trading data confirms that the adaptive LLM agent contributes to the formation of a more
255 realistic market microstructure. The trade intensity, bid-ask spread, and price volatility in the LLM-
256 enhanced simulation more closely match historical data, showcasing the emergent complexity that
257 arises from an intelligent agent learning within a competitive environment.

258 Moreover, as shown in Figure 4, our LLMON configuration aligns more closely with the real
259 market environment than either the Baseline or LLMOFF settings. Using identical initial conditions
260 and calibration to AMZN on June 21, 2012, LLMON reproduces key microstructure regularities,
261 including co-movement between spread and volatility, realistic order book depth and imbalance




Figure 4: **Left (a):** Real vs. simulated intraday mid-price overlays. **Right (b):** Full-session trajectories for the three simulated configurations. Each panel shows the complete trading day aligned to market open. LLMON exhibits a stronger directional drift, LLMOFF is more range-bound, and Baseline shows a mild decline.




Figure 5: We visualize simulator price responses to exogenous-impact injection using a heatmap. The LLMON setting exhibits a clearer and more accurate reflection of externally driven movements than the alternative modes.

262 profiles, and a stronger and timely price response to the injected news events, whereas LLMOFF and
263 Baseline show muted or delayed reactions and systematic deviations from the empirical benchmarks.

264 5 Conclusion

265 We present a high-fidelity, LLM-enhanced financial market simulator that leverages synthetic data
266 and a novel Brain-Memory Architecture to model realistic dynamics. Through an online learning
267 loop, we show that LLM agents can move beyond simple heuristics to exhibit complex, adaptive
268 behaviors, demonstrating the feasibility of using generative AI to simulate emergent phenomena in
269 real markets. Beyond finance, this framework offers a foundation for studying intelligent agents
270 in other high-frequency, adaptive domains such as bonds, emerging markets, and related trading
271 contexts.

272 Nevertheless, limitations remain: synthetic data cannot fully capture real-world subtleties, and current
273 LLMs face constraints in scale, consistency, and generalization. Addressing these challenges will
274 require integrating heterogeneous data sources, enhancing memory mechanisms, and developing
275 stronger evaluation metrics for emergent behaviors.

276 **References**

- 277 [1] Haitham Bou Ammar et al. “An automated measure of mdp similarity for transfer in reinforcement learning”. In: *28th AAAI Conference on Artificial Intelligence, AAAI 2014*. AI Access Foundation. 2014, pp. 31–37.
- 278 [2] Selim Amrouni et al. “ABIDES-Gym: Gym Environments for Multi-Agent Discrete Event Simulation and Application to Financial Markets”. In: *CoRR* abs/2110.14771 (2021). arXiv: 2110.14771. URL: <https://arxiv.org/abs/2110.14771>.
- 279 [3] W. Brian Arthur. “Complexity and the economy”. In: *Science* 284.5411 (1999), pp. 107–109.
- 280 [4] Hans Bühler et al. *Deep Hedging*. 2018. arXiv: 1802.03042 [q-fin.CP]. URL: <https://arxiv.org/abs/1802.03042>.
- 281 [5] David Byrd, Maria Hybinette, and Tucker Balch. “ABIDES: Towards High-Fidelity Market Simulation for AI Research”. In: *CoRR* abs/1904.12066 (2019). arXiv: 1904.12066. URL: <http://arxiv.org/abs/1904.12066>.
- 282 [6] Kshama Dwarakanath, Tucker Balch, and Svitlana Vyettrenko. *ABIDES-Economist: Agent-Based Simulator of Economic Systems with Learning Agents*. 2025. arXiv: 2402.09563 [cs.MA]. URL: <https://arxiv.org/abs/2402.09563>.
- 283 [7] Annalisa Fabretti. “On the problem of calibrating an agent based model for financial markets”. In: *Journal of Economic Interaction and Coordination* 8.2 (2013), pp. 277–293.
- 284 [8] Fernando Fernández and Manuela Veloso. “Learning domain structure through probabilistic policy reuse in reinforcement learning”. In: *Progress in Artificial Intelligence* 2.1 (2013), pp. 13–27.
- 285 [9] Dhananjay K Gode and Shyam Sunder. “Allocative efficiency of markets with zero-intelligence traders: Market as a partial substitute for individual rationality”. In: *Journal of political economy* 101.1 (1993), pp. 119–137.
- 286 [10] Mandeep Goyal and Qusay H. Mahmoud. “A Systematic Review of Synthetic Data Generation Techniques Using Generative AI”. In: *Electronics* 13.17 (2024). ISSN: 2079-9292. DOI: 10.3390/electronics13173509. URL: <https://www.mdpi.com/2079-9292/13/17/3509>.
- 287 [11] John H Holland. “Studying complex adaptive systems”. In: *Journal of systems science and complexity* 19.1 (2006), pp. 1–8.
- 288 [12] Narasimhan Jegadeesh and Sheridan Titman. “Returns to buying winners and selling losers: Implications for stock market efficiency”. In: *The Journal of finance* 48.1 (1993), pp. 65–91.
- 289 [13] Michaël Karpe et al. “Multi-agent reinforcement learning in a realistic limit order book market simulation”. In: *Proceedings of the First ACM International Conference on AI in Finance. ICAIF ’20*. ACM, Oct. 2020, pp. 1–7. DOI: 10.1145/3383455.3422570. URL: <http://dx.doi.org/10.1145/3383455.3422570>.
- 290 [14] David W. Lu. *Agent Inspired Trading Using Recurrent Reinforcement Learning and LSTM Neural Networks*. 2017. arXiv: 1707.07338 [q-fin.CP]. URL: <https://arxiv.org/abs/1707.07338>.
- 291 [15] Mohammad Mani, Steve Phelps, and Simon Parsons. “Applications of reinforcement learning in automated market-making”. In: *Proceedings of the GAIW: Games, Agents and Incentives Workshops, Montreal, Canada*. IFAAMS Np. 2019, pp. 13–14.
- 292 [16] Volodymyr Mnih et al. “Playing Atari with Deep Reinforcement Learning”. In: *CoRR* abs/1312.5602 (2013). arXiv: 1312.5602. URL: <http://arxiv.org/abs/1312.5602>.
- 293 [17] Yuriy Nevmyvaka, Yi Feng, and Michael Kearns. “Reinforcement learning for optimized trade execution”. In: *Proceedings of the 23rd International Conference on Machine Learning. ICML ’06*. Pittsburgh, Pennsylvania, USA: Association for Computing Machinery, 2006, pp. 673–680. ISBN: 1595933832. DOI: 10.1145/1143844.1143929. URL: <https://doi.org/10.1145/1143844.1143929>.
- 294 [18] Diego Pino et al. *Similarity metrics for Different Market Scenarios in Abides*. 2021. arXiv: 2107.09352 [cs.CE]. URL: <https://arxiv.org/abs/2107.09352>.
- 295 [19] Vamsi K. Potluru et al. *Synthetic Data Applications in Finance*. 2024. arXiv: 2401.00081 [cs.LG]. URL: <https://arxiv.org/abs/2401.00081>.
- 296 [20] Thomas Spooner et al. “Market Making via Reinforcement Learning”. In: *CoRR* abs/1804.04216 (2018). arXiv: 1804.04216. URL: <http://arxiv.org/abs/1804.04216>.

- 330 [21] R.S. Sutton and A.G. Barto. “Reinforcement Learning: An Introduction”. In: *IEEE Transactions on Neural Networks* 9.5 (1998), pp. 1054–1054. DOI: 10.1109/TNN.1998.712192.
- 331 [22] Jie Tan et al. *Sim-to-Real: Learning Agile Locomotion For Quadruped Robots*. 2018. arXiv: 1804.10332 [cs.RO]. URL: <https://arxiv.org/abs/1804.10332>.
- 332 [23] Lisa Torrey and Jude Shavlik. “Transfer learning”. In: *Handbook of research on machine learning applications and trends: algorithms, methods, and techniques*. IGI Global Scientific Publishing, 2010, pp. 242–264.
- 333 [24] Svitlana Vyettrenko et al. *Get Real: Realism Metrics for Robust Limit Order Book Market Simulations*. 2019. arXiv: 1912.04941 [q-fin.TR]. URL: <https://arxiv.org/abs/1912.04941>.
- 334
- 335
- 336
- 337
- 338
- 339

340 **Agents4Science AI Involvement Checklist**

- 341 1. **Hypothesis development:** Hypothesis development includes the process by which you
342 came to explore this research topic and research question. This can involve the background
343 research performed by either researchers or by AI. This can also involve whether the idea
344 was proposed by researchers or by AI.

345 Answer: **[A]**

346 Explanation: The idea of this research are entirely unique to humans

- 347 2. **Experimental design and implementation:** This category includes design of experiments
348 that are used to test the hypotheses, coding and implementation of computational methods,
349 and the execution of these experiments.

350 Answer: **[B]**

351 Explanation: Mostly humans, but in some scenarios, we use ai to assist.

- 352 3. **Analysis of data and interpretation of results:** This category encompasses any process to
353 organize and process data for the experiments in the paper. It also includes interpretations of
354 the results of the study.

355 Answer: **[B]**

356 Explanation: Mostly humans, but in some scenarios, we use ai to assist.

- 357 4. **Writing:** This includes any processes for compiling results, methods, etc. into the final
358 paper form. This can involve not only writing of the main text but also figure-making,
359 improving layout of the manuscript, and formulation of narrative.

360 Answer: **[B]**

361 Explanation: Mostly humans, but in some scenarios, we use ai to assist.

- 362 5. **Observed AI Limitations:** What limitations have you found when using AI as a partner or
363 lead author?

364 Description: AI can generate fluent, well-structured text, but it may struggle with critical
365 perspective and novel ideas.

366 **Agents4Science Paper Checklist**

367 **1. Claims**

368 Question: Do the main claims made in the abstract and introduction accurately reflect the
369 paper's contributions and scope?

370 Answer: [Yes]

371 Justification: yes, the main claim aligns with contribution and scope.

372 Guidelines:

- 373 • The answer NA means that the abstract and introduction do not include the claims
374 made in the paper.
- 375 • The abstract and/or introduction should clearly state the claims made, including the
376 contributions made in the paper and important assumptions and limitations. A No or
377 NA answer to this question will not be perceived well by the reviewers.
- 378 • The claims made should match theoretical and experimental results, and reflect how
379 much the results can be expected to generalize to other settings.
- 380 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
381 are not attained by the paper.

382 **2. Limitations**

383 Question: Does the paper discuss the limitations of the work performed by the authors?

384 Answer: [Yes]

385 Justification: yes, we have discussed it in the collusion.

386 Guidelines:

- 387 • The answer NA means that the paper has no limitation while the answer No means that
388 the paper has limitations, but those are not discussed in the paper.
- 389 • The authors are encouraged to create a separate "Limitations" section in their paper.
- 390 • The paper should point out any strong assumptions and how robust the results are to
391 violations of these assumptions (e.g., independence assumptions, noiseless settings,
392 model well-specification, asymptotic approximations only holding locally). The authors
393 should reflect on how these assumptions might be violated in practice and what the
394 implications would be.
- 395 • The authors should reflect on the scope of the claims made, e.g., if the approach was
396 only tested on a few datasets or with a few runs. In general, empirical results often
397 depend on implicit assumptions, which should be articulated.
- 398 • The authors should reflect on the factors that influence the performance of the approach.
399 For example, a facial recognition algorithm may perform poorly when image resolution
400 is low or images are taken in low lighting.
- 401 • The authors should discuss the computational efficiency of the proposed algorithms
402 and how they scale with dataset size.
- 403 • If applicable, the authors should discuss possible limitations of their approach to
404 address problems of privacy and fairness.
- 405 • While the authors might fear that complete honesty about limitations might be used by
406 reviewers as grounds for rejection, a worse outcome might be that reviewers discover
407 limitations that aren't acknowledged in the paper. Reviewers will be specifically
408 instructed to not penalize honesty concerning limitations.

409 **3. Theory assumptions and proofs**

410 Question: For each theoretical result, does the paper provide the full set of assumptions and
411 a complete (and correct) proof?

412 Answer: [NA]

413 Justification: the paper does not include theoretical results

414 Guidelines:

- 415 • The answer NA means that the paper does not include theoretical results.

- 416 • All the theorems, formulas, and proofs in the paper should be numbered and cross-
 417 referenced.
 418 • All assumptions should be clearly stated or referenced in the statement of any theorems.
 419 • The proofs can either appear in the main paper or the supplemental material, but if
 420 they appear in the supplemental material, the authors are encouraged to provide a short
 421 proof sketch to provide intuition.

422 **4. Experimental result reproducibility**

423 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
 424 perimental results of the paper to the extent that it affects the main claims and/or conclusions
 425 of the paper (regardless of whether the code and data are provided or not)?

426 Answer: [Yes]

427 Justification: we include experiments and can be reproduced.

428 Guidelines:

- 429 • The answer NA means that the paper does not include experiments.
 430 • If the paper includes experiments, a No answer to this question will not be perceived
 431 well by the reviewers: Making the paper reproducible is important.
 432 • If the contribution is a dataset and/or model, the authors should describe the steps taken
 433 to make their results reproducible or verifiable.
 434 • We recognize that reproducibility may be tricky in some cases, in which case authors
 435 are welcome to describe the particular way they provide for reproducibility. In the case
 436 of closed-source models, it may be that access to the model is limited in some way
 437 (e.g., to registered users), but it should be possible for other researchers to have some
 438 path to reproducing or verifying the results.

439 **5. Open access to data and code**

440 Question: Does the paper provide open access to the data and code, with sufficient instruc-
 441 tions to faithfully reproduce the main experimental results, as described in supplemental
 442 material?

443 Answer: [Yes]

444 Justification: it is open source

445 Guidelines:

- 446 • The answer NA means that paper does not include experiments requiring code.
 447 • Please see the Agents4Science code and data submission guidelines on the conference
 448 website for more details.
 449 • While we encourage the release of code and data, we understand that this might not be
 450 possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not
 451 including code, unless this is central to the contribution (e.g., for a new open-source
 452 benchmark).
 453 • The instructions should contain the exact command and environment needed to run to
 454 reproduce the results.
 455 • At submission time, to preserve anonymity, the authors should release anonymized
 456 versions (if applicable).

457 **6. Experimental setting/details**

458 Question: Does the paper specify all the training and test details (e.g., data splits, hyper-
 459 parameters, how they were chosen, type of optimizer, etc.) necessary to understand the
 460 results?

461 Answer: [Yes]

462 Justification: yes,it does.

463 Guidelines:

- 464 • The answer NA means that the paper does not include experiments.
 465 • The experimental setting should be presented in the core of the paper to a level of detail
 466 that is necessary to appreciate the results and make sense of them.

- 467 • The full details can be provided either with the code, in appendix, or as supplemental
468 material.

469 **7. Experiment statistical significance**

470 Question: Does the paper report error bars suitably and correctly defined or other appropriate
471 information about the statistical significance of the experiments?

472 Answer: [\[Yes\]](#)

473 Justification: yes,it does.

474 Guidelines:

- 475 • The answer NA means that the paper does not include experiments.
476 • The authors should answer "Yes" if the results are accompanied by error bars, confi-
477 dence intervals, or statistical significance tests, at least for the experiments that support
478 the main claims of the paper.
479 • The factors of variability that the error bars are capturing should be clearly stated
480 (for example, train/test split, initialization, or overall run with given experimental
481 conditions).

482 **8. Experiments compute resources**

483 Question: For each experiment, does the paper provide sufficient information on the com-
484 puter resources (type of compute workers, memory, time of execution) needed to reproduce
485 the experiments?

486 Answer: [\[Yes\]](#)

487 Justification: yes,it does.

488 Guidelines:

- 489 • The answer NA means that the paper does not include experiments.
490 • The paper should indicate the type of compute workers CPU or GPU, internal cluster,
491 or cloud provider, including relevant memory and storage.
492 • The paper should provide the amount of compute required for each of the individual
493 experimental runs as well as estimate the total compute.

494 **9. Code of ethics**

495 Question: Does the research conducted in the paper conform, in every respect, with the
496 Agents4Science Code of Ethics (see conference website)?

497 Answer: [\[Yes\]](#)

498 Justification: yes,it does.

499 Guidelines:

- 500 • The answer NA means that the authors have not reviewed the Agents4Science Code of
501 Ethics.
502 • If the authors answer No, they should explain the special circumstances that require a
503 deviation from the Code of Ethics.

504 **10. Broader impacts**

505 Question: Does the paper discuss both potential positive societal impacts and negative
506 societal impacts of the work performed?

507 Answer: [\[Yes\]](#)

508 Justification: yes,it does.

509 Guidelines:

- 510 • The answer NA means that there is no societal impact of the work performed.
511 • If the authors answer NA or No, they should explain why their work has no societal
512 impact or why the paper does not address societal impact.
513 • Examples of negative societal impacts include potential malicious or unintended uses
514 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations,
515 privacy considerations, and security considerations.
516 • If there are negative societal impacts, the authors could also discuss possible mitigation
517 strategies.