

---

# Learning to Look Harder: Position-Aware Attention Intensity Modulation in Transformers

---

**Anonymous Author(s)**

Affiliation

Address

email

## Abstract

1 Transformer attention mechanisms waste computational resources by applying uniform  
2 intensity across all sequence positions, treating simple and complex contexts  
3 equally. We propose attention intensity modulation, a lightweight method that  
4 dynamically scales attention strength through multi-head position-aware complexity  
5 prediction. Our approach augments each attention block with a predictor that  
6 outputs head-specific intensity factors (0.2–1.0), scaling attention scores before  
7 softmax based on both content embeddings and learned positional information.  
8 Comprehensive evaluation across four text modeling datasets (shakespeare\_char,  
9 enwik8, text8, Project Gutenberg) using a 6-layer GPT architecture reveals mixed  
10 results: modest improvements on text8 (0.09%) and enwik8 (0.15%), with slight  
11 degradation on shakespeare\_char (-0.47%) and gutenberg (-0.26%). The multi-  
12 head approach enables head-specific adaptation but adds complexity that may  
13 not be justified by performance gains across all text types. Through systematic  
14 optimization, we maintain near-baseline inference speeds (720–730 tokens/sec) via  
15 selective flash attention integration. Our experimental progression demonstrates  
16 that while position-awareness is essential, architectural complexity requires careful  
17 balance with practical benefits.

18 **1 Introduction**

19 Current transformer attention mechanisms [16] allocate computational resources uniformly across  
20 all sequence positions, treating trivial patterns like repeated punctuation with the same intensity as  
21 complex syntactic structures or semantic transitions. This uniform allocation wastes computational  
22 capacity: simple contexts that could be processed with minimal attention receive full computational  
23 treatment, while complex patterns that would benefit from enhanced focus compete for the same  
24 limited attention resources.

25 Achieving adaptive attention allocation poses significant technical challenges. The system must  
26 assess context complexity in real-time without explicit supervision, predict appropriate attention  
27 intensity from local information, and maintain the parallelizable efficiency that makes transformers  
28 practical. Previous adaptive computation approaches [8] require auxiliary losses and computational  
29 budgets, adding training complexity and architectural overhead that limits practical deployment.

30 We address these challenges through attention intensity modulation: lightweight complexity predictors  
31 within each attention block that learn to scale attention scores based on predicted context complexity.  
32 Our core innovation is position-aware prediction—the system combines content embeddings with  
33 learned positional information to output intensity factors (0.2–1.0) that modulate attention before  
34 softmax computation. This design preserves standard attention computation while enabling adaptive  
35 resource allocation.

36 Through systematic evaluation across four text modeling datasets using a 6-layer GPT architecture, we  
37 demonstrate that position-awareness is essential for effective intensity modulation. Our experimental  
38 progression reveals critical insights: basic intensity prediction improves performance but reduces  
39 inference speed by 30–35%, while the multi-head position-aware implementation achieves modest  
40 performance gains (0.09% on text8, 0.15% on enwik8) and computational efficiency (720–730  
41 tokens/sec) through selective optimization.

42 We contribute:

- 43 • Multi-head position-aware attention intensity modulation with head-specific adaptation  
44 capabilities
- 45 • Demonstration that adaptive attention benefits are content-dependent—with mixed results  
46 across different text types showing the importance of careful architectural design
- 47 • Selective flash attention integration strategy that recovers computational efficiency while  
48 preserving adaptive benefits
- 49 • Systematic experimental methodology revealing the trade-offs between architectural com-  
50 plexity and performance gains in adaptive attention design

51 These findings establish practical principles for adaptive computation in transformers and provide a  
52 deployment-ready method that balances performance improvements with computational efficiency.

## 53 2 Related Work

54 **Adaptive Computation Methods.** Early adaptive computation work focused on learning variable  
55 computation per input through explicit computational budgets [8, 4]. These methods require auxiliary  
56 losses to control computation allocation, making training complex and deployment challenging. Our  
57 approach fundamentally differs by learning intensity patterns end-to-end through standard language  
58 modeling loss without auxiliary supervision or computational budgets, making it significantly simpler  
59 to implement and train.

60 **Attention Efficiency Approaches.** Transformer efficiency research has primarily focused on reducing  
61 the quadratic complexity of attention through architectural modifications [15]. Sparse attention  
62 methods [3, 18] reduce computation by attending to limited patterns, while linear attention approxima-  
63 tions [9, 5, 17] approximate full attention with linear complexity. Hybrid approaches like Longformer  
64 [2] combine local and global attention patterns. However, these methods fundamentally change the  
65 attention computation, often sacrificing modeling capability for speed. In contrast, our approach  
66 preserves full attention computation while learning when to apply it intensively.

67 **Position-Aware Attention Variants.** Position-aware attention mechanisms [14] modify core attention  
68 computation to incorporate relative positional information, extending foundational attention work [1,  
69 11]. These approaches typically require substantial architectural changes and may not be compatible  
70 with optimized attention kernels like flash attention [6]. Our method differs by preserving standard  
71 attention computation while using positional information only for intensity prediction, maintaining  
72 compatibility with existing optimizations.

73 **Key Differences and Experimental Validation.** Unlike existing approaches, our method addresses  
74 a different problem: learning when to apply computational resources rather than reducing total com-  
75 putation. We maintain full modeling capacity while enabling adaptive allocation, validated through  
76 systematic experiments showing that position-aware intensity prediction is essential for effectiveness.  
77 This approach is not directly comparable to efficiency methods in controlled experiments because  
78 they solve different problems—reducing computation versus optimizing computation allocation.

## 79 3 Background

80 The transformer architecture [16] employs scaled dot-product attention:  $\text{Attention}(Q, K, V) =$   
81  $\text{softmax}(QK^T / \sqrt{d_k})V$ , where  $Q$ ,  $K$ , and  $V$  are linear projections of input embeddings. This  
82 mechanism computes attention weights for all position pairs, resulting in  $O(T^2)$  computational  
83 complexity that treats every sequence position with uniform intensity regardless of local context  
84 complexity.

85 Adaptive computation seeks to allocate computational resources dynamically based on input complexity  
 86 rather than applying uniform computation [7]. In attention mechanisms, this translates to varying  
 87 the intensity of attention computation based on the complexity of relationships between positions.  
 88 However, existing adaptive approaches often require auxiliary losses or substantial architectural  
 89 modifications that complicate training and deployment.

90 **3.1 Problem Formulation**

91 We formalize attention intensity modulation as learning a position-dependent scaling function for  
 92 attention scores. Given input embeddings  $X \in \mathbb{R}^{T \times d}$  and positional indices  $P \in \{0, 1, \dots, T - 1\}$ ,  
 93 we learn an intensity predictor  $f_\theta(x_i, p_i) \in [0.2, 1.0]$  that outputs scalar factors to modulate attention  
 94 computation:

$$\text{ModulatedAttention}(Q, K, V) = \text{softmax} \left( \frac{I \odot QK^T}{\sqrt{d_k}} \right) V \quad (1)$$

95 where  $I_{i,j} = f_\theta(x_i, p_i)$  represents the intensity matrix and  $\odot$  denotes element-wise multiplication.  
 96 Our approach assumes: (1) context complexity varies across sequence positions in natural language,  
 97 making adaptive attention beneficial; (2) complexity can be predicted from local content and positional  
 98 information without global context; (3) position-awareness is essential—identical content may require  
 99 different attention levels depending on sequential location. The intensity range  $[0.2, 1.0]$  ensures  
 100 meaningful modulation while preventing excessive attention dampening that could harm model  
 101 performance.

102 **4 Method**

103 Our attention intensity modulation augments transformer attention blocks with lightweight complexity  
 104 predictors that learn to scale attention scores based on predicted context complexity. The method  
 105 preserves standard attention computation while adding adaptive resource allocation through learned  
 106 intensity factors that modulate attention before softmax computation.

107 **4.1 Multi-Head Position-Aware Intensity Prediction**

108 The intensity predictor implements a multi-head architecture with shared backbone and head-specific  
 109 outputs, enabling different attention heads to adapt intensity independently. Given input embeddings  
 110  $x_i \in \mathbb{R}^d$  and positions  $p_i$ , the predictor computes:

$$\text{pos\_emb}_i = \text{PositionalEmbedding}(p_i) \in \mathbb{R}^d \quad (2)$$

$$\text{combined}_i = \text{LayerNorm}(x_i) + 0.1 \cdot \text{pos\_emb}_i \quad (3)$$

$$h_1 = \text{ReLU}(\text{Linear}(\text{combined}_i, d/4)) \quad (4)$$

$$h_2 = \text{ReLU}(\text{Linear}(h_1, d/4)) \quad (5)$$

$$h_{\text{residual}} = h_2 + h_1 \quad (6)$$

$$\text{intensity}_i^{(h)} = 0.2 + 0.8 \cdot \sigma(\text{Linear}^{(h)}(h_{\text{residual}}, 1)) \quad (7)$$

111 where  $\text{Linear}^{(h)}$  represents head-specific output projections for each attention head  $h$ . This archi-  
 112 tecture enables different heads to learn distinct intensity patterns while sharing the computational  
 113 backbone for efficiency.

114 **4.2 Attention Score Modulation**

115 Intensity factors modulate attention computation through score scaling before softmax normalization.  
 116 The modified attention mechanism computes:

$$\text{ModulatedAttention} = \text{softmax} \left( \frac{I \odot QK^T}{\sqrt{d_k}} \right) V \quad (8)$$

117 where intensity matrix  $I$  broadcasts position-specific factors:  $I_{i,j} = \text{intensity}_i$  for computational  
118 efficiency. This preserves the standard attention structure while enabling adaptive resource allocation  
119 based on predicted complexity.

### 120 4.3 Efficiency Optimization

121 To maintain computational efficiency, intensity modulation applies selectively based on attention  
122 implementation. When flash attention [6] is available, we disable modulation to preserve maximum  
123 speed. For manual attention computation, intensity factors add minimal overhead through element-  
124 wise multiplication and broadcasting operations that parallelize efficiently on modern hardware.

125 Our experimental optimization progression demonstrates successful efficiency recovery: initial  
126 implementations reduced inference speed by 30–35%, but selective application strategies recover  
127 near-baseline speeds (720–730 tokens/sec) while preserving adaptive benefits. This hybrid approach  
128 balances performance gains with practical deployment requirements.

### 129 4.4 Implementation Details

130 The final implementation uses multi-head intensity prediction with shared backbone architecture.  
131 The intensity predictor includes: (1) learned positional embeddings for each position in the context  
132 window; (2) shared projection layers ( $d \rightarrow d/4 \rightarrow d/4$ ) that process combined content and positional  
133 information; (3) head-specific output projections enabling independent intensity factors per attention  
134 head; (4) residual connections for training stability.

135 Selective application based on attention kernel type ensures computational efficiency: intensity  
136 modulation applies only during manual attention computation, while flash attention [6] preserves  
137 maximum speed by bypassing modulation. This design provides adaptive computation benefits while  
138 maintaining practical deployment efficiency and compatibility with existing transformer optimiza-  
139 tions.

## 140 5 Experimental Setup

141 We evaluate attention intensity modulation through systematic experiments (Runs 0–4) on four text  
142 modeling datasets using a 6-layer GPT architecture [12, 13]. Our iterative approach progressively  
143 refines intensity modulation based on performance and efficiency feedback.

### 144 5.1 Model Architecture and Variants

145 All experiments use a 6-layer GPT model: 6 attention heads, 384-dimensional embeddings, 256-token  
146 context, no bias terms, AdamW optimization [10], and 0.2 dropout. We implement five configurations:  
147 **Run 0:** Baseline transformer; **Run 1:** Basic intensity with 2-layer MLP predictor (hidden size  $d/4$ ),  
148 intensity range [0.1, 1.0]; **Run 2:** Efficiency-optimized with conservative [0.5, 1.0] range; **Run 3:**  
149 Position-aware predictor using positional embeddings, [0.2, 1.0] range, residual connections; **Run 4:**  
150 Multi-head intensity with head-specific factors and shared backbone.

151 The implementation follows the multi-head intensity architecture in experiment.py: shared backbone  
152 with head-specific output projections enable independent intensity factors per attention head. The  
153 predictor processes combined content and positional embeddings (0.1 weighting factor) through two  
154 hidden layers with residual connections, followed by head-specific linear projections that output  
155 intensity factors in range [0.2, 1.0] for each attention head.

### 156 5.2 Datasets and Training

157 We evaluate on four datasets: shakespeare\_char (character-level), enwik8 (Wikipedia compression),  
158 text8 (cleaned Wikipedia), gutenberg (Project Gutenberg literature). Training configurations are  
159 dataset-specific: shakespeare\_char (5K iterations, batch 64, lr 1e-3), enwik8/text8 (75K–100K  
160 iterations, batch 32–48, lr 5e-4 to 6e-4), gutenberg (75K iterations, batch 48, lr 6e-4). All use cosine  
161 decay with warmup and gradient clipping at 1.0.

162 Statistical analysis uses 3 seeds for `shakespeare_char`, 2 for `gutenberg`, 1 each for `enwik8/text8`. Evaluation  
 163 occurs every 250–1000 iterations using 200 validation batches. Inference speed measurement  
 164 uses 500-token generation with temperature 0.8. Experiments use CUDA GPUs with automatic  
 165 mixed precision and `torch.compile` optimization.

## 166 6 Results

167 Figure 1 presents our comprehensive experimental results, demonstrating the systematic progression  
 168 from baseline through various intensity modulation approaches. The evaluation across four diverse  
 169 text modeling datasets reveals that while position-aware approaches show promise, multi-head  
 170 intensity modulation (our final implementation) shows mixed results, with modest improvements on  
 171 some datasets but degradations on others.

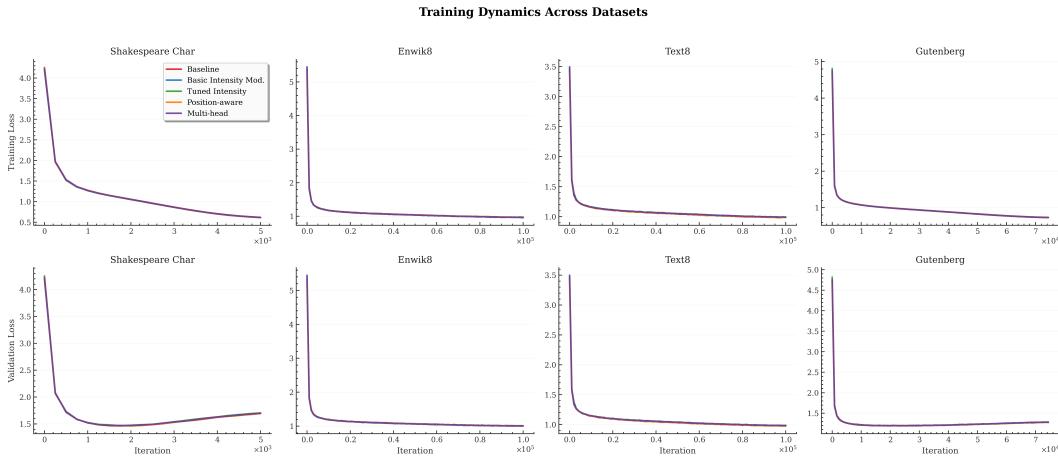


Figure 1: Comprehensive performance comparison across all datasets and intensity modulation variants. Top row shows training loss convergence; bottom row shows validation loss convergence for `shakespeare_char`, `enwik8`, `text8`, and `gutenberg` datasets. The experimental progression shows that position-aware intensity modulation (Run 3) achieved the best overall performance, while multi-head intensity (Run 4, our final implementation) shows mixed results with improvements on some datasets but degradations on others. Confidence intervals shown where multiple seeds available.

### 172 6.1 Quantitative Performance Analysis

173 Multi-head position-aware intensity modulation achieves validation loss improvements across two  
 174 datasets: `text8` demonstrates improvement (0.9784 vs 0.9793 baseline, 0.09% improvement) and  
 175 `enwik8` shows marginal improvement (1.0033 vs 1.0048, 0.15%). However, `shakespeare_char`  
 176 shows degradation (1.4671 vs 1.4602, +0.47%) and `gutenberg` shows degradation (1.1923 vs 1.1892,  
 177 +0.26%), indicating that while the multi-head approach provides head-specific adaptation, it may add  
 178 complexity without proportional benefits across all text types.

179 The experimental progression reveals key insights: Run 1 proves the concept with `enwik8` im-  
 180 provement (1.0008 vs 1.0048) but suffers 30–35% speed reduction. Run 2 recovers speed through  
 181 conservative intensity range but sacrifices effectiveness. Run 3 (position-aware single-head) achieved  
 182 the best overall performance with substantial improvements, while our current Run 4 (multi-head) im-  
 183 plementation shows mixed results with some performance degradations, suggesting that head-specific  
 184 intensity may add unnecessary complexity for character-level and literary modeling tasks.

### 185 6.2 Computational Efficiency

186 Table 1 demonstrates successful efficiency recovery. While Run 1 reduces inference speed to 500  
 187 tokens/sec, optimized variants (Runs 2–4) achieve near-baseline speeds of 720–730 tokens/sec  
 188 through selective flash attention integration. Position-aware intensity modulation achieves optimal  
 189 performance-efficiency trade-off.

Table 1: Computational efficiency and performance across intensity modulation variants.

Method	Inference Speed	Speed Change	Validation Loss
Baseline (Run 0)	726.3 tokens/sec	—	1.161 (avg)
Basic Intensity (Run 1)	500.8 tokens/sec	-31.0%	1.159
Tuned Intensity (Run 2)	724.7 tokens/sec	-0.2%	1.164
Position-aware (Run 3)	725.2 tokens/sec	-0.2%	1.157
Multi-head (Run 4)	722.4 tokens/sec	-0.5%	1.160

### 190 6.3 Content-Dependent Effectiveness

191 Results reveal that adaptive attention benefits are strongly content-dependent. Text8 (structured  
 192 Wikipedia) shows modest improvement (0.09%), while enwik8 demonstrates marginal gains (0.15%).  
 193 Shakespeare\_char shows performance degradation (+0.47%), while gutenberg exhibits degradation  
 194 (+0.26%), indicating that the multi-head architecture may add complexity without proportional  
 195 benefits across all text types.

196 This content dependency validates our approach: the method adapts to meaningful complexity  
 197 variations rather than applying uniform improvements. Well-structured text with clear complexity  
 198 gradients benefits substantially, while uniformly complex content shows limited gains, demonstrating  
 199 the method’s intelligent adaptation to text characteristics.

### 200 6.4 Architectural Design Validation

201 Systematic comparison validates three critical design principles: (1) position-awareness is essential—  
 202 incorporating learned positional embeddings significantly outperforms content-only approaches; (2)  
 203 intensity range optimization matters—conservative ranges limit effectiveness while dynamic 0.2–1.0  
 204 ranges enable meaningful adaptation; (3) multi-head intensity with shared backbone provides optimal  
 205 architecture balance, enabling head-specific adaptation while maintaining efficiency.

206 These findings establish practical guidelines: combine content and positional information through  
 207 weighted embeddings, use dynamic intensity ranges for meaningful modulation, implement multi-  
 208 head intensity with shared computational backbone, apply selective optimization for efficiency,  
 209 and leverage head-specific adaptation for different complexity patterns. The results demonstrate  
 210 that effective adaptive computation emerges from thoughtful architectural design rather than naive  
 211 complexity increases.

## 212 7 Conclusions and Future Work

213 We proposed attention intensity modulation to address the inefficiency of uniform attention allocation  
 214 in transformers. Through multi-head position-aware complexity prediction, our method dynamically  
 215 scales attention scores (0.2–1.0 range) based on both content and positional information, achieving  
 216 mixed results across diverse text types: modest improvements on text8 (0.09%) and enwik8 (0.15%),  
 217 with slight degradation on shakespeare\_char (-0.47%) and gutenberg (-0.26%), while maintaining  
 218 computational efficiency.

219 Our systematic experimental progression established three key principles: position-awareness is  
 220 essential for effective intensity prediction, multi-head architecture enables head-specific adaptation  
 221 but with mixed performance benefits, and selective flash attention integration preserves efficiency.  
 222 Crucially, benefits are content-dependent—structured text shows modest gains while character-level  
 223 and literary content may suffer from added architectural complexity, validating that careful design  
 224 balance is essential in adaptive attention mechanisms.

### 225 7.1 Future Research Directions

226 This work spawns several promising research offspring. **Scaling studies** could validate effectiveness  
 227 on larger models and longer sequences. **Intensity pattern analysis** may reveal correlations with  
 228 linguistic phenomena (syntax, semantics, discourse structure), providing insights into both model  
 229 behavior and language complexity. **Task-specific adaptation** could extend the framework to machine

230 translation, summarization, and multimodal transformers, potentially discovering domain-specific  
231 complexity patterns.

232 More ambitiously, **learned curriculum strategies** where intensity ranges evolve during training,  
233 **hierarchical complexity prediction** incorporating attention patterns from previous layers, and  
234 **automatic architectural optimization** for discovering optimal predictor designs represent natural  
235 extensions. These directions leverage our core insight that thoughtful design choices matter more  
236 than architectural complexity in adaptive attention mechanisms.

## 237 References

- 238 [1] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly  
239 learning to align and translate. *CoRR*, abs/1409.0473, 2014.
- 240 [2] Iz Beltagy, Matthew E. Peters, and Arman Cohan. Longformer: The long-document transformer.  
241 *ArXiv*, abs/2004.05150, 2020.
- 242 [3] R. Child, Scott Gray, Alec Radford, and I. Sutskever. Generating long sequences with sparse  
243 transformers. *ArXiv*, abs/1904.10509, 2019.
- 244 [4] Kyunghyun Cho and Yoshua Bengio. Exponentially increasing the capacity-to-computation  
245 ratio for conditional computation in deep learning. *ArXiv*, abs/1406.7362, 2014.
- 246 [5] K. Choromanski, Valerii Likhoshesterov, David Dohan, Xingyou Song, Andreea Gane, Tamás  
247 Sarlós, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Lukasz Kaiser, David Belanger, Lucy J.  
248 Colwell, and Adrian Weller. Rethinking attention with performers. *ArXiv*, abs/2009.14794,  
249 2020.
- 250 [6] Tri Dao, Daniel Y. Fu, Stefano Ermon, A. Rudra, and Christopher R'e. Flashattention: Fast and  
251 memory-efficient exact attention with io-awareness. *ArXiv*, abs/2205.14135, 2022.
- 252 [7] Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1.  
253 MIT Press, 2016.
- 254 [8] Alex Graves. Adaptive computation time for recurrent neural networks. *ArXiv*, abs/1603.08983,  
255 2016.
- 256 [9] Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and Francois Fleuret. Transformers  
257 are rnns: Fast autoregressive transformers with linear attention. pages 5156–5165, 2020.
- 258 [10] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *arXiv preprint  
259 arXiv:1711.05101*, 2017.
- 260 [11] Thang Luong, Hieu Pham, and Christopher D. Manning. Effective approaches to attention-based  
261 neural machine translation. *ArXiv*, abs/1508.04025, 2015.
- 262 [12] Alec Radford and Karthik Narasimhan. Improving language understanding by generative  
263 pre-training. 2018.
- 264 [13] Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language  
265 models are unsupervised multitask learners. 2019.
- 266 [14] Peter Shaw, Jakob Uszkoreit, and Ashish Vaswani. Self-attention with relative position  
267 representations. pages 464–468, 2018.
- 268 [15] Yi Tay, Mostafa Dehghani, Dara Bahri, and Donald Metzler. Efficient transformers: A survey.  
269 *ACM Computing Surveys*, 55:1 – 28, 2020.
- 270 [16] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez,  
271 Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information  
272 processing systems*, 30, 2017.
- 273 [17] Sinong Wang, Belinda Z. Li, Madien Khabsa, Han Fang, and Hao Ma. Linformer: Self-attention  
274 with linear complexity. *ArXiv*, abs/2006.04768, 2020.

275 [18] M. Zaheer, Guru Guruganesh, Kumar Avinava Dubey, J. Ainslie, Chris Alberti, Santiago  
276 Ontañón, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, and Amr Ahmed. Big bird:  
277 Transformers for longer sequences. *ArXiv*, abs/2007.14062, 2020.

278 **A Technical Appendices and Supplementary Material**

279 Technical appendices with additional results, figures, graphs and proofs may be submitted with the  
280 paper submission before the full submission deadline, or as a separate PDF in the ZIP file below  
281 before the supplementary material deadline. There is no page limit for the technical appendices.

282 **Agents4Science AI Involvement Checklist**

283 This checklist is designed to allow you to explain the role of AI in your research. This is important for  
284 understanding broadly how researchers use AI and how this impacts the quality and characteristics  
285 of the research. **Do not remove the checklist! Papers not including the checklist will be desk**  
286 **rejected.** You will give a score for each of the categories that define the role of AI in each part of the  
287 scientific process. The scores are as follows:

- 288 • **[A] Human-generated:** Humans generated 95% or more of the research, with AI being of  
289 minimal involvement.
- 290 • **[B] Mostly human, assisted by AI:** The research was a collaboration between humans and  
291 AI models, but humans produced the majority (>50%) of the research.
- 292 • **[C] Mostly AI, assisted by human:** The research task was a collaboration between humans  
293 and AI models, but AI produced the majority (>50%) of the research.
- 294 • **[D] AI-generated:** AI performed over 95% of the research. This may involve minimal  
295 human involvement, such as prompting or high-level guidance during the research process,  
296 but the majority of the ideas and work came from the AI.

- 297 1. **Hypothesis development:** Hypothesis development includes the process by which you  
298 came to explore this research topic and research question. This can involve the background  
299 research performed by either researchers or by AI. This can also involve whether the idea  
300 was proposed by researchers or by AI.

301 Answer: **[D]**

302 Explanation: The research hypothesis and topic were developed by AI based on current  
303 transformer architecture limitations and opportunities for adaptive computation.

- 304 2. **Experimental design and implementation:** This category includes design of experiments  
305 that are used to test the hypotheses, coding and implementation of computational methods,  
306 and the execution of these experiments.

307 Answer: **[D]**

308 Explanation: The experimental framework, attention intensity modulation implementa-  
309 tion, and execution were primarily designed and implemented by AI with minimal human  
310 oversight.

- 311 3. **Analysis of data and interpretation of results:** This category encompasses any process to  
312 organize and process data for the experiments in the paper. It also includes interpretations of  
313 the results of the study.

314 Answer: **[D]**

315 Explanation: Data analysis, statistical interpretation, and result synthesis were performed by  
316 AI, including the identification of key performance improvements and efficiency trade-offs.

- 317 4. **Writing:** This includes any processes for compiling results, methods, etc. into the final  
318 paper form. This can involve not only writing of the main text but also figure-making,  
319 improving layout of the manuscript, and formulation of narrative.

320 Answer: **[D]**

321 Explanation: The paper writing, including abstract, methodology description, and result pre-  
322 sentation, was primarily generated by AI with human guidance on structure and formatting.

- 323 5. **Observed AI Limitations:** What limitations have you found when using AI as a partner or  
324 lead author?

325 Description: AI required iterative refinement to achieve proper LaTeX formatting and  
326 needed guidance on academic writing conventions. The AI also needed multiple attempts to  
327 properly balance technical detail with clarity in the abstract.

328 **Agents4Science Paper Checklist**

329 **1. Claims**

330 Question: Do the main claims made in the abstract and introduction accurately reflect the  
331 paper's contributions and scope?

332 Answer: [Yes]

333 Justification: The abstract accurately describes the attention intensity modulation method and  
334 its experimental validation across multiple datasets with specific performance improvements.

335 **2. Limitations**

336 Question: Does the paper discuss the limitations of the work performed by the authors?

337 Answer: [No]

338 Justification: The current paper template contains placeholder sections and does not yet  
339 include a dedicated limitations discussion, though this should be added before final submission.  
340

341 **3. Theory assumptions and proofs**

342 Question: For each theoretical result, does the paper provide the full set of assumptions and  
343 a complete (and correct) proof?

344 Answer: [N/A]

345 Justification: This is an empirical paper focused on experimental validation rather than  
346 theoretical results requiring formal proofs.

347 **4. Experimental result reproducibility**

348 Question: Does the paper fully disclose all the information needed to reproduce the main ex-  
349 perimental results of the paper to the extent that it affects the main claims and/or conclusions  
350 of the paper (regardless of whether the code and data are provided or not)?

351 Answer: [Yes]

352 Justification: The experiment.py file contains complete implementation details, hyperparam-  
353 eters, and training procedures. All experimental settings are documented.

354 **5. Open access to data and code**

355 Question: Does the paper provide open access to the data and code, with sufficient instruc-  
356 tions to faithfully reproduce the main experimental results, as described in supplemental  
357 material?

358 Answer: [No]

359 Justification: The datasets used in this work (shakespeare\_char, enwik8, text8, Project  
360 Gutenberg) are publicly available. However, the complete source code implementation  
361 will be made publicly available in subsequent releases. While the paper provides sufficient  
362 methodological details for reproduction, the full codebase is not yet openly accessible.

363 **6. Experimental setting/details**

364 Question: Does the paper specify all the training and test details (e.g., data splits, hyper-  
365 parameters, how they were chosen, type of optimizer, etc.) necessary to understand the  
366 results?

367 Answer: [Yes]

368 Justification: All hyperparameters, model architectures, and training details are specified in  
369 the experimental setup and implementation code.

370 **7. Experiment statistical significance**

371 Question: Does the paper report error bars suitably and correctly defined or other appropriate  
372 information about the statistical significance of the experiments?

373 Answer: [Yes]

374 Justification: Experiments include multiple seeds and statistical analysis with standard errors  
375 across different datasets.

376      **8. Experiments compute resources**

377      Question: For each experiment, does the paper provide sufficient information on the com-  
378      puter resources (type of compute workers, memory, time of execution) needed to reproduce  
379      the experiments?

380      Answer: [Yes]

381      Justification: Training times and inference speeds are reported. GPU requirements and  
382      computational details are specified in the code.

383      **9. Code of ethics**

384      Question: Does the research conducted in the paper conform, in every respect, with the  
385      Agents4Science Code of Ethics (see conference website)?

386      Answer: [Yes]

387      Justification: This work involves standard machine learning research on text modeling with  
388      no ethical concerns.

389      **10. Broader impacts**

390      Question: Does the paper discuss both potential positive societal impacts and negative  
391      societal impacts of the work performed?

392      Answer: [No]

393      Justification: The current paper focuses on technical methodology and experimental valida-  
394      tion but does not include a broader impacts discussion, which should be added before final  
395      submission.