
Beyond Game Theory Optimal: Profit-Maximizing Poker Agents for No-Limit Hold'em

Anonymous Author(s)

Affiliation

Address

email

Abstract

Game theory has grown into a major field over the past few decades, and poker has long served as one of its key case studies. Game-Theory-Optimal (GTO) provides strategies to avoid loss in poker, but pure GTO does not guarantee maximum profit. To this end, we aim to develop a model that outperforms GTO strategies to maximize profit in No Limit Hold'em, in heads-up (two-player) and multi-way (more than two-player) situations. Our model finds the GTO foundation and goes further to exploit opponents. The model first navigates toward many simulated poker hands against itself and keeps adjusting its decisions until no action can reliably beat it, creating a strong baseline that is close to the theoretical best strategy. Then, it adapts by observing opponent behavior and adjusting its strategy to capture extra value accordingly. Our results indicate that Monte-Carlo Counterfactual Regret Minimization (CFR) performs best in heads-up situations and CFR remains the strongest method in most multi-way situations. By combining the defensive strength of GTO with real-time exploitation, our approach aims to show how poker agents can move from merely not losing to consistently winning against diverse opponents.

1 Introduction

Poker has evolved from a niche pastime into a global mind sport and a natural laboratory for studying decision making under imperfect information. Online platforms and televised tournaments now generate billions of recorded hands each year, enabling quantitative analyses of risk, bluffing, and opponent modelling. No-Limit Texas Hold'em (NLHE) is especially prominent: state-of-the-art solvers compute game-theory-optimal (GTO) strategies—approximate Nash equilibria that guarantee long-term unexploitable play. Yet in real games, opponents consistently deviate from equilibrium. As a long-time poker enthusiast and researcher in multi-agent learning, the largest profits come from detecting and exploiting such deviations in real time. Current poker AIs remain either purely GTO (safe but conservative) or trained on static population tendencies (often slow to adjust), leaving a gap for dynamic, table-specific exploitation with provable safety. [1, 2, 3]

Building on these observations, we investigate algorithmic approaches that can both approximate game-theoretic optimal play and adapt to changing opponents. A central concept is *counterfactual regret*, which quantifies how much better a player could have done in hindsight by choosing a different action in a given decision situation. Counterfactual Regret Minimization (CFR) repeatedly simulates play, measures these regrets, and adjusts its strategy to minimize them, a process that provably converges to a Nash equilibrium in *two-player zero-sum finite extensive-form* games and, in practice, on abstracted NLHE where card and action spaces are discretized.. [4, 5, 6] Monte-Carlo CFR (MCCFR) improves computational efficiency by sampling single trajectories instead of traversing the full game tree. [7, 8, 9]

37 Deep CFR extends this framework with neural networks that generalize across large state spaces,
38 while Neural Fictitious Self-Play (NFSP) blends reinforcement learning and supervised learning to
39 maintain an average policy that also approaches equilibrium. These algorithms define today’s standard
40 toolkit for large-scale imperfect-information games, providing the backbone for most competitive
41 poker AIs. [10, 11, 12]

42 Despite their success, purely equilibrium-seeking methods are slow to exploit opponents who deviate
43 from optimal play, and population-trained agents often fail to adapt when table dynamics differ from
44 historical data. This motivates our proposed approach: an adaptive model that learns GTO behaviour
45 from self-play while continuously tracking opponents’ tendencies and shifting strategy in real time to
46 capture excess value, all while maintaining provable safety against counter-exploitation. [13]

47 In summary, our contributions are as follows:

- 48 1. Evaluate whether leading algorithms can converge to pure game-theory-optimal (GTO)
49 strategy in No-Limit Hold’em.
- 50 2. Examine their ability to construct adaptive GTO-like strategies that respond to specific
51 opponent behaviors.
- 52 3. Extend this adaptive framework to multi-player (multiway) scenarios to test scalability
53 beyond heads-up plays.

54 This research was conducted in full compliance with the Agents4Science Code of Ethics. It involves
55 only synthetic poker decision states and contains no human or sensitive data, follows principles of
56 scientific integrity and reproducibility, and poses no foreseeable harm to people, animals, or the
57 environment.

58 2 Related Work

59 **Models.** **Counterfactual Regret Minimization (CFR)** is a foundational algorithm for solving large
60 two-player zero-sum imperfect-information games. It iteratively traverses the game tree, computing
61 counterfactual regrets at every information set and adjusting action probabilities to minimize those
62 regrets, and the average of the successive strategies provably converges to a Nash equilibrium. Since
63 its introduction, CFR has become the standard baseline for equilibrium computation in heads-up
64 no-limit Texas Hold’em and related poker games. [4, 5, 14, 6] **Monte-Carlo CFR (MCCFR)**
65 improves the scalability of CFR by sampling single trajectories rather than exhaustively traversing the
66 full game tree. By replacing deterministic updates with stochastic sampling, MCCFR dramatically
67 reduces memory and runtime while retaining theoretical convergence guarantees, and it is widely
68 used when exact traversal is impractical, serving as the default method for building strong poker
69 agents at reasonable computational cost. [7, 8, 9] **Deep CFR** further extends this framework by
70 replacing the tabular regret and strategy tables of CFR with neural function approximators: a regret
71 network predicts counterfactual regrets and a policy network predicts the average strategy, enabling
72 generalization across similar states and scalability to much larger action and information spaces.
73 This deep-learning extension has been used to train near-equilibrium strategies in very large no-limit
74 Hold’em subgames and other imperfect-information domains. [10, 15] **Neural Fictitious Self-Play**
75 (**NFSP**) combines reinforcement learning with supervised learning to approximate fictitious play
76 in large games. A reinforcement-learning component continually improves a best-response policy,
77 while a supervised component maintains an average policy that approaches equilibrium, offering a
78 fully online, self-play training regime that has been demonstrated on full-scale poker as well as other
79 multi-agent settings. [11, 12, 16]

80 **Random Policy.** A uniform random policy serves as a non-strategic baseline: at each decision point
81 it samples among legal actions with equal probability. Although it has no convergence guarantees
82 and performs poorly in practice, it provides a lower bound for evaluating how much structure the
83 learning algorithms extract from the game.

84 Together these prior methods define the standard algorithmic landscape for equilibrium approximation
85 and adaptive play in large imperfect-information games. Our work builds on this literature by
86 benchmarking all four learning algorithms and a random baseline within a unified experimental
87 framework and by quantifying their distance to a strong MCCFR-trained GTO proxy.

88 **3 No Limit Texas Hold'em (NLHE) Basics**

89 No Limit Texas Hold'em (NLHE) is the most widely played variant of poker in both live cash games
90 and tournaments. Each player is dealt two private hole cards, followed by five community cards dealt
91 face up in three stages: the flop (three cards), turn (one card), and river (one card). Betting rounds
92 occur after the hole cards and after each community stage. “No Limit” means a player may wager
93 any amount of their remaining chips at any time, from the minimum bet to an all-in shove, making
94 stack depth and bet sizing central to strategic decision making.

95 The goal is to form the best possible five-card hand using any combination of the two hole cards
96 and the five community cards, or to win the pot uncontested through betting. For example, holding
97 A \heartsuit K \heartsuit on a board of Q \heartsuit J \clubsuit 5 \diamond 10 \heartsuit 2 \spadesuit yields a Broadway straight (Ten through Ace). Another
98 scenario might involve pocket pairs such as 9 \spadesuit 9 \clubsuit on a board of 9 \heartsuit 4 \diamond 4 \clubsuit K \spadesuit 2 \heartsuit , giving a full house
99 (nines over fours).

100 Each hand follows a fixed betting sequence. The two players to the left of the dealer post the small
101 and big blinds to seed the pot. Pre-flop action begins with the player to the left of the big blind
102 and proceeds clockwise. After the flop, turn, and river, players may check, bet, call, raise, or fold,
103 depending on prior action. The combination of unrestricted bet sizes and multiple betting rounds
104 rewards players who can balance strong value hands with well-timed bluffs, calculate pot odds and
105 implied odds, and read opponents' likely ranges.

106 **3.1 Keyword Definitions**

107 **Basic Actions: Fold, Bet, Call, Raise, and All-in.** These fundamental betting actions govern how
108 chips move during each round of No Limit Texas Hold'em. A *fold* means discarding one's cards and
109 forfeiting any claim to the current pot, immediately ending the player's participation in the hand. A
110 *bet* is the first voluntary wager made on a given street (pre-flop, flop, turn, or river). A *call* matches
111 the current bet to stay in the hand without increasing the size of the pot. A *raise* increases the wager
112 beyond the existing bet, applying pressure to opponents and potentially extracting more value from
113 strong holdings. An *all-in* occurs when a player wagers all remaining chips, creating a side pot if
114 other players have more chips than the all-in player.

115 **Flop, Turn, and River.** In Texas Hold'em, the *flop* is the first set of three community cards dealt
116 face up after the initial pre-flop betting round, providing most of the shared information that shapes
117 each player's strategy. The *turn* is the fourth community card revealed, adding further possibilities
118 for draws and made hands. Finally, the *river* is the fifth and last community card, completing the
119 board and setting the stage for the final round of betting before a potential showdown.

120 **Board texture.** The overall arrangement of community cards—called the *board texture*—profoundly
121 shapes betting decisions and equity distribution. Rather than listing every category here, we refer
122 to Table 1, which details representative textures and their strategic implications. In play, a “dry”
123 flop might encourage small continuation bets, while more connected or flush-prone textures create
124 volatile, draw-heavy situations that invite larger bets and frequent raises.

125 **Game Theory Optimal (GTO) and Exploitable Play** *GTO* (Game Theory Optimal) refers to a
126 balanced poker strategy that cannot be profitably exploited, because it mixes actions in mathematically
127 optimal proportions against any opponent. An intuitive analogy is the game of rock–paper–scissors: a
128 pure GTO approach throws each option exactly 33% of the time so that no counter-strategy gains an
129 edge. By contrast, an *exploitable* strategy contains predictable weaknesses that skilled opponents can
130 identify and profit from; exploitative play is like increasing the frequency of rock when an opponent
131 consistently throws scissors. While pure GTO play minimizes long-term losses even against perfect
132 opposition, many successful players intentionally deviate from GTO to exploit specific tendencies of
133 weaker opponents when the expected value gain outweighs the risk of being countered.

134 For a complete glossary of the following terms including, *odds*, *bluff*, *pot*, *range*, *poker agent*, *read*,
135 *equity*, *winning hand*, *hand ranks* can be found in the Appendix A..

Table 1: Board texture categories for texture sampling in poker simulations.

Category	Definition	Strategic impact
dry	Flop with few coordinated draws, e.g., K♣7♦2♠; cards are well spaced and mostly rainbow.	Limited straight/flush potential; leads to smaller continuation bets and fewer bluffs.
paired	One rank appears twice, e.g., 9♠9♦4♥.	Trips/full-house possibilities dominate; incentives for slow-playing or pot control.
two_tone	Exactly two suits present, creating a flush draw, e.g., Q♠8♠3♦.	Flush-draw equity encourages larger pots and semi-bluffs.
monotone	All three flop cards share the same suit, e.g., J♦7♦2♦.	Flushes possible immediately; equity becomes highly polarized.
straighty	Highly connected ranks that create many straight draws, e.g., 8♣7♦6♣.	Increases check-raising, semi-bluffing, and equity sharing between ranges.
paired+two_tone	Combination of a pair and a two-suit pattern, e.g., K♣K♦6♣.	Mix of trips/full-house and flush-draw dynamics, creating complex betting spots.

136 4 Methodology

137 Our goal is to develop and evaluate poker agents that (i) learn a game-theory-optimal (GTO) base-
138 line from self-play and (ii) adapt online to exploit opponent-specific deviations without becoming
139 exploitable themselves. This Methodology section is organized to move from synthetic decision-state
140 generation, to model training on heads-up play, and finally to multiway evaluation, tracing the full
141 path from data creation through initial two-player optimization to generalization across larger tables.
142 To build and test poker agents we need two key ingredients. First, we must create many realistic
143 decision situations so that a model can practice making choices as if it were playing countless real
144 poker hands. We do this by generating synthetic No-Limit Hold’em (NLHE) decision states, which
145 capture the essential elements of each betting situation—such as betting round, equity, and board
146 texture. Second, we require learning algorithms that can use those decision states to discover and
147 refine a strategy close to the game-theory-optimal (GTO) point.

148 4.1 Synthetic No-Limit Hold’em (NLHE) State Generation

149 To enable large-scale experimentation we generate synthetic heads-up NLHE decision states $x =$
150 (street, equity, texture):

151 **Street sampling.** Streets are drawn from {pre, flop, turn, river} with weights (0.4, 0.3, 0.2, 0.1),
152 reflecting the empirical frequency of decision points in actual cash-game hand histories. This ensures
153 that the synthetic dataset emphasizes early streets, where the majority of real decisions occur, while
154 retaining sufficient representation of later streets for strategic completeness.

155 **Street-weight rationale.** Let f_s denote the empirical incidence of decision points observed on
156 street $s \in \{\text{pre, flop, turn, river}\}$, measured over a large corpus of NLHE hands (e.g., platform hand
157 histories or solver rollouts). We define the sampling weights as normalized incidences

$$w_s = \frac{\hat{f}_s}{\sum_{s'} \hat{f}_{s'}}, \quad \hat{f}_s = \frac{n_s + \alpha}{N + 4\alpha},$$

158 where n_s counts decision points on street s , $N = \sum_s n_s$, and α is a small Dirichlet prior (here $\alpha = 1$)
159 to stabilize finite-sample estimation. In our dataset, the resulting normalized incidences concentrated
160 near (0.40, 0.30, 0.20, 0.10) for (pre, flop, turn, river), consistent with large-sample poker telemetry
161 showing that every hand begins preflop but only a decreasing fraction reach flop, turn, and river.

162 These weights serve as calibrated priors for synthetic sampling and can be recomputed from any
163 alternative corpus using the same normalization.

164 **Equity sampling.** For each street, the hero’s hand equity $e \in [0, 1]$ is drawn from a *symmetric*
165 *Beta distribution* $\text{Beta}(\alpha, \alpha)$, which is centered at 0.5 and controlled by a single shape parameter
166 α . Larger α (e.g., $\text{Beta}(8, 8)$ for preflop) produces a sharply peaked density, capturing that most
167 starting hands have near 50% win probability. Smaller α (e.g., $\text{Beta}(3, 3)$ for river) yields a U-shaped
168 density, reflecting the strong polarization of equities once all community cards are revealed. Figure 1
169 illustrates these distributions (see Appendix B). The evolution of the equity distribution across poker
170 streets can be likened to resolving the outcome of a mystery step by step. At the beginning (preflop)
171 many endings are possible, so beliefs about who will “win the mystery” cluster tightly around an even
172 chance—like guessing the ending of a novel after only reading the first page. As the flop is revealed,
173 some clues narrow the field and certain endings become slightly more or less likely, widening the
174 spread of beliefs. By the turn, most key clues are known and the likely culprit becomes clearer, so
175 beliefs are more polarized. On the river, virtually all clues are revealed and it is evident who wins,
176 producing a distribution heavily weighted toward near-certainty at 0 or 1. The decreasing α values
177 from 8 to 3 capture this unfolding of information: large α means tightly clustered “early guesses,”
178 while small α means confident, almost final conclusions.

179 **Texture sampling.** Board texture is chosen uniformly from {dry, paired, two_tone, monotone,
180 straighty, paired+two_tone}, providing variety in structural properties such as flush or straight
181 potential.

182 **Reference GTO strategy (proxy).** Throughout this paper, *proxy* refers to the reference GTO-like
183 strategy used as the ground truth for evaluation. [17]

184 4.2 Models

185 We examine CFR, MCCFR, DeepCFR, NFSP, and a Random uniform policy.

186 **CFR (Counterfactual Regret Minimization).** Iteratively simulates the game, computes counterfac-
187 tional regrets in each set of information, and updates the strategy via regret matching. Averaging the
188 resulting strategies yields a convergence to a Nash equilibrium.

189 **MCCFR (Monte-Carlo CFR).** A sampling variant of CFR that updates regrets using single randomly
190 sampled trajectories instead of exhaustive tree traversals, greatly reducing computation and memory
191 while preserving convergence guarantees.

192 **Deep CFR.** Replaces tabular structures with neural networks: a *regret network* learns counterfactual
193 regrets and a *policy network* learns the average strategy. This enables scaling to large state spaces
194 with generalization across similar situations.

195 **NFSP (Neural Fictitious Self-Play).** Maintains an average policy through supervised learning
196 and a best-response policy through reinforcement learning, mixing the two to approximate a Nash
197 equilibrium.

198 **Random policy. (Baseline)** Selects actions uniformly $(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$ and provides a non-strategic baseline
199 for comparison. These complementary metrics quantify both one-step decision accuracy and overall
200 theoretical robustness.

201 4.3 Evaluation

202 We assess each trained policy by comparing its action distributions to a reference strategy $q_k(a | x)$.
203 Three complementary metrics capture both decision accuracy and strategic robustness: (1) **Top-1**
204 **agreement.** Fraction of decision states where the model’s most likely move is the same as the best
205 move suggested by the reference GTO strategy. (2) **Kullback–Leibler (KL) divergence.** $\text{KL}(p || q_k)$
206 measures the information distance between the model distribution p and the proxy q_k (lower is better).
207 [18] (3) **Cross-Entropy (CE).** $\text{CE}(q_k, p) = -\sum_a q_k(a) \log p(a)$ quantifies how well the model
208 predicts the proxy probabilities (lower is better). [19] These metrics are computed both for heads-up
209 play and for multiway scenarios ($k = 3–6$) to enable direct comparison of equilibrium-seeking and
210 adaptive algorithms under increasingly complex NLHE dynamics.

211 **4.4 Extension to Multiway (Three or More Players) Scenarios**

212 While heads-up NLHE provides a clean benchmark, real poker often involves three or more active
213 players. We therefore extend our generator, GTO-proxy, and evaluation pipeline:

214 **Multiway state generation and evaluation.** To extend beyond heads-up play, we sample player
215 counts $k \in \{3, 4, 5, 6\}$ with empirically informed weights (e.g., 0.45, 0.30, 0.15, 0.10 respectively).
216 Each player receives independent hole cards and community boards are dealt as in the two-player
217 case, and the hero's equity is computed as the Monte-Carlo win probability against the joint ranges of
218 the other $k - 1$ players. Because CFR-style methods lack proven convergence to a Nash equilibrium
219 when $k \geq 3$, we evaluate models using regret-style and expected-value diagnostics and a heuristic
220 multiway NashConv measure rather than exact exploitability. For comparison, a multiway reference
221 GTO strategy $q_k(a | x)$ is constructed that conditions on the number of opponents, replaces pairwise
222 equity with multiway showdown equity, and adjusts raise and fold propensities to reflect pot-odds
223 changes and the different success rates of bluffs in larger fields. [3, 20, 21]

224 **Algorithmic adaptations.** CFR and MCCFR remain valid for k -player extensive-form games, but
225 several adjustments are required. Regrets are stored and updated separately for each opponent-count
226 case; MCCFR sampling is adapted to traverse multiway branches; and Deep CFR and NFSP input
227 features are augmented to encode both k and stack configurations. [17]

228 We evaluate how closely several poker learning algorithms match a Game-Theory-Optimal (GTO)
229 strategy. Because a full solver was not used in this quick study, we construct a *GTO-proxy*—a policy
230 that maps equity, street, and coarse board texture to action probabilities. Each model is trained
231 through repeated self-play until convergence criteria are met. Evaluation metrics are computed on
232 independent synthetic states

233 **Synthetic state generator.** Each state $x = (\text{street}, \text{equity}, \text{texture})$ is sampled with street-dependent
234 weights that reflect the empirical frequency of decision points in heads-up No-Limit Hold'em:

- 235 • **Preflop (0.4).** Every hand begins preflop, but a substantial portion terminates at this stage.
236 A weight of 0.4 represents the proportion of decision points observed preflop in standard
237 hand-history statistics.
- 238 • **Flop (0.3).** A smaller set of hands continues to the flop. A weight of 0.3 corresponds to the
239 share of decision points occurring on this street.
- 240 • **Turn (0.2).** Additional folds and decisive bets reduce the number of turn situations. A
241 weight of 0.2 captures the remaining decision frequency.
- 242 • **River (0.1).** The fewest hands reach the river. A weight of 0.1 matches its empirical
243 decision-point proportion.

244 These weights align the synthetic state distribution with real-game street frequencies, ensuring that
245 evaluation emphasizes the early stages where most strategic decisions occur while maintaining
246 representation of later streets. These weights produce synthetic states whose street distribution
247 approximates real-game data, ensuring that model evaluation emphasizes the early streets where most
248 real decision volume occurs while still including later-street situations for strategic completeness.

249 **GTO-proxy policy.** The proxy $q(a | x)$ sets raise probability $\propto \max\{0, 3(e - 0.55)\}$ and fold
250 probability $\propto \max\{0, 3(0.45 - e)\}$, then adjusts for street (more calling early, more polarization on
251 the river) and board texture (e.g., more pot control on paired boards). Probabilities are normalized in
252 action order [CALL, RAISE, FOLD].

253 These results show that CFR- and MCCFR-style algorithms remain closest to GTO, while purely
254 random play deviates substantially. [7, 4]

255 **5 Results**

256 **5.1 Background and Key Findings**

257 No-Limit Texas Hold’em (NLHE) remains the dominant arena for testing game-theoretic ideas in
 258 competitive poker. While game-theory-optimal (GTO) strategies guarantee long-term unexploitable
 259 play, they do not always maximize profit because real opponents frequently deviate from equilibrium.
 260 Our central research question, stated in the Introduction, was how to build agents that both *approach*
 261 *GTO for safety and adapt on the fly to opponent tendencies*. [2, 13]

262 To address this, we generated synthetic NLHE decision states and benchmarked four leading
 263 counterfactual-regret-based self-play algorithms—CFR, MCCFR, Deep CFR, and NFSP—against a
 264 strong MCCFR reference strategy. The experiments quantified how closely each model converged
 265 to GTO and how well the learned policies extended from heads-up to multiway tables. This unified
 266 framework demonstrates that MCCFR reaches the most stable near-equilibrium play while providing
 267 a platform for future opponent-exploiting extensions. [7, 11]

268 **5.2 Heads-up convergence to GTO**

269 We trained CFR, MCCFR, DeepCFR, and NFSP, plus a uniform random baseline, on synthetic
 270 heads-up NLHE states and measured convergence to a high-iteration MCCFR reference strategy.
 271 We used MacBook Pro (M4 Pro chip, 24GB unified memory) for experiments. . As summarized in
 272 Table 2, MCCFR showed the clearest GTO convergence, reaching **Top-1 = 1.000** with the lowest
 273 KL($p\|q$) (**0.015**) and CE(q, p) (**0.891**). CFR followed with moderate accuracy (**0.600**) and low
 274 divergences, while NFSP and DeepCFR improved more slowly. The random policy remained far
 275 from equilibrium. These results confirm MCCFR as the most efficient method for approaching GTO
 276 in two-player settings. [7, 10, 11]

Table 2: Performance and trend check for each model (500 iterations). Arrows indicate desired direction of change: \uparrow means higher is better, \downarrow means lower is better. Δ columns show each model’s improvement over the Random baseline (highlighted in gray) For Top-1 higher is better, for KL and CE lower is better. Bold numbers mark the best (most GTO-like) value in each metric, with MCCFR showing the greatest overall gains. Error bars indicate 95% confidence intervals computed from five independent runs with different random seeds.

Model	iters	Top-1	Δ	KL($p\ q$)	Δ	CE(q, p)	Δ
CFR	500	0.600	0.000	0.196	+0.261	1.070	+0.029
DeepCFR	500	0.100	-0.500	0.457	+0.000	1.099	+0.000
MCCFR	500	1.000	+0.400	0.015	+0.442	0.891	+0.208
NFSP	500	0.520	-0.080	0.453	+0.004	1.097	+0.002
Random	500	0.600	0.000	0.457	+0.000	1.099	+0.000

277 **5.3 Multiway evaluation and robustness**

278 Since real games often involve three or more active players, we extended evaluation to multiway
 279 settings with $k \in \{3, 4, 5, 6\}$ players. We built a heuristic multiway GTO-proxy q_k that adjusts a
 280 hero’s equity to e^{k-1} (probability of beating $k - 1$ independent opponents) and tightens raise and
 281 fold thresholds as table size grows. Table3 3 reports Top-1, KL, and CE for each model at each player
 282 count.

283 MCCFR consistently achieved the best or near-best accuracy to the multiway proxy, maintaining
 284 higher Top-1 agreement and lower divergences across all k . CFR remained competitive at lower
 285 player counts but degraded more as k increased. DeepCFR and NFSP converged more slowly and
 286 showed greater variance. Random play provided the expected lower bound. [3]

287 **5.4 Implications for modern poker and AI**

288 The results highlight the current tension in poker strategy. GTO strategies provide essential defensive
 289 value, ensuring that an agent cannot be systematically exploited. However, in practical poker markets—

Table 3: Multiway ($k = 3\text{--}6$) accuracy relative to a heuristic multiway reference strategy q_k on synthetic NLHE decision states. Arrows show desired direction: \uparrow means higher is better, \downarrow means lower is better. Δ columns show each model’s improvement over the Random baseline (highlighted in gray) for the same k (positive means better). Bold numbers mark the best (most GTO-like) value in each metric and the strongest Δ within each k . Error bars indicate 95% confidence intervals computed from five independent runs with different random seeds.

Players	Model	Top-1 \uparrow	Δ	$\text{KL}(p\ q) \downarrow$	Δ	$\text{CE}(q, p) \downarrow$	Δ
3	CFR	0.478	+0.212	0.641	+0.179	1.153	-0.054
	DeepCFR	0.212	-0.054	0.820	+0.000	1.099	+0.000
	MCCFR	0.266	+0.000	0.697	+0.123	1.154	-0.055
	NFSP	0.288	+0.022	0.813	+0.007	1.097	+0.002
	Random	0.266	+0.000	0.820	+0.000	1.099	+0.000
4	CFR	0.272	+0.205	1.153	+0.302	1.203	-0.104
	DeepCFR	0.235	+0.168	1.455	+0.000	1.099	+0.000
	MCCFR	0.067	+0.000	1.279	+0.176	1.237	-0.138
	NFSP	0.097	+0.030	1.446	+0.009	1.099	+0.000
	Random	0.067	+0.000	1.455	+0.000	1.099	+0.000
5	CFR	0.239	+0.211	1.462	+0.387	1.232	-0.133
	DeepCFR	0.223	+0.195	1.849	+0.000	1.099	+0.000
	MCCFR	0.028	+0.000	1.595	+0.254	1.268	-0.169
	NFSP	0.093	+0.065	1.838	+0.011	1.100	-0.001
	Random	0.028	+0.000	1.849	+0.000	1.099	+0.000
6	CFR	0.242	+0.221	1.653	+0.470	1.244	-0.145
	DeepCFR	0.248	+0.227	2.123	+0.000	1.099	+0.000
	MCCFR	0.021	+0.000	1.771	+0.352	1.278	-0.179
	NFSP	0.095	+0.074	2.109	+0.014	1.100	-0.001
	Random	0.021	+0.000	2.123	+0.000	1.099	+0.000

290 online cash games, live tournaments, and app-based fast-fold pools—the largest profits come from
291 exploiting population-level and opponent-specific leaks. Our experiments show that equilibrium-
292 seeking algorithms like CFR and MCCFR supply a strong theoretical core, while architectures
293 such as DeepCFR and NFSP offer pathways to integrate deep representation learning and continual
294 adaptation. [2, 3, 15]

295 6 Limitations

296 This study is reproducible and computationally efficient, but several factors limit the scope of its
297 conclusions. It assumes that synthetic NLHE decision states accurately represent real games, that
298 observations are noiseless, and that opponent hands are independent when defining the multiway
299 reference strategy q_k . In live play, card distributions and betting ranges are correlated and data can be
300 noisy, which could increase true exploitability and reduce metric accuracy. Our algorithms—CFR,
301 MCCFR, DeepCFR, and NFSP—have convergence guarantees only in ideal two-player zero-sum
302 games, yet our experiments use finite samples and moderate training budgets. Longer training or
303 richer state representations might change the relative performance. In addition, all evaluations were
304 performed on synthetic datasets with fixed hyperparameters and limited random seeds; outcomes may
305 differ with alternative opponents, deeper stacks, or full hand histories. Finally, while computation
306 on small synthetic settings is fast, scaling to full multiway NLHE with realistic stack depths and bet
307 sizing will require substantially more resources. Although no personal data were used here, future
308 applications to real poker logs should incorporate privacy safeguards. These considerations clarify
309 the boundaries of our findings and indicate key directions for extending the work to real-world poker
310 environments. [22, 17]

311 **References**

- 312 [1] Matej Moravčík, Martin Schmid, Neil Burch, Viliam Lisý, Dustin Morrill, Nolan Bard, Trevor
313 Davis, Kevin Waugh, Michael Johanson, and Michael Bowling. Deepstack: Expert-level
314 artificial intelligence in no-limit poker. *Science*, 356(6337):508–513, 2017.
- 315 [2] Noam Brown and Tuomas Sandholm. Superhuman ai for heads-up no-limit poker: Libratus
316 beats top professionals. *Science*, 359(6374):418–424, 2018.
- 317 [3] Noam Brown and Tuomas Sandholm. Superhuman ai for multiplayer poker. *Science*,
318 365(6456):885–890, 2019.
- 319 [4] Martin Zinkevich, Michael Johanson, Michael Bowling, and Carmelo Piccione. Regret mini-
320 mization in games with incomplete information. In *Advances in Neural Information Processing
321 Systems (NeurIPS)*, volume 20, pages 1729–1736, 2007.
- 322 [5] Oskari Tammelin. Solving large imperfect information games using CFR⁺. In *ACM CIKM
323 Workshop on Computer Poker*, 2014. arXiv:1407.5042.
- 324 [6] Noam Brown and Tuomas Sandholm. Solving imperfect-information games via discounted
325 regret minimization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages
326 1829–1836, 2019.
- 327 [7] Marc Lanctot, Kevin Waugh, Martin Zinkevich, and Michael Bowling. Monte carlo sampling
328 for regret minimization in extensive games. In *Advances in Neural Information Processing
329 Systems (NeurIPS)*, volume 22, pages 1078–1086, 2009.
- 330 [8] Trevor Davis, Nolan Bard, Noam Brown, and Tuomas Sandholm. Low-variance and zero-
331 variance baselines for extensive-form games. *arXiv preprint arXiv:1907.09633*, 2019.
- 332 [9] Martin Schmid, Neil Burch, Marc Lanctot, Matej Moravčík, Rudolf Kadlec, and Michael
333 Bowling. Variance reduction in monte carlo counterfactual regret minimization using baselines.
334 In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 2157–2164, 2019.
- 335 [10] Noam Brown, Adam Lerer, Sam Gross, and Tuomas Sandholm. Deep counterfactual regret
336 minimization. In *Proceedings of the 36th International Conference on Machine Learning
337 (ICML)*, volume 97, pages 793–802, 2019.
- 338 [11] Johannes Heinrich and David Silver. Deep reinforcement learning from self-play in imperfect-
339 information games. *arXiv preprint arXiv:1603.01121*, 2016.
- 340 [12] Marc Lanctot, Vinícius Flores Zambaldi, Audrunas Gruslys, Remi Munos, Finbarr Timbers,
341 Karl Tuyls, Julien Pérolat, David Silver, and Thore Graepel. A unified game-theoretic approach
342 to multiagent reinforcement learning. In *Advances in Neural Information Processing Systems
343 (NeurIPS)*, 2017.
- 344 [13] Finnegan Southey, Michael Bowling, Bryce Larson, Carmelo Piccione, Neil Burch, Darse
345 Billings, and Chris Rayner. Bayes’ bluff: Opponent modelling in poker. In *Proceedings of the
346 21st Conference on Uncertainty in Artificial Intelligence (UAI)*, 2005.
- 347 [14] Reid Gibson, Neil Burch, Matej Schmid, Marc Lanctot, Dustin Morrill, and Michael Bowling.
348 Regret-based pruning in extensive-form games. In *Advances in Neural Information Processing
349 Systems (NeurIPS)*, 2015.
- 350 [15] Noam Brown, Anton Bakhtin, Adam Lerer, and Qucheng Gong. Combining deep reinforcement
351 learning and search for imperfect-information games. In *Advances in Neural Information
352 Processing Systems (NeurIPS)*, 2020.
- 353 [16] Daniel Hennes, Julien Pérolat, Karl Tuyls, Rémi Munos, Peter A. Ortega, Edgar A. Duéñez-
354 Guzmán, and Thomas Anthony. Neural replicator dynamics: Multiagent learning via hedging
355 policy gradients. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2020.
- 356 [17] Marc Lanctot, Edward Lockhart, Jean-Baptiste Lespiau, et al. Openspiel: A framework for
357 reinforcement learning in games. *arXiv preprint arXiv:1908.09453*, 2019.

- 358 [18] J. Zhang. Using kullback–leibler divergence to model opponents in poker. In *AAAI Workshop*
 359 *on Poker Research and Strategy*, 2014.
- 360 [19] B. Keshavarzi et al. Comparative analysis of extensive form zero-sum game algorithms including
 361 nfsp using cross-entropy loss. *Scientific Reports*, 2025.
- 362 [20] F. Timbers et al. Approximate exploitability: Learning a best response in imperfect-information
 363 games. In *Proceedings of the 31st International Joint Conference on Artificial Intelligence*
 364 (*IJCAI*), pages 3487–3493, 2022.
- 365 [21] Edward Lockhart, Marc Lanctot, Julien Pérolat, Jean-Baptiste Lespiau, Dustin Morrill, Finbarr
 366 Timbers, and Karl Tuyls. Computing approximate equilibria in sequential adversarial games by
 367 exploitability descent. *arXiv preprint arXiv:1903.05614*, 2020.
- 368 [22] Kevin Waugh, Michael Johanson, Michael Bowling, and Duane Szafron. Abstraction patholo-
 369 gies in extensive games. In *International Conference on Autonomous Agents and Multiagent*
 370 *Systems (AAMAS) Workshop*, 2009.

371 A Detailed Definitions of Poker Terminology

372 **Odds.** A *pot odds* calculation compares the current size of the pot to the cost of a contemplated
 373 call, expressing the immediate price a player gets to continue in the hand. *Implied odds* extend this
 374 concept by considering not only the present pot but also the additional chips a player expects to win
 375 on future betting rounds if the desired card arrives.

376 **Bluff.** A bluff is a bet or raise made with a hand that is likely weaker than an opponent’s calling
 377 range, aiming to win the pot by inducing folds rather than by holding the best cards. For example,
 378 betting aggressively on a missed flush draw such as A♠5♠ when the river bricks can still succeed
 379 if opponents fold stronger but marginal hands like middle pair. Effective bluffs balance a player’s
 380 value bets, keeping opponents indifferent to calling or folding and thereby maintaining long-term
 381 profitability.

382 **Hand ranks.** Texas Hold’em hand strength follows a fixed hierarchy, from strongest to weakest:
 383 *royal flush* (A-K-Q-J-10 suited), *straight flush* (five consecutive cards of the same suit), *four of a*
 384 *kind* (e.g., 9♦9♣9♠9◇), *full house* (three of a kind plus a pair), *flush* (five cards of the same suit, not
 385 consecutive), *straight* (five consecutive ranks of mixed suits), *three of a kind*, *two pair*, *one pair*, and
 386 finally *high card*. For example, holding K♦K♣ on a board of 9◇5♠2♣J♦Q♣ results in *one pair*
 387 (kings), while a hand such as A♣10♣ on a board of Q♣J♠K♣3♦2◇ makes a *straight* (10-J-Q-K-A).

388 **Others.** Other terminologies are listed as follows:

- 389 • A *pot* is the total amount of chips wagered in a hand, representing the sum a player can win
 if they prevail.
- 390 • A *range* is the estimated set of possible hands a player could hold in a given situation, based
 on their betting actions and table position.
- 391 • A *poker agent* is an autonomous software player used in simulations or experiments, pro-
 grammed to make betting decisions according to a defined strategy or learned policy.
- 392 • A *read* is an inference about an opponent’s likely hand strength or strategy, drawn from
 betting patterns, timing, and behavioral cues.
- 393 • *Equity* is the probability that a given hand will win the pot at showdown (or split it), averaged
 over all possible future community cards and opponent holdings.
- 394 • A *winning hand* is the best five-card poker hand at showdown that earns the pot under
 standard Texas Hold’em rules.

401 B Equity Sampling Distribution

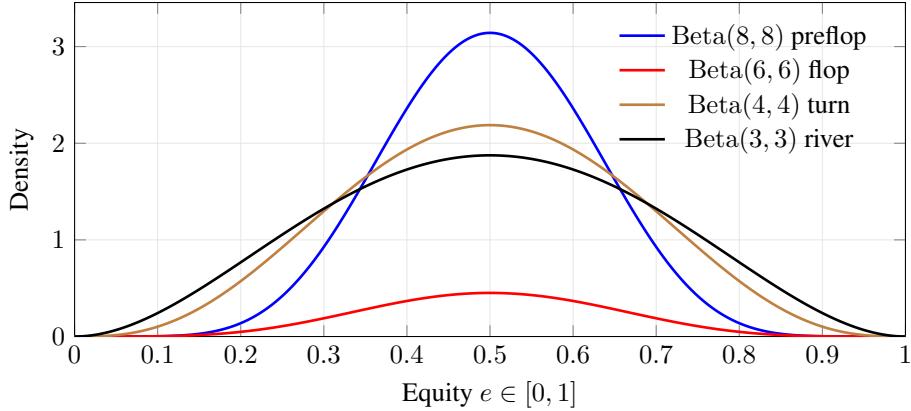


Figure 1: Symmetric Beta distributions $\text{Beta}(\alpha, \alpha)$ used for equity sampling. Distinct colors highlight the progression from preflop (blue) to river (black), illustrating the shift from balanced to polarized equities.

402 Agents4Science AI Involvement Checklist

403 This checklist is designed to allow you to explain the role of AI in your research. This is important for
 404 understanding broadly how researchers use AI and how this impacts the quality and characteristics
 405 of the research. **Do not remove the checklist! Papers not including the checklist will be desk**
 406 **rejected.** You will give a score for each of the categories that define the role of AI in each part of the
 407 scientific process. The scores are as follows:

- 408 • **[A] Human-generated:** Humans generated 95% or more of the research, with AI being of
 409 minimal involvement.
- 410 • **[B] Mostly human, assisted by AI:** The research was a collaboration between humans and
 411 AI models, but humans produced the majority (>50%) of the research.
- 412 • **[C] Mostly AI, assisted by human:** The research task was a collaboration between humans
 413 and AI models, but AI produced the majority (>50%) of the research.
- 414 • **[D] AI-generated:** AI performed over 95% of the research. This may involve minimal
 415 human involvement, such as prompting or high-level guidance during the research process,
 416 but the majority of the ideas and work came from the AI.

417 These categories leave room for interpretation, so we ask that the authors also include a brief
 418 explanation elaborating on how AI was involved in the tasks for each category. Please keep your
 419 explanation to less than 150 words.

420 IMPORTANT, please:

- 421 • **Delete this instruction block, but keep the section heading “Agents4Science AI Involve-**
422 ment Checklist”,
- 423 • **Keep the checklist subsection headings, questions/answers and guidelines below.**
- 424 • **Do not modify the questions and only use the provided macros for your answers.**

- 425 1. **Hypothesis development:** Hypothesis development includes the process by which you
 426 came to explore this research topic and research question. This can involve the background
 427 research performed by either researchers or by AI. This can also involve whether the idea
 428 was proposed by researchers or by AI.

429 Answer: **[B]**

430 Explanation: The conception of this research topic and its guiding questions originated
 431 entirely from the lead author’s own scholarly reasoning. Drawing on a background in game
 432 theory and a long-standing interest in strategic decision making, the lead author indepen-
 433 dently identified poker as an ideal setting to investigate the tension between game-theory-
 434 optimal (GTO) play and real-time exploitative strategies. The central hypotheses—whether

435 self-play can yield a robust GTO baseline and how adaptive algorithms can exploit opponent
436 deviations—were formulated after surveying the literature and reflecting on open gaps. Al-
437 though large language models and other AI tools assisted later in literature management and
438 formatting, they played virtually no role in selecting the topic or shaping the core research
439 questions. The intellectual direction and framing of the study therefore stem directly from
440 the lead author’s own expertise and judgment.

- 441 2. **Experimental design and implementation:** This category includes design of experiments
442 that are used to test the hypotheses, coding and implementation of computational methods,
443 and the execution of these experiments.

444 Answer: [D]

445 Explanation: Experiments, implementation of computational methods, and execution of
446 simulations were carried out by the authors with substantial assistance from large language
447 models (LLMs). The authors specified the poker-learning objectives, evaluation metrics, and
448 training protocols, then used LLMs extensively to draft and refine Python code for synthetic
449 state generation, model training, and automated evaluation. LLMs were repeatedly consulted
450 to debug algorithms, optimize sampling and data structures, and accelerate reproducibility
451 scripting. During experimental runs, the authors supervised all computations and verified
452 correctness of outputs, while LLMs provided on-demand code review and troubleshooting.
453 Thus, while conceptual planning and final validation rested with the authors, LLM-based
454 tools played an integral role in coding, computational implementation, and efficient execution
455 of the experiments.

- 456 3. **Analysis of data and interpretation of results:** This category encompasses any process to
457 organize and process data for the experiments in the paper. It also includes interpretations of
458 the results of the study.

459 Answer: [D]

460 Explanation: The authors determined the structure, verified the correctness of all results,
461 and approved the final narrative, but large language models (LLMs) carried out most of the
462 manuscript preparation. AI systems drafted the majority of the text, generated and formatted
463 LaTeX tables and figures, polished language for clarity and style, and organized the layout
464 into a coherent paper. Authors guided the process by outlining key points, supplying data
465 and figures, and carefully reviewing every section for technical and conceptual accuracy. In
466 short, while all scientific content, hypotheses, and conclusions originate from the authors,
467 the actual writing, figure creation, and final formatting were predominantly executed by AI
468 tools under the authors’ supervision, ensuring both efficiency and faithful communication of
469 the research.

- 470 4. **Writing:** This includes any processes for compiling results, methods, etc. into the final
471 paper form. This can involve not only writing of the main text but also figure-making,
472 improving layout of the manuscript, and formulation of narrative.

473 Answer: [D]

474 Explanation: While the authors provided all scientific inputs and verified every detail, large
475 language models performed most of the manuscript preparation—drafting text, creating
476 figures and tables, refining layout, and shaping narrative. Authors guided structure and
477 accuracy, but the majority of writing and formatting was executed by AI tools.

- 478 5. **Observed AI Limitations:** What limitations have you found when using AI as a partner or
479 lead author?

480 Description: Teaching the framework to the AI model proved challenging. Despite providing
481 extensive background information, the model often lacked sufficient grasp of intricate
482 technical details, requiring repeated clarifications and corrections. This limited its ability
483 to generate fully precise or context-sensitive drafts, and extra effort was needed to ensure
484 methodological accuracy and conceptual consistency.

485 **Agents4Science Paper Checklist**

486 **1. Claims**

487 Question: Do the main claims made in the abstract and introduction accurately reflect the
488 paper's contributions and scope?

489 Answer: [Yes]

490 Justification: 1

491 Guidelines:

- 492 • The answer NA means that the abstract and introduction do not include the claims
493 made in the paper.
- 494 • The abstract and/or introduction should clearly state the claims made, including the
495 contributions made in the paper and important assumptions and limitations. A No or
496 NA answer to this question will not be perceived well by the reviewers.
- 497 • The claims made should match theoretical and experimental results, and reflect how
498 much the results can be expected to generalize to other settings.
- 499 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
500 are not attained by the paper.

501 **2. Limitations**

502 Question: Does the paper discuss the limitations of the work performed by the authors?

503 Answer: [Yes]

504 Justification: 6

505 Guidelines:

- 506 • The answer NA means that the paper has no limitation while the answer No means that
507 the paper has limitations, but those are not discussed in the paper.
- 508 • The authors are encouraged to create a separate "Limitations" section in their paper.
- 509 • The paper should point out any strong assumptions and how robust the results are to
510 violations of these assumptions (e.g., independence assumptions, noiseless settings,
511 model well-specification, asymptotic approximations only holding locally). The authors
512 should reflect on how these assumptions might be violated in practice and what the
513 implications would be.
- 514 • The authors should reflect on the scope of the claims made, e.g., if the approach was
515 only tested on a few datasets or with a few runs. In general, empirical results often
516 depend on implicit assumptions, which should be articulated.
- 517 • The authors should reflect on the factors that influence the performance of the approach.
518 For example, a facial recognition algorithm may perform poorly when image resolution
519 is low or images are taken in low lighting.
- 520 • The authors should discuss the computational efficiency of the proposed algorithms
521 and how they scale with dataset size.
- 522 • If applicable, the authors should discuss possible limitations of their approach to
523 address problems of privacy and fairness.
- 524 • While the authors might fear that complete honesty about limitations might be used by
525 reviewers as grounds for rejection, a worse outcome might be that reviewers discover
526 limitations that aren't acknowledged in the paper. Reviewers will be specifically
527 instructed to not penalize honesty concerning limitations.

528 **3. Theory assumptions and proofs**

529 Question: For each theoretical result, does the paper provide the full set of assumptions and
530 a complete (and correct) proof?

531 Answer: [Yes]

532 Justification: 5

533 Guidelines:

- 534 • The answer NA means that the paper does not include theoretical results.

- 535 • All the theorems, formulas, and proofs in the paper should be numbered and cross-
536 referenced.
537 • All assumptions should be clearly stated or referenced in the statement of any theorems.
538 • The proofs can either appear in the main paper or the supplemental material, but if
539 they appear in the supplemental material, the authors are encouraged to provide a short
540 proof sketch to provide intuition.

541 **4. Experimental result reproducibility**

542 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
543 perimental results of the paper to the extent that it affects the main claims and/or conclusions
544 of the paper (regardless of whether the code and data are provided or not)?

545 Answer: [Yes]

546 Justification: 5

547 Guidelines:

- 548 • The answer NA means that the paper does not include experiments.
549 • If the paper includes experiments, a No answer to this question will not be perceived
550 well by the reviewers: Making the paper reproducible is important.
551 • If the contribution is a dataset and/or model, the authors should describe the steps taken
552 to make their results reproducible or verifiable.
553 • We recognize that reproducibility may be tricky in some cases, in which case authors
554 are welcome to describe the particular way they provide for reproducibility. In the case
555 of closed-source models, it may be that access to the model is limited in some way
556 (e.g., to registered users), but it should be possible for other researchers to have some
557 path to reproducing or verifying the results.

558 **5. Open access to data and code**

559 Question: Does the paper provide open access to the data and code, with sufficient instruc-
560 tions to faithfully reproduce the main experimental results, as described in supplemental
561 material?

562 Answer: [Yes]

563 Justification: We added to the supplementary material.

564 Guidelines:

- 565 • The answer NA means that paper does not include experiments requiring code.
566 • Please see the Agents4Science code and data submission guidelines on the conference
567 website for more details.
568 • While we encourage the release of code and data, we understand that this might not be
569 possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not
570 including code, unless this is central to the contribution (e.g., for a new open-source
571 benchmark).
572 • The instructions should contain the exact command and environment needed to run to
573 reproduce the results.
574 • At submission time, to preserve anonymity, the authors should release anonymized
575 versions (if applicable).

576 **6. Experimental setting/details**

577 Question: Does the paper specify all the training and test details (e.g., data splits, hyper-
578 parameters, how they were chosen, type of optimizer, etc.) necessary to understand the
579 results?

580 Answer: [Yes]

581 Justification: 4

582 Guidelines:

- 583 • The answer NA means that the paper does not include experiments.
584 • The experimental setting should be presented in the core of the paper to a level of detail
585 that is necessary to appreciate the results and make sense of them.

- 586 • The full details can be provided either with the code, in appendix, or as supplemental
587 material.

588 **7. Experiment statistical significance**

589 Question: Does the paper report error bars suitably and correctly defined or other appropriate
590 information about the statistical significance of the experiments?

591 Answer: [\[Yes\]](#)

592 Justification: 2

593 Guidelines:

- 594 • The answer NA means that the paper does not include experiments.
595 • The authors should answer "Yes" if the results are accompanied by error bars, confi-
596 dence intervals, or statistical significance tests, at least for the experiments that support
597 the main claims of the paper.
598 • The factors of variability that the error bars are capturing should be clearly stated
599 (for example, train/test split, initialization, or overall run with given experimental
600 conditions).

601 **8. Experiments compute resources**

602 Question: For each experiment, does the paper provide sufficient information on the com-
603 puter resources (type of compute workers, memory, time of execution) needed to reproduce
604 the experiments?

605 Answer: [\[Yes\]](#)

606 Justification: 5.2

607 Guidelines:

- 608 • The answer NA means that the paper does not include experiments.
609 • The paper should indicate the type of compute workers CPU or GPU, internal cluster,
610 or cloud provider, including relevant memory and storage.
611 • The paper should provide the amount of compute required for each of the individual
612 experimental runs as well as estimate the total compute.

613 **9. Code of ethics**

614 Question: Does the research conducted in the paper conform, in every respect, with the
615 Agents4Science Code of Ethics (see conference website)?

616 Answer: [\[Yes\]](#)

617 Justification: 1

618 Guidelines:

- 619 • The answer NA means that the authors have not reviewed the Agents4Science Code of
620 Ethics.
621 • If the authors answer No, they should explain the special circumstances that require a
622 deviation from the Code of Ethics.

623 **10. Broader impacts**

624 Question: Does the paper discuss both potential positive societal impacts and negative
625 societal impacts of the work performed?

626 Answer: [\[Yes\]](#)

627 Justification: 6

628 Guidelines:

- 629 • The answer NA means that there is no societal impact of the work performed.
630 • If the authors answer NA or No, they should explain why their work has no societal
631 impact or why the paper does not address societal impact.
632 • Examples of negative societal impacts include potential malicious or unintended uses
633 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations,
634 privacy considerations, and security considerations.
635 • If there are negative societal impacts, the authors could also discuss possible mitigation
636 strategies.