

---

# Analysis of Mock Conversations Across Large Language Models

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 The rapid advancement of large language models (LLMs) has enabled increasingly sophisticated conversational agents, and systematic comparisons of their  
2 conversational behaviors are of great importance. In this study, we generated  
3 mock conversations between two people using four LLMs—ChatGPT Free version  
4 without account, Gemini-2.0-flash, GPT-5 thinking model, and Claude Opus  
5 4.1—prompted to produce 30-turn interactions each. We quantitatively analyzed  
6 multiple conversation-level features, including structural metrics (e.g., number of  
7 turns, utterance length), lexical and linguistic properties (e.g., type-token ratio,  
8 noun/verb ratios, lexical alignment), sentiment and emotion, repetition and novelty,  
9 question-response patterns, speaker balance, and linguistic complexity measured  
10 via perplexity. Statistical tests (Kruskal-Wallis), feature importance analyses using  
11 random forests, and dimensionality reduction (PCA) were employed to identify  
12 discriminative features and uncover patterns across models. Results revealed that  
13 GPT-5 exhibited high novelty, lexical diversity, and complexity but shorter utterances,  
14 whereas ChatGPT Free produced longer, more positive utterances with higher question rates.  
15 Claude Opus 4.1 generated the longest conversations with balanced linguistic profiles,  
16 and Gemini-2.0-flash was generally intermediate. Our work provides a multi-dimensional understanding of AI conversational behavior  
17 within single-agent interactions. This can offer insights into model selection,  
18 fine-tuning, and the design of future human-AI dialogue systems.

21 

## 1 Introduction

22 The rapid evolution of conversational AI agents has significantly impacted various domains. Recent  
23 advancements have led to the development of different models, each exhibiting unique conversational  
24 characteristics. Understanding these differences is crucial for selecting the appropriate model for  
25 specific applications. There are many studies dealing with AI and conversation. For example,  
26 Ebubechukwu, et al. (2024) generated artificial conversations using GPT-4o and compared human  
27 and GPT-4 evaluations of predefined key performance indicators. Xu, et al. (2024) discussed LLMs  
28 in the context of second language learning and evaluated their dialogues based on response success  
29 rate, suggestion success rate, and session success rate. In addition to these, there are a large number of  
30 evaluation methods for LLMs in multi-turn conversational settings (Guan, et al. (2025)). Rzadeczka,  
31 et al. (2025) discussed conversational AI in mental health interventions, especially in the context  
32 of cognitive biases. Anderson, et al. (2025) discussed LLM-associated words commonly used by  
33 LLMs and pointed out that they might be shifting our language patterns. These studies from multiple  
34 perspectives and contexts show a growing interest in AI conversation patterns.

35 To better understand conversations generated by LLMs, in this study, we systematically evaluate and  
36 compare the conversational characteristics of four LLMs; ChatGPT Free version, Gemini-2.0-flash,  
37 GPT-5 thinking model, and Claude Opus 4.1. Our analysis encompasses various aspects, including

38 conversation structure, lexical and linguistic features, sentiment and emotion, repetition and novelty,  
39 question and interaction patterns, speaker balance, and complexity. By employing statistical tests,  
40 feature importance analysis, and dimensionality reduction techniques, we provide a comprehensive  
41 overview of how these models differ in their conversational behaviors. The findings from this study  
42 offer valuable insights for developers, researchers, and practitioners seeking to understand the nuanced  
43 differences between contemporary conversational AI agents. By highlighting these distinctions, we  
44 aim to inform the selection and optimization of AI models for specific tasks, ultimately enhancing  
45 user experience and application efficacy.

## 46 2 Methods

### 47 2.1 Conversation Generation

48 We generated a synthetic dataset of multi-turn conversations by prompting four large language models  
49 (LLMs): ChatGPT (Free version without account), Gemini-2.0-flash, GPT-5 (thinking model), Claude  
50 Opus 4.1

51 Each model was prompted to produce five independent conversations, with 60 utterances per conver-  
52 sation (30 per speaker). We used the following controlled prompt to standardize conversation length  
53 and format:

54 Please make a mock conversation between two people A and B. The output format  
55 should be as follows: conversation = [ ("A", "XXX"), ("B", "YYY"), ("A", "ZZZ"),  
56 ("B", "WWW"), ] Please make A and B speak 30 times each.

### 57 2.2 Feature Extraction

58 We computed a rich set of linguistic, semantic, and interactional features for each conversation.

59 Libraries and models used in this study was as follows:

- 60 • spaCy (en\_core\_web\_sm): tokenization, POS tagging, sentence segmentation
- 61 • SentenceTransformers (all-MiniLM-L6-v2): utterance embeddings, cosine similarity for  
62 topical consistency
- 63 • TextBlob: sentiment polarity
- 64 • HuggingFace Transformers (GPT-2): token-level perplexity for measuring predictive diffi-  
65 culty

66 For each conversation, we extracted:

67 **Basic statistics:** Number of turns (num\_turns), mean utterance length (avg\_len), type-token ratio  
68 (ttr), per-speaker word balance

69 **Syntactic features:** Question rate (fraction of utterances ending with "?"), noun-to-token ratios,  
70 verb-to-token ratios, average sentence length

71 **Semantic consistency:** Topical consistency (mean cosine similarity between consecutive utter-  
72 ances), embedding variance (dispersion of utterance embeddings), novelty (complement of topical  
73 consistency, reflecting informational change)

74 **Discourse dynamics:** Self-repetition rates, partner word reuse rates, lexical alignment between  
75 speakers (shared vocabulary proportion), lexical entropy (Shannon entropy over word distribution),  
76 lexical convergence over conversation progression

77 **Pragmatic and affective features** Mean of sentiment polarity, variance of sentiment polarity,  
78 sentiment shift between first and second halves, pattern rate (frequency of common conversational  
79 clichés), questionresponse matching rate (proportion of questions followed by an answer), pronoun  
80 balance (I vs. you ratio)

81 **Complexity:** Average utterance perplexity under GPT-2 (an estimate of predictability)

82 This feature set captures a diverse set of conversation properties, from lexical richness to semantic  
83 coherence and pragmatic interaction quality.

84 **2.3 Statistical Analysis and Visualization**

85 **Exploratory Visualization:** To compare distributions across agents, we produced box plots of each  
86 feature grouped by agent, feature correlation heatmap (Pearson correlations across all features), and  
87 Principal Component Analysis (PCA) with two components, to visualize overall separation between  
88 agents in a reduced feature space. These visualizations allowed qualitative assessment of whether  
89 certain features systematically varied between LLMs.

90 **Inferential Statistics:** For each feature, we performed a KruskalWallis H-test (a non-parametric test  
91 appropriate for small samples) to assess whether feature distributions differed significantly across  
92 agents. If a feature contained fewer than two unique values, it was excluded from testing. Test  
93 statistics and p-values were recorded and sorted by significance.

94 **Feature Importance Analysis:** To identify which features most strongly discriminated between  
95 agents, we trained a Random Forest classifier (scikit-learn, `random_state=0`) using standardized  
96 feature vectors (StandardScaler). Feature importance scores were extracted and visualized as a ranked  
97 bar plot.

98 **2.4 Computational Resource**

99 All analyses were implemented in Python 3.10. The analysis was conducted on a server equipped  
100 with an Intel Xeon Gold 6130 CPU 2.10GHz, three NVIDIA Quadro GV100 GPUs, and 256 GB of  
101 RAM. Also, we placed our code on Google Colaboratory (see Reproducibility Statement for the link)  
102 for easier reproducibility. We were able to extract features within 5 minutes on Google Colaboratory  
103 without GPU. All the other analyses were almost immediate.

104 **3 Results**

105 **3.1 Patterns Observed from Exploratory Visualization**

106 We generated box plots of each feature grouped by agent (Figure 1).

107 **Conversation Structure:** Analysis of conversation length and utterance properties revealed sys-  
108 tematic differences across agents. Claude Opus 4.1 produced the longest conversations in terms of  
109 number of turns, while ChatGPT Free exhibited the widest variability. In contrast, average utterance  
110 length and average sentence length were highest for ChatGPT Free, whereas GPT-5 thinking model  
111 tended to produce shorter utterances and sentences. These findings suggest that ChatGPT Free  
112 emphasizes longer, more elaborate responses, while GPT-5 favors conciseness, and Claude maintains  
113 extended but balanced interactions.

114 **Lexical and Linguistic Features:** Lexical diversity and syntactic patterns differed notably between  
115 agents. GPT-5 thinking model exhibited the highest typetoken ratio (TTR) and novelty, indicating  
116 more varied and creative language, along with a higher noun ratio but lower verb ratio, suggesting  
117 noun-heavy phrasing. In contrast, ChatGPT Free and Claude Opus 4.1 displayed more balanced  
118 syntactic distributions. Claude Opus 4.1 also showed the highest lexical entropy, reflecting broader  
119 word usage, while lexical alignment and convergence were elevated for ChatGPT Free and Claude,  
120 indicating stronger mirroring of conversational partners.

121 **Sentiment and Emotional Patterns:** Sentiment analysis revealed distinct affective profiles. ChatGPT  
122 Free and Claude Opus 4.1 tended toward more positive sentiment, whereas GPT-5 thinking model  
123 had lower average sentiment and demonstrated more negative sentiment shifts over the conversation.  
124 Variation in sentiment was slightly lower for GPT-5, consistent with its concise style.

125 **Repetition and Novelty:** Measures of self-repetition and partner word reuse highlighted differences  
126 in conversational creativity. GPT-5 thinking model had much lower self-repetition and partner reuse  
127 rates, suggesting less recycling of prior phrases. This agent also scored highest in novelty, consistent  
128 with its more diverse and unpredictable language production. ChatGPT Free and Claude Opus 4.1  
129 were moderate in both repetition and novelty.

130 **Questions and Interaction:** Question-asking behavior varied by agent. ChatGPT Free asked  
131 more questions on average, whereas all agents maintained a questionresponse rate of 1.0, ensuring  
132 that every question was answered. Pattern usage was negligible across agents, with only GPT-5

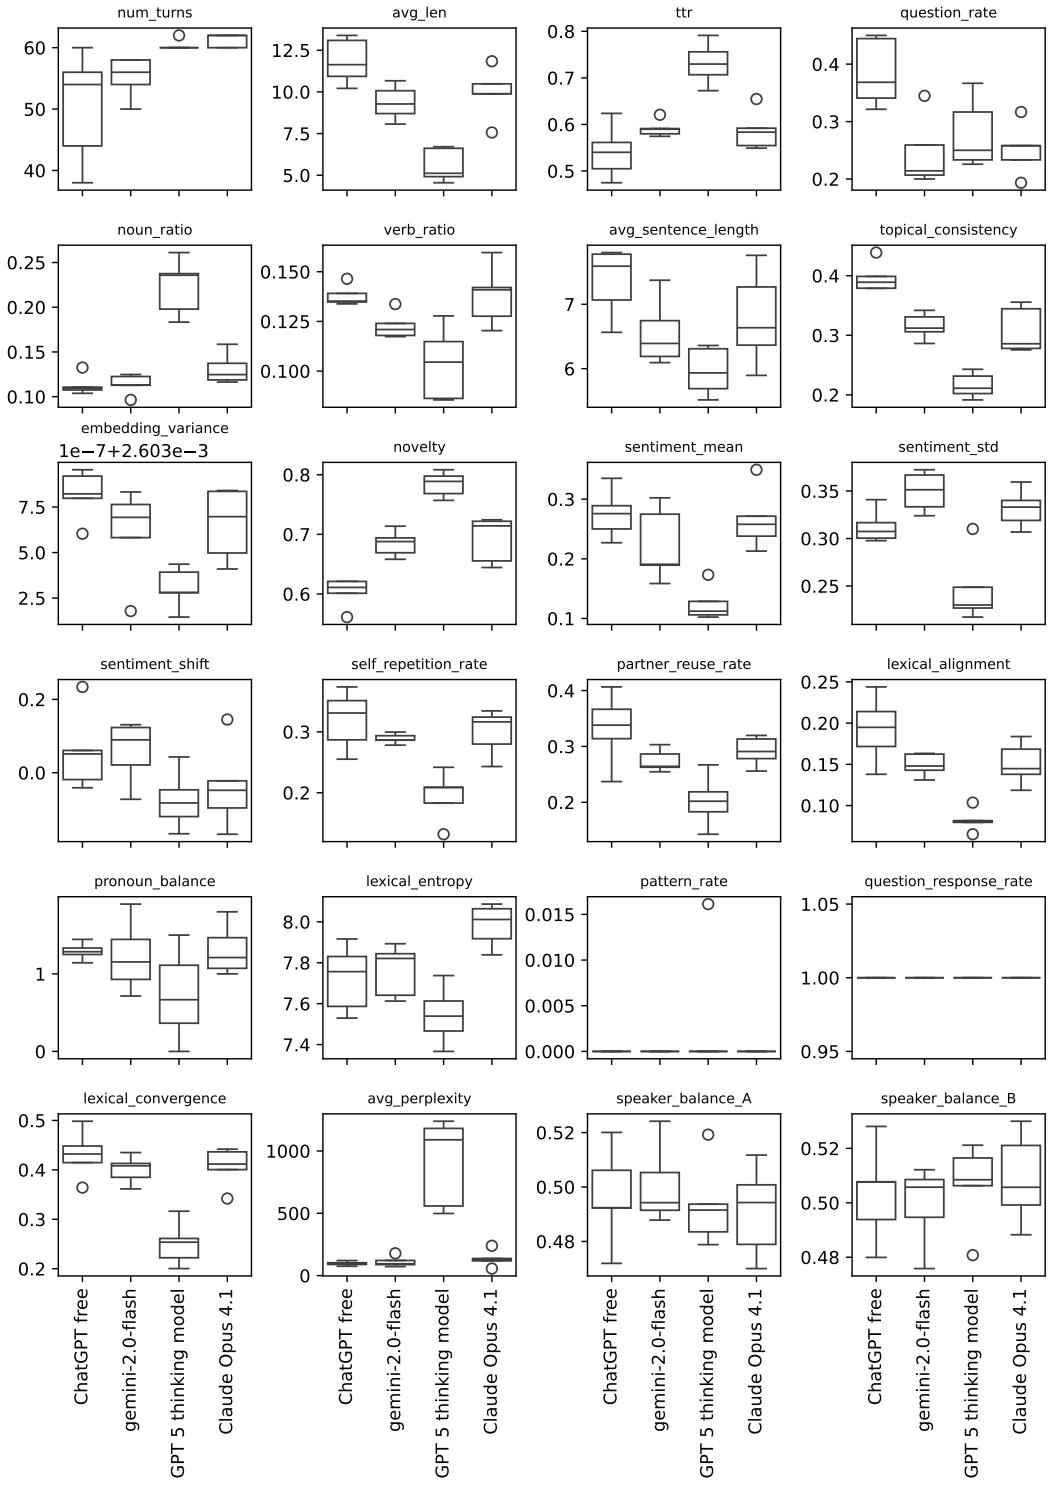


Figure 1: Box plots of each feature grouped by agent. We can qualitatively see the differences across agents.

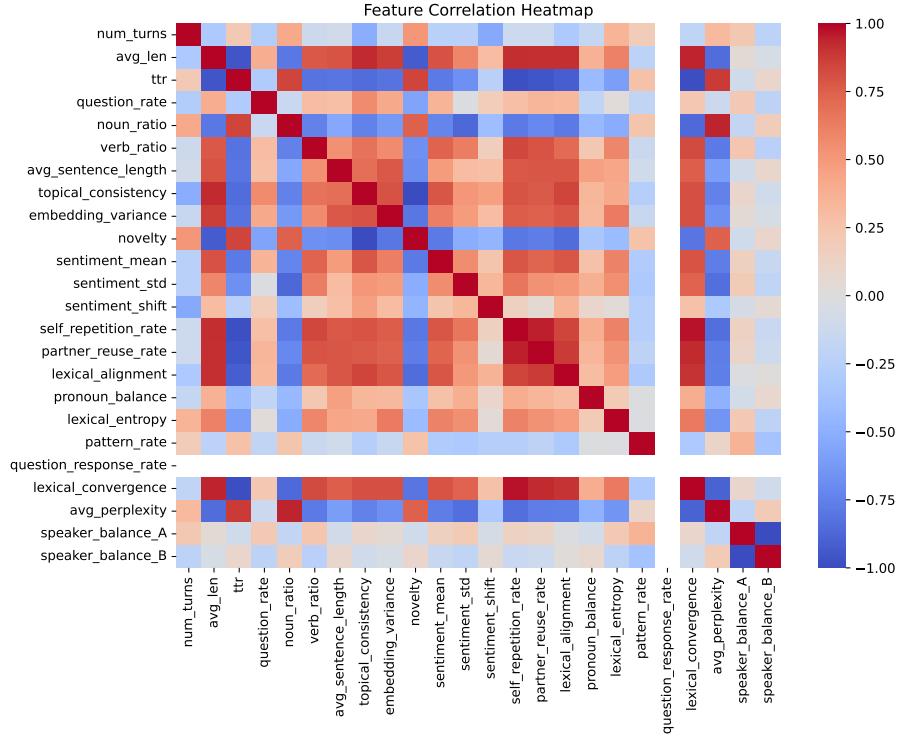


Figure 2: A heatmap of feature correlations. Several groups of highly correlated variables were identified.

133 showing a single outlier. Speaker balance remained around 0.5 for all agents, confirming that both  
 134 participants contributed equally. GPT-5 exhibited a slightly lower pronoun balance, reflecting reduced  
 135 self-referencing.

136 **Complexity:** Average perplexity analysis highlighted linguistic complexity. GPT-5 thinking model  
 137 exhibited extreme median perplexity ( $\sim 1000$ ), indicating highly unpredictable or complex language,  
 138 while other agents produced more moderate perplexity values.

139 Here is the summary of the characteristics described above. **GPT-5 thinking model:** High novelty,  
 140 high lexical diversity, low repetition, shorter utterances, and extreme perplexity. **ChatGPT Free:**  
 141 Longer utterances, higher question rate, more positive sentiment, stronger alignment with partner.  
 142 **Claude Opus 4.1:** Longest conversations, balanced linguistic style, moderate novelty and entropy.  
 143 **Gemini-2.0-flash:** Generally intermediate across most metrics, with less extreme behaviors.

144 We also visualized feature correlations in Figure 2. It revealed several groups of highly correlated vari-  
 145 ables ( $|r| > 0.85$ ), including (i) avg\_len, ttr, topical\_consistency, embedding\_variance,  
 146 novelty, and several lexical reuse measures, (ii) noun\_ratio, sentiment\_std, and  
 147 avg\_perplexity, and (iii) speaker\_balance\_A and speaker\_balance\_B.

### 148 3.2 Statistical Comparison and Feature Importance Analysis

149 KruskalWallis tests confirmed that many features differed significantly between agents (Table 1). The  
 150 most significant differences included topical consistency, novelty, average utterance length, number  
 151 of turns, noun ratio, lexical alignment, sentiment standard deviation, TTR, and verb ratio. Features  
 152 related to turn-taking balance, question-response adherence, and fixed conversational patterns were  
 153 not significantly different. Random forest analysis identified topical consistency, novelty, noun  
 154 ratio, number of turns, and lexical entropy as the most important features for distinguishing agents  
 155 (Figure 3). Lexical diversity, self-repetition rate, and alignment metrics also contributed substantially,  
 156 whereas pragmatic features such as speaker balance and question-response adherence had minimal

Table 1: Results of Kruskal-Wallis tests for all the features. The most significant differences included topical consistency, novelty, average utterance length, etc. Features related to turn-taking balance, question-response adherence, and fixed conversational patterns were not significantly different. Nominal p values are presented.

feature	stat	p value
topical_consistency	16.142857	0.001060
novelty	16.142857	0.001060
avg_len	14.325714	0.002494
num_turns	13.726408	0.003302
noun_ratio	13.422857	0.003806
lexical_alignment	13.064108	0.004500
sentiment_std	12.462857	0.005955
ttr	12.440000	0.006018
verb_ratio	12.165714	0.006837
lexical_entropy	12.028571	0.007286
lexical_convergence	11.891429	0.007764
self_repetition_rate	11.754286	0.008274
avg_perplexity	11.617143	0.008817
sentiment_mean	11.457143	0.009494
partner_reuse_rate	10.451429	0.015094
embedding_variance	10.177143	0.017119
avg_sentence_length	10.108571	0.017665
question_rate	9.100528	0.027984
sentiment_shift	6.508571	0.089325
pronoun_balance	3.439729	0.328664
pattern_rate	3.000000	0.391625
speaker_balance_A	0.782857	0.853563
speaker_balance_B	0.782857	0.853563
question_response_rate	0.000000	1.000000

importance. These results highlight that semantic, lexical, and stylistic properties are the strongest markers of agent-specific conversational behavior.

PCA using all features provided partial separation of agents (Figure 4 left). Restricting PCA to features with Kruskal-Wallis  $p < 0.05$  and random forest importance  $> 0.03$  enhanced clustering (Figure 4 right), suggesting that a subset of discriminative features captures the majority of agent-specific variation.

## 4 Discussions

This study systematically compared conversations generated by four large language models: ChatGPT Free, Gemini-2.0-flash, GPT-5 thinking model, and Claude Opus 4.1. Across 23 linguistic, semantic, and interactional features, we observed clear and consistent differences in conversation style, complexity, and lexical diversity: GPT-5 thinking model emphasized novelty, lexical diversity, and unpredictability, producing shorter, highly varied utterances with low self- and partner-repetition rates. ChatGPT Free favored longer utterances, higher question frequency, and positive sentiment, with strong alignment to conversational partners. Claude Opus 4.1 generated the longest conversations, with balanced syntactic and lexical properties, moderate novelty, and high lexical entropy. Gemini-2.0-flash remained intermediate across most metrics, without extreme tendencies in any feature.

These patterns were corroborated by statistical tests, feature importance rankings, and PCA visualizations, indicating that content-related (topical consistency, novelty, lexical entropy, etc.) and interactional (lexical alignment, partner reuse rate, etc.) features reliably distinguish agent-generated conversations, while more basic features such as speaker balance and question response rate do not differ significantly among different agents.

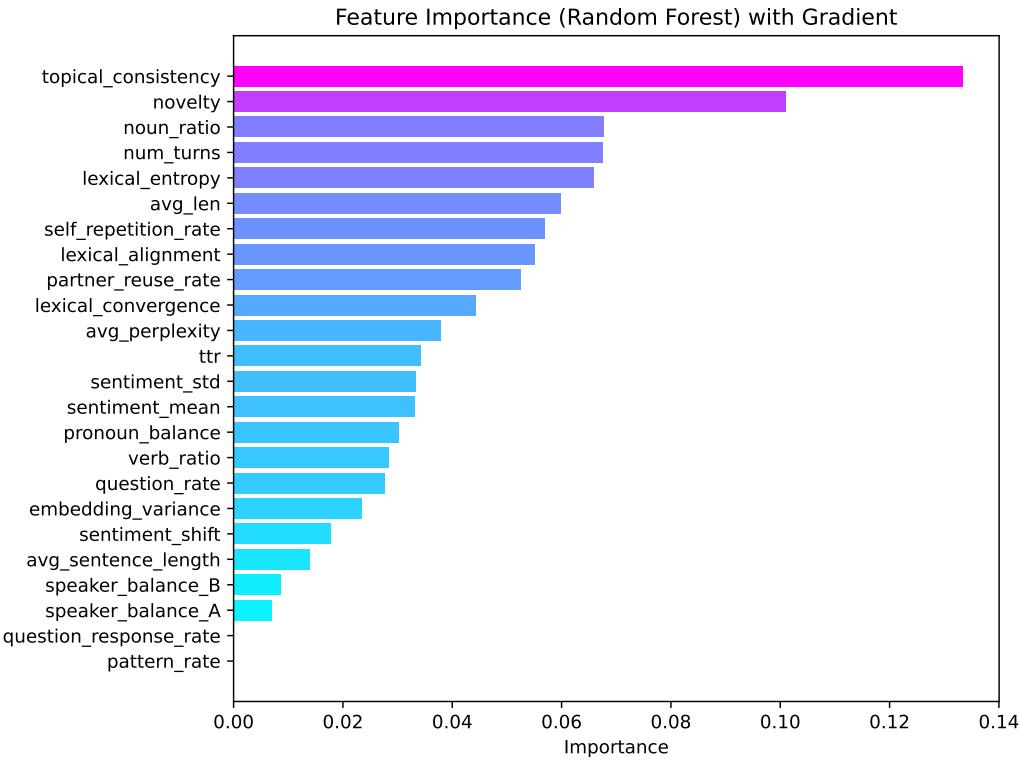


Figure 3: Important features identified by random forest analysis. Topical consistency, novelty, noun ratio, number of turns, and lexical entropy had the highest importance.

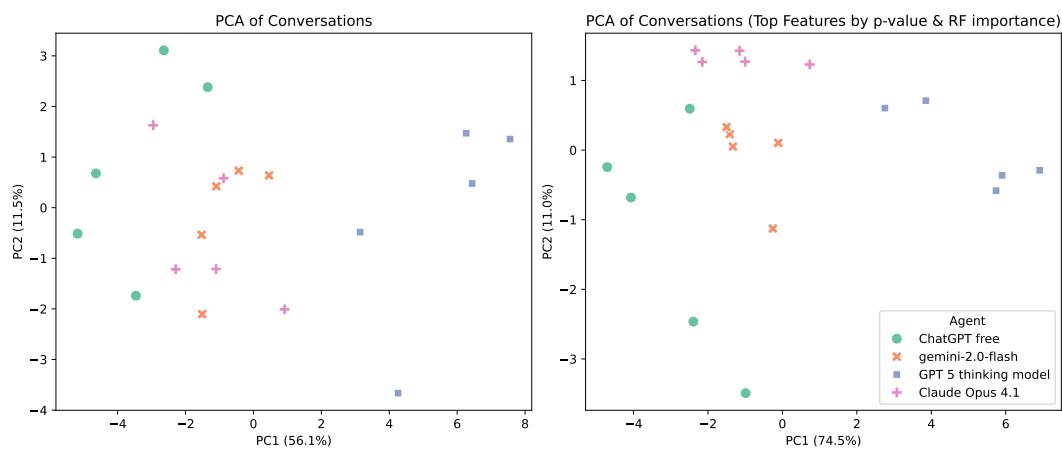


Figure 4: PCA using all features (left) and pre-selected features (right). The right PCA has an enhanced clustering compared to the left PCA.

179 Our results provide several key insights. Firstly, even under controlled prompts, LLMs exhibited  
180 distinct conversational fingerprints. This suggests that conversation-level metrics can serve as  
181 quantitative probes of model behavior. Secondly, features such as topical consistency, novelty, lexical  
182 entropy, and noun ratio were the most informative for distinguishing agents, whereas turn-taking and  
183 question-response adherence were uniform. This implies that traditional interaction metrics alone are  
184 insufficient to capture qualitative differences in AI-generated conversations. Thirdly, GPT-5s highly  
185 novel but unpredictable utterances highlight a trade-off between creativity and coherence. In contrast,  
186 ChatGPT and Claude maintain more predictable, partner-aligned responses, suggesting different  
187 optimization priorities in model training.

188 Our approach in this study systematically combines multi-turn conversation generation, quantitative  
189 feature extraction, and statistical discrimination across multiple LLMs. By providing both lexical-  
190 semantic measures and interactional metrics, this study offers a reproducible methodology for cross-  
191 agent comparison, insights into agent-specific stylistic tendencies, and a framework for evaluating  
192 novelty, complexity, and alignment in synthetic dialogues. These contributions are relevant for  
193 researchers seeking to benchmark LLMs beyond traditional metrics, and for practitioners designing  
194 conversational AI that meets specific interactional goals.

195 Several limitations should be noted:

- 196 • Dataset Size: Each agent generated only five conversations per prompt, which may limit  
197 generalizability. Future work could scale this to dozens or hundreds of conversations.
- 198 • Controlled Prompting: While standardized prompts enabled comparability, real-world usage  
199 involves more diverse inputs, which may alter agent behavior.
- 200 • Feature Scope: Although 23 features capture lexical, syntactic, and semantic properties, other  
201 aspects such as pragmatic reasoning, humor, or subtle discourse cues were not measured.
- 202 • Perplexity Interpretation: Extreme GPT-5 perplexity may reflect model idiosyncrasies rather  
203 than true communicative complexity; further exploration is warranted.
- 204 • Redundant Features: This was an exploratory study and there were many correlated features.  
205 These redundancy should be reduced in future studies for further analysis.

206 Future research could extend this framework by including human evaluation, longer dialogue horizons,  
207 and task-oriented scenarios to probe model behavior in more ecologically valid settings.

208 In conclusion, this study demonstrates that large language models produce quantifiably distinct  
209 conversation styles, with measurable differences in novelty, lexical diversity, sentiment, and inter-  
210 actional alignment. Our framework combining feature extraction, statistical testing, and machine  
211 learning-based discrimination provides a reproducible and interpretable approach for evaluating  
212 multi-turn AI dialogues. By highlighting the trade-offs between creativity, coherence, and alignment,  
213 these findings inform both LLM evaluation and design of more human-like conversational agents.

## 214 Reproducibility Statement

215 We have placed all the data and codes to reproduce this study in the Google Colaboratory:  
216 [https://colab.research.google.com/drive/1JobFFniam7HXbCmIiBgLuLC8t5dGFTct?](https://colab.research.google.com/drive/1JobFFniam7HXbCmIiBgLuLC8t5dGFTct?usp=sharing)  
217 Readers can reproduce the whole analysis results without any environment building  
218 on their machines.

## 219 References

- 220 [1] Ebubechukwu, I. et al. (2024) Dialogue you can trust: Human and AI perspectives on generated conversations.  
221 arXiv: 2409.01808
- 222 [2] Xu, S. et al. (2024) Large language model based situational dialogues for second language learning. arXiv:  
223 2403.20005
- 224 [3] Guan, S. et al. (2025) Evaluating LLM-based agents for multi-turn conversations: A survey. arXiv:  
225 2503.22458
- 226 [4] Rzadeczka, M. et al. (2025) The efficacy of conversational AI in rectifying the theory-of-mind and autonomy  
227 biases: comparative analysis. *JMIR Mental Health.* 12:e64396.

228 **Agents4Science AI Involvement Checklist**

- 229     1. **Hypothesis development:** Hypothesis development includes the process by which you  
230       came to explore this research topic and research question. This can involve the background  
231       research performed by either researchers or by AI. This can also involve whether the idea  
232       was proposed by researchers or by AI.

233       Answer: **[C]**

234       Explanation: Human 1st author asked AI 1st author what research they can do without  
235       external data. AI 1st author replied multiple possibilities, among which was AI agents  
236       repeating talks. Human 1st author elaborated this into the current research question.

- 237     2. **Experimental design and implementation:** This category includes design of experiments  
238       that are used to test the hypotheses, coding and implementation of computational methods,  
239       and the execution of these experiments.

240       Answer: **[D]**

241       Explanation: Human 1st author prompted Chat GPT Free version and Gemini-2.0-flash for  
242       data generation. Human 2nd author prompted GPT-5 thinking model for data generation.  
243       Human 3rd author prompted Claude Opus 4.1 for data generation. AI 1st author suggested  
244       the list of features to evaluate. Human 1st author prompted AI 1st author to include more  
245       features, and AI 1st author suggested additional features. Visualization, Statistical analysis,  
246       and machine learning-based evaluation were all suggested by AI 1st author.

- 247     3. **Analysis of data and interpretation of results:** This category encompasses any process to  
248       organize and process data for the experiments in the paper. It also includes interpretations of  
249       the results of the study.

250       Answer: **[D]**

251       Explanation: AI 1st author wrote original codes to analyze data according to the study  
252       design. Human 1st author ran the code to generate the results. Human 1st author sent  
253       numerical results to AI 1st author and AI 1st author interpreted the results. Human 1st  
254       author sent box plot image to AI 1st author and AI 1st author described the results. Human  
255       1st author reviewed the AI 1st author's interpretation but only fixed obvious mistakes (e.g.,  
256       higher/lower).

- 257     4. **Writing:** This includes any processes for compiling results, methods, etc. into the final  
258       paper form. This can involve not only writing of the main text but also figure-making,  
259       improving layout of the manuscript, and formulation of narrative.

260       Answer: **[C]**

261       Explanation: AI 1st author drafted the Results, Discussions, and Methods section. Human  
262       1st author converted them into L<sup>A</sup>T<sub>E</sub>X template compatible style. Human 1st author prompted  
263       AI 1st author to improve the figure layout. AI 2nd author did literature search for Introduction  
264       section prompted by Human 1st author. Human 1st author used the information to write  
265       Introduction with AI 1st author.

- 266     5. **Observed AI Limitations:** What limitations have you found when using AI as a partner or  
267       lead author?

268       Description: In this study, we tried to have AI agents do as much part of research as  
269       possible. Although it allowed us to generate research paper rapidly, there were several  
270       notable limitations. Firstly, hypothesis generation by AI agents strongly depended on human  
271       prompts. Making good prompts felt harder than making relevant hypotheses without AI  
272       agents. This was also true for study design. In both cases, if prompts were not specific, AI  
273       agents tended to give very standard suggestions. Also, there were some obvious mistakes in  
274       the result interpretations, such as high vs low values. Also, AI agents sometimes generated  
275       result sentences that we did not input. For example, the PCA result interpretations were  
276       generated before we input actual results.

277 **Agents4Science Paper Checklist**

278 **1. Claims**

279 Question: Do the main claims made in the abstract and introduction accurately reflect the  
280 paper's contributions and scope?

281 Answer: [Yes]

282 Justification: The claims made in the abstract and introduction are concretely addressed in  
283 the Results section.

284 Guidelines:

- 285 • The answer NA means that the abstract and introduction do not include the claims  
286 made in the paper.
- 287 • The abstract and/or introduction should clearly state the claims made, including the  
288 contributions made in the paper and important assumptions and limitations. A No or  
289 NA answer to this question will not be perceived well by the reviewers.
- 290 • The claims made should match theoretical and experimental results, and reflect how  
291 much the results can be expected to generalize to other settings.
- 292 • It is fine to include aspirational goals as motivation as long as it is clear that these goals  
293 are not attained by the paper.

294 **2. Limitations**

295 Question: Does the paper discuss the limitations of the work performed by the authors?

296 Answer: [Yes]

297 Justification: We have discussed limitations in the Discussions section.

298 Guidelines:

- 299 • The answer NA means that the paper has no limitation while the answer No means that  
300 the paper has limitations, but those are not discussed in the paper.
- 301 • The authors are encouraged to create a separate "Limitations" section in their paper.
- 302 • The paper should point out any strong assumptions and how robust the results are to  
303 violations of these assumptions (e.g., independence assumptions, noiseless settings,  
304 model well-specification, asymptotic approximations only holding locally). The authors  
305 should reflect on how these assumptions might be violated in practice and what the  
306 implications would be.
- 307 • The authors should reflect on the scope of the claims made, e.g., if the approach was  
308 only tested on a few datasets or with a few runs. In general, empirical results often  
309 depend on implicit assumptions, which should be articulated.
- 310 • The authors should reflect on the factors that influence the performance of the approach.  
311 For example, a facial recognition algorithm may perform poorly when image resolution  
312 is low or images are taken in low lighting.
- 313 • The authors should discuss the computational efficiency of the proposed algorithms  
314 and how they scale with dataset size.
- 315 • If applicable, the authors should discuss possible limitations of their approach to  
316 address problems of privacy and fairness.
- 317 • While the authors might fear that complete honesty about limitations might be used by  
318 reviewers as grounds for rejection, a worse outcome might be that reviewers discover  
319 limitations that aren't acknowledged in the paper. Reviewers will be specifically  
320 instructed to not penalize honesty concerning limitations.

321 **3. Theory assumptions and proofs**

322 Question: For each theoretical result, does the paper provide the full set of assumptions and  
323 a complete (and correct) proof?

324 Answer: [NA]

325 Justification: The paper does not include theoretical results.

326 Guidelines:

- 327 • The answer NA means that the paper does not include theoretical results.

- 328           • All the theorems, formulas, and proofs in the paper should be numbered and cross-  
329           referenced.  
330           • All assumptions should be clearly stated or referenced in the statement of any theorems.  
331           • The proofs can either appear in the main paper or the supplemental material, but if  
332           they appear in the supplemental material, the authors are encouraged to provide a short  
333           proof sketch to provide intuition.

334          **4. Experimental result reproducibility**

335          Question: Does the paper fully disclose all the information needed to reproduce the main ex-  
336          perimental results of the paper to the extent that it affects the main claims and/or conclusions  
337          of the paper (regardless of whether the code and data are provided or not)?

338          Answer: [Yes]

339          Justification: The full data and code are made open on Google Colaboratory. Therefore, all  
340          the results can be reproduced by simply running the cells in Google Colaboratory.

341          Guidelines:

- 342           • The answer NA means that the paper does not include experiments.  
343           • If the paper includes experiments, a No answer to this question will not be perceived  
344           well by the reviewers: Making the paper reproducible is important.  
345           • If the contribution is a dataset and/or model, the authors should describe the steps taken  
346           to make their results reproducible or verifiable.  
347           • We recognize that reproducibility may be tricky in some cases, in which case authors  
348           are welcome to describe the particular way they provide for reproducibility. In the case  
349           of closed-source models, it may be that access to the model is limited in some way  
350           (e.g., to registered users), but it should be possible for other researchers to have some  
351           path to reproducing or verifying the results.

352          **5. Open access to data and code**

353          Question: Does the paper provide open access to the data and code, with sufficient instruc-  
354          tions to faithfully reproduce the main experimental results, as described in supplemental  
355          material?

356          Answer: [Yes]

357          Justification: The full data and code are made open on Google Colaboratory. Therefore, all  
358          the results can be reproduced by simply running the cells in Google Colaboratory.

359          Guidelines:

- 360           • The answer NA means that paper does not include experiments requiring code.  
361           • Please see the Agents4Science code and data submission guidelines on the conference  
362           website for more details.  
363           • While we encourage the release of code and data, we understand that this might not  
364           be possible, so No is an acceptable answer. Papers cannot be rejected simply for not  
365           including code, unless this is central to the contribution (e.g., for a new open-source  
366           benchmark).  
367           • The instructions should contain the exact command and environment needed to run to  
368           reproduce the results.  
369           • At submission time, to preserve anonymity, the authors should release anonymized  
370           versions (if applicable).

371          **6. Experimental setting/details**

372          Question: Does the paper specify all the training and test details (e.g., data splits, hyper-  
373          parameters, how they were chosen, type of optimizer, etc.) necessary to understand the  
374          results?

375          Answer: [Yes]

376          Justification: All the data generation processes are explained in the Methods section.

377          Guidelines:

- 378           • The answer NA means that the paper does not include experiments.

- 379           • The experimental setting should be presented in the core of the paper to a level of detail  
380           that is necessary to appreciate the results and make sense of them.  
381           • The full details can be provided either with the code, in appendix, or as supplemental  
382           material.

383      **7. Experiment statistical significance**

384      Question: Does the paper report error bars suitably and correctly defined or other appropriate  
385      information about the statistical significance of the experiments?

386      Answer: [Yes]

387      Justification: We conducted appropriate statistical tests and visualizations to support our  
388      results.

389      Guidelines:

- 390           • The answer NA means that the paper does not include experiments.  
391           • The authors should answer "Yes" if the results are accompanied by error bars, confi-  
392           dence intervals, or statistical significance tests, at least for the experiments that support  
393           the main claims of the paper.  
394           • The factors of variability that the error bars are capturing should be clearly stated  
395           (for example, train/test split, initialization, or overall run with given experimental  
396           conditions).

397      **8. Experiments compute resources**

398      Question: For each experiment, does the paper provide sufficient information on the com-  
399      puter resources (type of compute workers, memory, time of execution) needed to reproduce  
400      the experiments?

401      Answer: [Yes]

402      Justification: Computer resource information can be found in Methods section.

403      Guidelines:

- 404           • The answer NA means that the paper does not include experiments.  
405           • The paper should indicate the type of compute workers CPU or GPU, internal cluster,  
406           or cloud provider, including relevant memory and storage.  
407           • The paper should provide the amount of compute required for each of the individual  
408           experimental runs as well as estimate the total compute.

409      **9. Code of ethics**

410      Question: Does the research conducted in the paper conform, in every respect, with the  
411      Agents4Science Code of Ethics (see conference website)?

412      Answer: [Yes]

413      Justification: Human authors reviewed the Code of Ethics.

414      Guidelines:

- 415           • The answer NA means that the authors have not reviewed the Agents4Science Code of  
416           Ethics.  
417           • If the authors answer No, they should explain the special circumstances that require a  
418           deviation from the Code of Ethics.

419      **10. Broader impacts**

420      Question: Does the paper discuss both potential positive societal impacts and negative  
421      societal impacts of the work performed?

422      Answer: [Yes]

423      Justification: In the Discussions section, we discussed the implication of our study in how  
424      we should understand LLMs' conversations.

425      Guidelines:

- 426           • The answer NA means that there is no societal impact of the work performed.  
427           • If the authors answer NA or No, they should explain why their work has no societal  
428           impact or why the paper does not address societal impact.

- 429
- Examples of negative societal impacts include potential malicious or unintended uses
- 430 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations,
- 431 privacy considerations, and security considerations.
- 432
- If there are negative societal impacts, the authors could also discuss possible mitigation
- 433 strategies.